

Betriebssysteme

Alle Tutoriumsfolien

Peter Bohner

6. Februar 2025

ITEC - Operating Systems Group

Abstraction / Standardization

- Devices and how to talk to them differ greatly
⇒ Remember the „Driver-CDs“ shipped with motherboards back in the days so you could properly configure and run things?
 - Strange user demands: Programs *should* run on more than one hardware configuration
- ⇒ Abstract away hardware details!

Resource management

- You want to print? Too bad, another program is already using that printer.
- You want to access the storage drive? Too bad, another program is already doing that.
- You want to get CPU time? Too bad, this `while (true)` loop is more important.
- ...

Security and Protection


- You are a good citizen and use a password manager. What could happen without your helpful OS? \Rightarrow Other programs may read its memory!
- You copied a password to your clipboard? \Rightarrow Ooops. You're on your own there. The Clipboard is not provided by the OS and mostly has no special protections.
- You write a cool little program that fills a buffer with a random value. Sadly you made a mistake and missed a bounds check. What happens? You crash, but your text editor doesn't suddenly have its memory overwritten!
- ...

Provide an execution environment for applications

- What does that mean? Basically all of the above combined and more. Make a homely place where applications like to live!

Kernel and user mode

What are the differences between a processor running in kernel or user mode?
Why are both modes needed?



[Wikipedia - Hertzprung](#)

What are the differences between a processor running in kernel or user mode?
Why are both modes needed?

- In kernel mode you have full access to privileged instructions.

Great, but what is a privileged instruction?

- Change control registers for memory mappings \Rightarrow Read other process's memory
- Disable / Enable interrupts \Rightarrow No preemption for you.
- Access platform devices (network card, storage, printer,...)
- Some nice registers: **LGDT** (Load Global Descriptor Table) or the **LLDT**, **INVD** (Invalidate cache), **HLT** (Halt processor!)
- Is **MOV** (Move) a privileged instruction? Yes, if moving to debug/control registers!

C - Basics

Geschichte

- Wie auch der Vorgänger von und mit Dennis Ritchie bei Bell Labs um 1972 rum entwickelt
- Ist der Nachfolger von B
- Ist eine der meistgenutzten Sprachen

Eigenschaften

- Imperativ, Prozedural, *nicht* Objektorientiert
- Low level
- Manuelle Speicherverwaltung, Spaß mit Pointern und allen möglichen Implementierungsdetails
- CVE-Factory (cough <https://kitctf.de/> cough)



```
1 #include <stdio.h>
2 #include "World.c"
3
4 int computeAnswer();
5
6 int main() {
7     printf("Hey, your answer is %d\n", computeAnswer());
8
9     return 0;
10 }
11
12 int computeAnswer() {
13     return answerMeWorld();
14 }
```

Primitive

Name	Minimale Größe in Bytes	Größe bei mir in Bytes
char	1	1
short	2	2
int	2	4
float		4
double		8
<hr/>		
long int	4	8
long long int	8	8
<hr/>		
signed int	2	4
unsigned int	2	4

Oder mit bestimmter Größe

```
#include <inttypes.h>
```

```
int8_t a; int16_t b; int32_t c; int64_t d;
```

Wie sieht ein boolean in C aus?

Gibt es erst seit C99. `0` ist **false**, alle anderen Zahlenwerte sind **true**

Was macht man da also? Selbst basteln (oder `0` / `1` nutzen)

```
1 // #include <stdbool.h>
2
3 typedef unsigned char bool;
4 #define false 0
5 #define true 1
6
7 int main() {
8     bool hey = true;
9     return 0;
10 }
```

```
1 int array[5] = {1, 2, 3, 4, 5};  
2 int array[] = {1, 2, 3, 4, 5};  
3 int array[5];
```

Was ist das?

- Ein zusammenhängender Speicherbereich
- Größe ist Teil des *Typs* (**sizeof** geht, wenn Typ bekannt)
- Wie bekommt man die Größe zur *Laufzeit*? \Rightarrow **sizeof(array)**? *Heh.*
GAR NICHT! Nur im Typ präsent.

```
1 char[] myString = "world";  
2 char[] myString = { 'w', 'o', 'r', 'l', 'd'};, '\0'};
```

Was sind strings?

- Ein `char`-Array mit *Null-Terminator* \Rightarrow Ein Byte länger!
- Nutzen den ASCII code

Die Sprache C - Structs

```
1 #include <stdbool.h>
2 #include <stdio.h>
3 struct User {
4     bool userIsGoat;
5     int age;
6     char name[20];
7 };
8 int main() {
9     struct User pete = { .userIsGoat = false, .age = 20, .name = "Pete" };
10    printf(
11        "I have a user %s (goaty: %s) of age %u\n",
12        pete.name,
13        pete.userIsGoat ? "true" : "false",
14        pete.age
15    );
16    return 0;
17 }
```


Was sind Structs?

- Komplexere Datentypen
- Zusammengesetzt aus anderen Datentypen
- **KEINE KLASSE**. Warum? *Keine Vererbung*, keine Methoden

Die Sprache C - Structs

```
1 #include <stdbool.h>
2 #include <stdio.h>
3 typedef struct User {
4     bool userIsGoat;
5     int age;
6     char name[20];
7 } User_t;
8 int main() {
9     User_t peter = { .userIsGoat = false, .age = 20, .name = "Peter" };
10    printf(
11        "I have a user %s (goaty: %s) of age %u\n",
12        peter.name,
13        peter.userIsGoat ? "true" : "false",
14        peter.age
15    );
16    return 0;
17 }
```

Die Sprache C - Call-By-Value

```
1 #include <stdio.h>
2 typedef struct User { int age; } User;
3
4 void foo(User user) {
5     user.age = 200;
6     printf("User age in foo is %d\n", user.age);
7 }
8
9 int main() {
10     User user = { .age = 20 };
11     foo(user);
12     printf("User age in main %d\n", user.age);
13     return 0;
14 }
```

Output?

```
User age in foo is 200
User age in main 20
```

Die Sprache C - Call-By-Reference

```
1 #include <stdio.h>
2 typedef struct User { int age; } User;
3
4 void foo(User* user) {
5     user->age = 200;
6     printf("User age in foo is %d\n", user->age);
7 }
8
9 int main() {
10     User user = { .age = 20 };
11     foo(&user);
12     printf("User age in main %d\n", user.age);
13     return 0;
14 }
```

Output?

```
User age in foo is 200
User age in main 200
```

Die Sprache C - Pointer

```
1 #include <stdio.h>
2 void swap(int* a, int* b) {
3     int tmp = *a;
4     *a = *b;
5     *b = tmp;
6 }
7 int main() {
8     int first  = 20;
9     int second = 5;
10    printf("Before swap - First: %d, Second: %d\n", first, second);
11    swap(&first, &second);
12    printf("After swap  - First: %d, Second: %d\n", first, second);
13    return 0;
14 }
```

Output?

Before swap - First: 20, Second: 5

After swap - First: 5, Second: 20

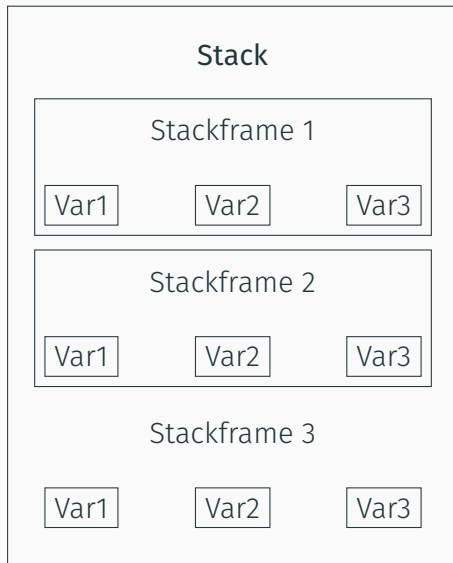
Programmbereiche für Variablen

Data Globale Variablen

Stack Lokale Variablen

Heap Explizit zur Laufzeit angeforderter Speicher

...



Speicher zur Laufzeit allokieren

```
1 int main(int argc, char** argv) {  
2     char** stringsCopy[??];  
3     return 0;  
4 }
```


Speicher zur Laufzeit allokieren

```
1 #include <stdlib.h>
2
3 int main(int argc, char** argv) {
4     char** stringsCopy = malloc(sizeof(char*) * argc);
5     for(int i = 0; i < argc; i++) {
6         stringsCopy[i] = argv[i]; // shallow copy
7     }
8     free(stringsCopy);
9     return 0;
10 }
```

Speicher zur Laufzeit allokieren

```
1 #include <stdbool.h>
2 #include <stdlib.h>
3 #include <string.h>
4 struct User { bool userIsGoat; int age; char name[20]; };
5 int main() {
6     struct User* pete = malloc(sizeof(struct User));
7
8     (*pete).userIsGoat = false;
9     pete->age = 20;
10    strcpy((*pete).name, "Pete");
11
12    free(pete);
13    return 0;
14 }
```

Warum?

- Funktionen müssen *vor dem Aufruf* **deklariert** sein
- *Implementierung* kann auch irgendwann später erfolgen
- Implementierung *zur Entwicklung* nicht notwendig!

```
1  int halfTruths(); // Deklaration
2
3  int main() {
4      return halfTruths();
5  }
6
7  int halfTruths() { // Implementierung
8      return 21;
9  }
```

Implementierung zur *Entwicklung* nicht notwendig!

- \Rightarrow Header files
- Aufsplitten in Deklarationen und Definitionen
- Nutzer bindet Header zum Programmieren ein, *linkt* dann gegen eine Implementierung

Listing 1: World.h

```
1 int answerMeWorld();
```

Listing 2: World.c

```
1 int answerMeWorld() {  
2     return 42;  
3 }
```

```
1 #include <stdio.h>
2 #include "World.c"
3
4 int computeAnswer();
5
6 int main() {
7     printf("Hey, your answer is %d\n", computeAnswer());
8
9     return 0;
10 }
11
12 int computeAnswer() {
13     return answerMeWorld();
14 }
```

What does that acronym even mean?

Application Binary Interface

And what does that specify?

- Interface of *binary* programs (i.e. *after* compilation)
- Instruction Set (e.g. x86, ARM)
- Calling convention (e.g. `cdecl` or `System V AMD64 ABI`)
- Basic data types and their size / alignment (`int`, `sizeof(int)`)
- How to perform System Calls

You might also know the term „ABI of a library“?

This is used when talking about *binary compatibility* of different library versions.

⇒ Do you need to recompile your code against the new version?

What's the usual calling convention on modern Linux?

System V AMD64 ABI

<https://godbolt.org/z/68xexn>

- Integer arguments in `rdi`, `rsi`, `rdx`, `rcx`, `r8`, `r9`, then stack
- FP arguments in `xmm0` to `xmm7`
- Integer return value in `rax`, `rdx`

Static and Dynamic Linking

What is the difference between static and dynamic linking?

- A static library is a collection of object files
⇒ The linker can treat it as normal code
- A dynamic library is loaded and linked *at runtime*

What are advantages of static / dynamic linking?

- S+ Unused references can be elided
- S+ Library calls just as fast as local ones
- S+ No runtime overhead for loading and relocation
- S- Library can not be shared ⇒ Memory overhead
- D+ Library (Code segment) can be shared

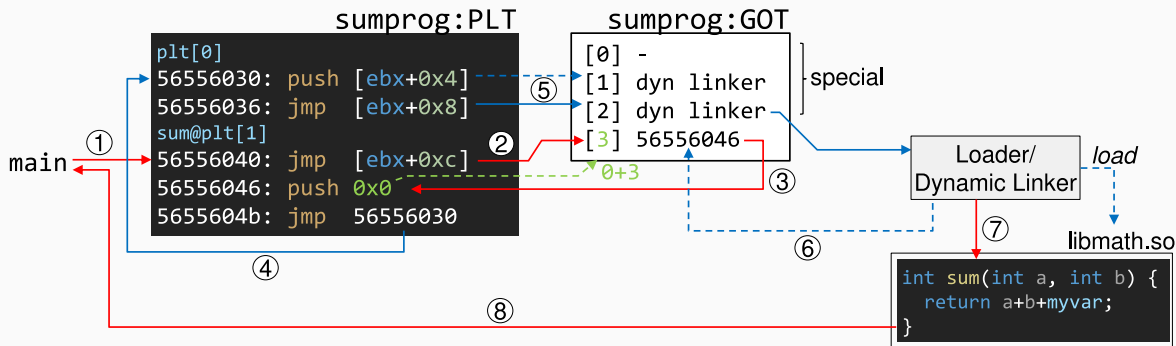
What does PIC stand for?

Position Independent Code

Independent to *what*?

It's position in the process' address space.

- Shared libraries are loaded *somewhere* in the address space of a process
- PIC uses *relative* addresses (to what?)(to the instruction pointer)only – and can therefore be loaded anywhere
- How do you find *global* symbols then? \Rightarrow GOT, the **G**lobal **O**ffset **T**able



Quelle: Vorlesungsfolien

Multi-Programming

What is that?

Allow multiple processes to run *concurrently* (in parallel?)

How can they share a CPU? How can you make them take turns?

Two possible broad categories: **Cooperative Multitasking** and **Preemption**

How does Cooperative Multitasking work? Do you know any system using it?

- Processes **voluntarily** yield control so another one can take over
- When? Typically when they are blocking on some I/O or have finished their job for now
- Do you know any language or system using that concept? e.g. **await** in JavaScript or Python
- Can you think of any problems? Any advantages?
 - A bad actor or buggy program can bring the whole system to its knees!
 - It can lead to easier programs, as you can just assume you are never interrupted ⇒ Usage on small resource-constrained embedded systems

How does Preemption work?

- Processes get periodically interrupted by **Timer-interrupts**
- This transfers control to the OS, which can choose to schedule another process instead

But why even bother with this?

- Process A runs while Process B waits for an I/O device to complete its task
- ⇒ Better CPU and I/O utilization if you have a mixed workload
- What do the I/O devices and CPU need to support for this to be an improvement?
- ⇒ Interrupts, else the CPU needs to *continuously poll*

Are there any drawbacks?

- More complex
- Protection. How do you separate them from each other?
- Scheduler: Fairness and accounting

What are CPU and IO-bound Processes?

- CPU-bound: Process that rarely invokes I/O, unlikely to block, likes your CPU very much
- I/O-bound: Often blocks to wait for I/O, performs short CPU-bursts in between

Address Spaces

Why can't a process read another process's memory?

- Let's have a look at two processes...
- Address `0xFF` in Process A is *not* the same as Address `0xFF` in Process B

⇒ „If you can't name it, you can't touch it“

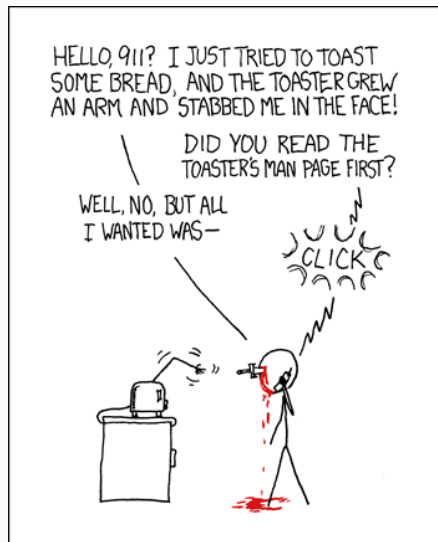
- You will learn about this in way more detail in later lectures

Fork



What does **fork** do?

- Creates a mostly identical child process
 - Same address space layout
 - Same data (as if it was all copied over)
 - Share open file descriptors
 - But the child has its own PID and address space
 - ... \Rightarrow **man fork**



Life is too short for man pages, and occasionally much too short without them.

Let's fork!

- Print `Fork failed` if the fork failed
- Print `I am the child` in the child
- Print `I am the parent of <child PID>` in the parent
- Wait for the child to die peacefully and print when it is dead in the parent

Let's fork! (More official solution)

```
1 #include <stdio.h>
2 #include <unistd.h>
3 #include <sys/wait.h>
4 #include <sys/types.h>
5
6 int main() {
7     pid_t pid;
8     switch( (pid = fork()) ) {
9         case -1:
10             printf( "Error. Fork failed\n" );
11             break;
12         case 0:
13             printf( "I am the child!\n" );
14             break;
15         default: // pid > 0
16             printf( "I am the parent!\nChild PID is %d\n", pid );
17             wait( NULL ); // wait and ignore the result
18             printf( "Child terminated\n" );
19     }
20     return 0;
21 }
```

Is `fork` sufficient to create a small shell?

- No, you also need to *load the other program* in the child

⇒ `exec(vp|v|...)`

Interrupts



What three events lead to an invocation of the kernel?

- Interrupts
- Syscalls
- Exceptions

What is an Interrupt?

- A signal generated *outside* the processor (typically from an I/O device)

What is an Exception?

- Caused by the CPU
- Inform the operating system of an error condition during execution
- E.g. **division by zero** or page faults

What is a System Call?

- Explicit call from user application into OS
- Used to interact with the OS (e.g. to read a file)

A simple categorisation

	Voluntary	Involuntary
Sync	System call	Exception
Async	???	Interrupt

Can you think of an exception the OS can't repair?

i.e. it will terminate the process :(

- Divide by Zero
- The wrong kind of page fault (access forbidden or invalid memory)

Can you think of an exception the OS *can* repair?

- The right kind of page fault (e.g. access memory currently paged out to disk)
- When the OS is run in a virtual environment, it might try to access privileged instructions

⇒ CPU throws an exception

⇒ Virtual machine monitor handles that and executes the action on the real physical device (if allowed)

How is the **trap** instruction related to system calls?

1. **trap** is executed and switches the CPU to kernel mode
 2. The CPU continues execution at a hardware dependent location (where exactly depends on your architecture and a few other things, but the important part is that the operating system registered itself there)
- ⇒ **trap** can be seen as a low-level instruction making a higher-level mechanism (syscalls) possible
- You could also synchronously trap into the kernel using a division by zero or other fun stuff. That isn't done normally though...

What is a system call number? Why do you need them?

The user-level application can't just call a kernel function!

⇒ The kernel somehow needs to dispatch to the correct function

⇒ Array of function pointers with system call numbers as indices

How can you pass parameters to the kernel when executing a syscall?

- Registers (used by Linux, see `man 2 syscall`)
- Stack (used by Windows)
- Main memory (and location in register)

How do syscalls relate to libraries?

Abstraction? How do you call them?

- Syscall mechanism is usually hidden behind library functions (**libc** on linux)
- The functions take care of preparing arguments, setting the syscall number and trapping in an efficient way
- Windows uses the **ntdll.dll** / **Win32 API**

What do we need to keep an eye on when performing a system call in the kernel?

- Validate all parameters! Twice! And a third time to be sure.
- Otherwise you might have some arbitrary reads/writes into system memory

Scheduling basics

What is a Scheduler? Why do we even need it?

- Maps processes to resources
- Ideally: Every process gets what it needs some day

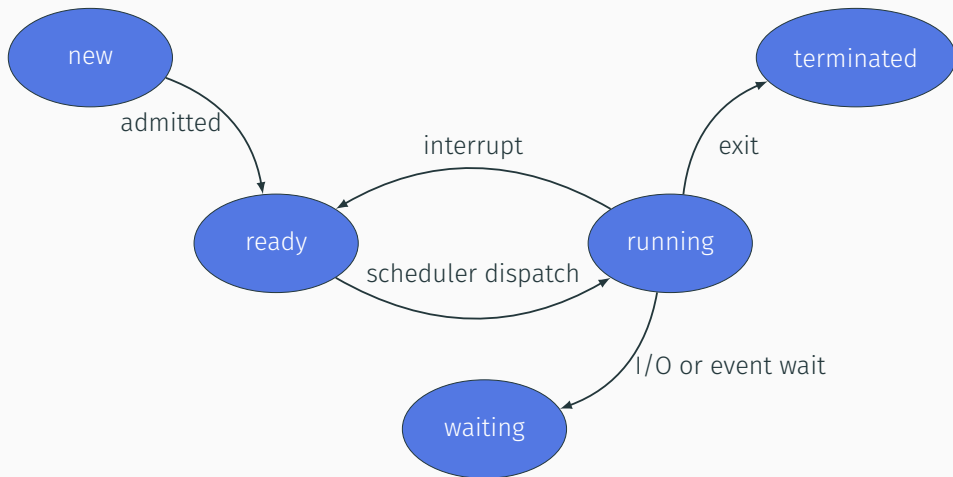
What Schedulers do you know?

- CPU-Scheduler: The classic
- Disk-Scheduler: Why have one? Multiplexing but also efficiency!
- Network I/O: When to send packets, which packets to drop, QoL,...

What are the differences? When are they used?

- LTS: Decide which processes to put in the *run queue*
 - STS: Decide which process runs on the *CPU*
 - MTS: Temporarily removes processes from main memory (and e.g. writes them out to disk)
- ⇒ Reduce degree of multiprogramming, make room in memory (and a few other reasons)

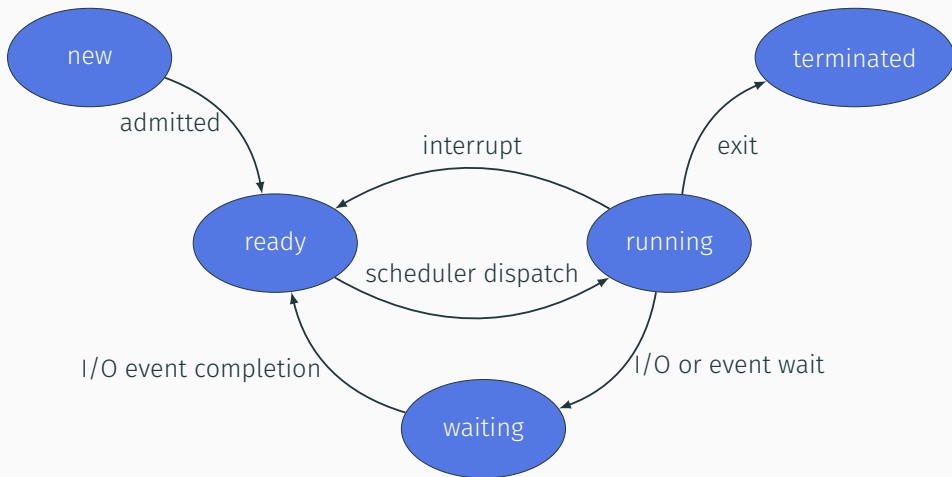
Process states



„I/O or event wait“? When does a process move from ready to waiting?

- Network / Disk I/O
- Mutex or other inter-process synchronisation
- Sleepyness

Process states



What makes a good Scheduler good?

Let's play scheduler!

Some metrics

- Processor utilization: Percentage of working time
- Throughput: How many jobs do you finish?
- Turnaround time: Wallclock-time from submission to finish
- Waiting time: How long did it spend in the ready queue
- Response time: Time between submission of a request and first response (e.g. key press to echo on screen)

What does your hardware need to support to allow non-cooperative scheduling?

Timer Interrupts! Waiting for a cosmic ray to hit, a network package to arrive, a system call or any other random interrupt gets old fast :)

Scheduling - When to interrupt

Any guesses for how long a timeslice usually is?

2ms - 200ms

- On windows it depends on the configuration (favor foreground / background processes)

Which setting has the longer timeslice?

20ms to 30ms for foreground, 180ms to 200ms for background

- Linux's „Completely Fair Scheduler“ adjusts them dynamically based on the priority, number of processes, ...

Benefits of shorter/longer timeslices?

- Short: High interactivity, higher overhead
- Long: Lower interactivity, smaller overhead

Pitfalls - How would you implement this?

- You would need to have future knowledge to figure out the job length!
How do you solve this?
- **Predict** the *future* based on *past* behaviour
- Does this work?

„THERE IS AS YET INSUFFICIENT DATA FOR A MEANINGFUL ANSWER.“

~ Isaac Asimov, *„The Last Question“* (Comic)


You need some balanced initial value. Not *that* big of a deal with preemption though. Why?

Interrupt the process after the estimated time is over.

What is priority scheduling? Why would you use it?

- Each process is assigned a priority
- The process with the highest priority is chosen

Multi-Level Feedback Queues



How it works

- All processes start in the highest queue
 - When they use up their timeslice and are preempted, they descend
 - If they block before, they stay in the level (optionally: Are moved up)
- ⇒ I/O bound processes rise to the top and react quickly, CPU bound processes get longer timeslices but less often

A Flawless Scheduling Algorithm?

Consider the waiting time

- A few I/O-bound jobs could saturate the CPU!
- ⇒ The lower-level processes starve

How could you fix this?

- E.g. reset the whole thing after a given interval, so all start in the highest level again
- "Boost" processes that waited for a long time

What metrics does it optimize?

- Utilization? Turnaround time? Throughput? Waiting time? Response time?
- Prefer I/O bound, prefer short jobs, group the rest based on their needs

Write an insert function

Task

```
1 void insert(int *a, size_t *len, int z);  
2  
3 void main() {  
4     int a[10] = {0, 1, 2, 3, 5, 6, 7, 8, 9};  
5     size_t len = 9;  
6     insert(a, &len, 4);  
7     assert(a[4] == 4);  
8     assert(len == 10);  
9 }
```

What happens when we insert another?

We write over the array boundary!

What methods does C offer to copy memory?

memcpy

- + Potentially more efficient
- Can not handle overlaps

memmove

- + *Can* handle overlaps
- Potentially less efficient

Process switching

What does the kernel need to do when switching processes?

- Adjust Instruction, Base and Frame Pointer – okay

The process shouldn't notice it was interrupted! What else do you need to adjust?

- General purpose register
- Address space
- And a bit of housekeeping: Open files, scheduling priorities, ...

⇒ These things make up the PCB - The Process Control Block

Where are the PCBs stored? In user or kernel space?


- Kernel space! Users shouldn't be able to modify them

Are the PCBs always valid?

No! Some parts (registers, PC, etc) only when the process is not running. Why?
They are saved when it is switched out.

Process switching

In pictures



As text

1. Transition to the kernel (How? Interrupt, Exception, Syscall)
2. Save the context of the process. What was that again? Registers, Program counter, Stack pointer, etc.
3. Where do you save it to? A per-process kernel stack typically (as it is often moved there automatically) and then move it to the PCB.
4. Restore the context of the next process
5. Leave kernel mode and transfer control to the PC of the next process

Threads



What are processes, address spaces and threads? How do they relate to each other?


- A thread is an *entity of execution*, the personification of control flow
- A thread lives in an address space, i.e. all the addresses that it can access and the data that is stored there
- Thread + Address Space = Process

Spawn a few threads using pthreads!

Write a small program that creates five threads using the pthread library. Each thread should print its number (e.g., **Hello, I am 4**) and the main program should wait for each thread to exit.

Thread models — One To One

One To One




Problems?

Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff
- + Blocking does not affect other threads
- + Can piggy-back onto the OS scheduler
- **Must** piggy-back onto the OS scheduler
- Relatively high overhead due to context switches
- Relatively high overhead *when creating one!!*

Thread models

Many to One



Thread models - Many To One

Problems and benefits of *Many To One*?

Do they improve anything?

- Dummy
- + **Scales with core count?**
 - Can only use one core
- + **Conceptually easy — the OS does the hard stuff?**
 - Harder to implement — the OS doesn't help you much
- + **Blocking does not affect other threads?**
 - Blocking *does* affect other threads
- + **Can piggy-back onto the OS scheduler?**
 - Can *not* piggy-back onto the OS scheduler
 - **Must** piggy-back onto the OS scheduler?
- + Can implement its **own** scheduler
 - **Relatively high overhead due to context switches?**

Thread models - Many To One

Do you know a programming language / runtime using that?

E.g. `node.js` using its „event loop“

A small excursion - Structured Programming

Control flow should fall into one of four patterns:

- Sequence: One block is executed after another
- Selection: One or more are executed (i.e. an `if`-statement)
- Iteration: A block is executed more than once (i.e. a loop)
- Recursion: A block calls itself until an exit condition is met (i.e. recursion!)

Do you know any keyword in C which *doesn't* quite adhere to that but can instead totally spaghetti your control flow? `goto`

Famous paper by a proponent of Structured Programming:

„*Go To Statement Considered Harmful*“ by Edsger W. Dijkstra


And what about threads?

- Can outlive the methods they were spawned in
- Can use variables and fields after they went out of scope in a method
- Can split up or transfer their control flow arbitrarily

So that might sound familiar...

Thread models - Many To One

Structured Concurrency



Taken from vorpus.org

Nice, but what does this have to do with ULTs?

- Spawning lots of threads for small operations *is too slow otherwise*


Further reading:

[Notes on Structured Concurrency](#)

[ULTs and Structured concurrency in Java - Project Loom](#)

Thread models - Many To Many

Many To Many



Thread models - Many To Many

Problems and benefits of *Many To Many*?


Important: The kernel *knows about the user level scheduler*

- Dummy
- + **Scales with core count?**
- + Scales with core count
- + **Conceptually easy — the OS does the hard stuff?**
- Harder to implement — the OS doesn't help you much
- + **Blocking does not affect other threads?**
- + Blocking does not affect other threads as the kernel informs the user scheduler (**upcalls**) and does *not* pause the whole kernel level thread
- + **Can piggy-back onto the OS scheduler?**
- Can *not* piggy-back onto the OS scheduler
- **Must** piggy-back onto the OS scheduler?
- + Can implement its **own** scheduler

Thread models - Many To Many

Problems the second

You have all attended SWT 1! So let's have a look.



And preemption is now possible, which might complicate user code.

What events can trigger a context switch?

- Voluntary: yield, blocking system call
- Involuntary:
 - interrupts, exceptions, syscalls
 - end of time-slice
 - high priority thread becoming ready


What events can trigger a context switch?

- Most libraries only support *cooperative scheduling*
- Why is switching with preemption, interrupts, blocking system calls hard?
Kernel is not aware of the ULTs and will return where it came from — *but not call out to the scheduler and carry on*
- The benefit of platforms: How can **Java** (using Project Loom) or **Node.js** switch on most of the above?
You can not execute syscalls directly, but need to call library methods!
Suspension points can be inserted there.

„Jobs are either I/O-bound or compute-bound. In neither case would user-level threads be a win. Why would one go for pure user-level threads at all?“

- Program structure (e.g. Structured Concurrency, channels or just easier pipelines)
- The same or higher I/O throughput if on an abstracted platform

Should a fork do this?



Threads — Forking

Who is aware of the fork?

- The thread that executed the fork and the child
- Who else? Nobody!

Can you foresee any problems?

- If you copy the thread it might do weird things
- Locks are memory and will be copied as well! What if a thread holds a critical lock when it is forked?

Do you need thread duplication for our shell example?

No, we `exec` anyways.

Summary

`fork` is not as simple as it once was. Is it still a good abstraction?

What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

Why would you need kernel *mode* threads?

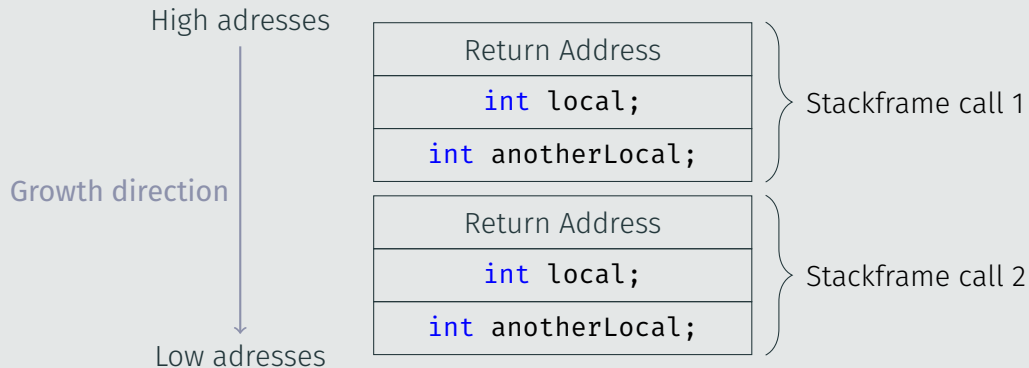
- Sometimes the OS needs to do some housekeeping that doesn't belong in a system call
- Free some pages, swap something in and out of memory
- Flush pages from the disk cache to the hard disk
- Perform VFS Checkpointing
- ...

Stackframes

Let's look at some code

```
1  int foo(int a, int b) {  
2    int local = a + b;  
3    int anotherLocal = 'a';  
4    return local + anotherLocal;  
5  }
```

What is this? What does it contain for `int foo(int,int)?`



Stackframe

How are arguments passed?

Registers! What happens for

```
1 void foo(int a, int b, int c, int d, int e, int f, int g, int h);
```

We do not have enough registers!

Fixing...

int h
int g
Return Address
int local;
int anotherLocal;

Popping frames

How could you release the space? I.e. figure out by how much to decrement the stackpointer?

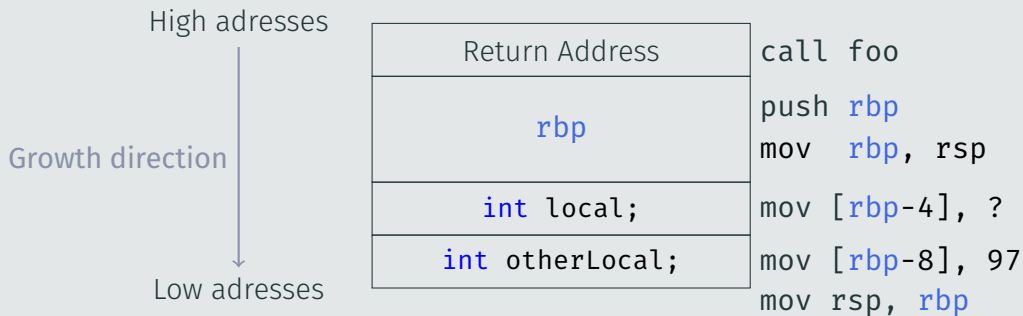
The compiler surely knows how much space the local variables take up, right?

No! C has e.g. *variable-length arrays* (VLA): `new int[someVariable]`.

⇒ Store the address at the start of the frame! Enter: `rbp`, the base pointer.

Base pointer and the function (Pro|Epi)log

Base pointer and the function (Pro|Epi)log




Base pointer and the function (Pro|Epi)log

Base pointer and the function (Pro|Epi)log

High addresses

Growth direction

Low addresses




```
mov rsp, rbp
pop rbp
ret
```

Memory Management Basics

And once again: What is the difference?


- We assume no 1:1 mapping (i.e. we have virtual memory)
- *All program addresses are virtual*
- Mapped to *physical* addresses as needed by the memory management unit

What is *internal* fragmentation?



Internal, i.e. *within* a block

What is *external* fragmentation?



External, i.e. due to *external factors* (different time-to-free)

Can you have *both* types at the same time?

Yes!

- Allocate in *chunks* by e.g. rounding up to 2^x
- Have different lifetimes

⇒ Wasteful allocations scattered throughout RAM

Fragmentation

What do we do now? This sounds bad!



This is called **Compaction**

Compaction - Is that even possible?

- C uses direct pointers
- ⇒ They are all garbage now!
- Works just fine in languages with indirections (e.g. garbage collection)
 - Also works for segments in physical memory! How? Update base addresses in MMU

Segmentation


Where have you seen that word before while sadly staring at your screen?

> Segmentation fault (core dumped)

Segmentation

A few example segments


0xFFFFFFFF



0x00000000

A Virtual Address

What does it look like?



$0xCAFE < 0xF00D$ $0xCAFE < 0xF00D$

$$0xB105 + 0xCAFE = 0x17C03$$

And let's try it

Segments

Segment Number	Base	Limit
0	0xdead	0x00ef
1	0xf154	0x013a
2	0x0000	0x0000
3	0x0000	0x3fff

Your task

Virtual Address	Segment Number	Offset	Valid?	Physical Address
0x2020	3	0x3999		
		0x0204	yes	
			yes	0xf15f

And let's try it

Solution

Virtual Address	Segment Number	Offset	Valid?	Physical Address
0xf999	3	0x3999	yes	0x3999
0x2020	0	0x2020	no	Offset outside limit
0xc204	3	0x0204	yes	0x0204
0x400b	1	0x000b	yes	0xf15f

Memory allocation policies

Some common strategies

Which strategies for finding free blocks do you know?

First Fit, Best Fit, Worst Fit

First Fit

Pick the first block that is large enough

Best Fit

Pick the *smallest* block that fits



Worst Fit

Pick the *largest* block that fits


Let's try them

Best fit

Allocate 500, 1200, and 200, fail if not possible.




And compact it to fit the next one!



What allocator would you make up for this?

You are a poor kernel and you need lots of inodes

Every inode has the same size, 64 Byte. Can you think of any fast allocation strategy that does not waste a single bit?



This is called a *Slab allocator*


So far we've seen

- Consistent large blocks \Rightarrow Low external, high internal fragmentation
- Fitted blocks \Rightarrow High external, low internal fragmentation

Can we do better for some applications? Any ideas?

Buddy Allocator


Allocator



How do you find a fitting Element? *Freelist!*

And if there is no such block? *Recursively split a higher-up block*


Merging



Buddy Allocator

How small/large can the free list be?

Allocate 2^m chunk of memory in a managed Block of 2^k (here: $k = 18$, as $256 \text{ KiB} = 2^{18}$)



⇒ Max size $\frac{1}{2} \cdot 2^{k-m}$

⇒ Min size 0

What kind of fragmentation can occur?

Internal fragmentation


- Power of two blocks

⇒ Request memory of size $2^k + 1, k \in \mathbb{N}_0$

External fragmentation

- Free every other block in a level

External fragmentation



But this works alright for larger sizes. So combine it with...

...the Slab allocator! Allocate large chunks with the buddy allocator and small chunks within them using the slab allocators

Paging



What is that? What is the difference to Segmentation?

- Virtual memory is broken into fixed-size chunks (**pages**)
- Physical memory is broken into fixed-size chunks of the same size (**frames**)
- Virtual pages are mapped to page frames Always?
No! (Paged out, zero pages, ...)

Benefits over Segmentation?

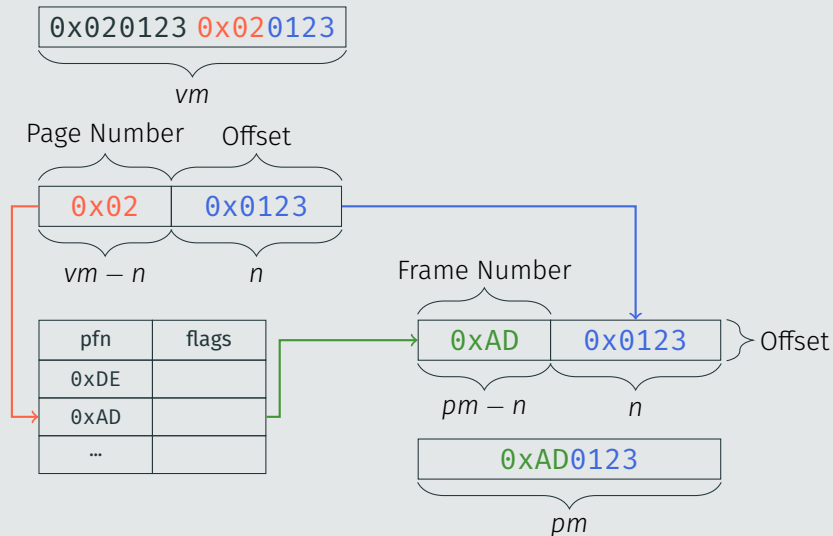
- Virtual memory does not need to map to *continuous* physical memory
- Swapping in/out is easier
- No external fragmentation, little internal

Segment and Page tables

Segment Number	Base	Limit
0	0xdead	0x00ef
1	0xf154	0x013a
2	0x0000	0x0000
3	0x0000	0x3fff

Paging - Single Level Page Table

Segment and Page tables



Single Level Page Table - Disadvantages

Why could that be a bit problematic?

- Increases with the size of the size of the *virtual* address space
- 64 Bit AS, 4KiB (2^{12}) pages $\Rightarrow n = 12 \Rightarrow 2^{vm-n} = 2^{64-12} = 2^{52}$

\Rightarrow If every entry was 1 Bit we'd need (asking **units...**)

You have: 2^{52} bit

You want: tebibyte

* 512

/ 0.001953125

- You might *not* have that much memory to spare :)

Single Level Page Table - Disadvantages

Math is fun, let's do some math

Calculate the space requirements for a single level page table with

- 32-bit virtual addresses, 4KiB pages, 4 bytes per page table entry
- 48-bit virtual addresses, 4KiB pages, 4 bytes per page table entry

32-bit


- $vm = 32, 4Kib = 2^{12} \Rightarrow n = 12$
- $2^{32-12} = 2^{20}$ entries $\Rightarrow 2^{20} \cdot 2^2 = 2^{22}$ Byte (4 MiB)

48-bit

- $vm = 48, 4Kib = 2^{12} \Rightarrow n = 12$
- $2^{48-12} = 2^{36}$ entries $\Rightarrow 2^{36} \cdot 2^2 = 2^{38}$ Byte (256 GiB)

Alternatives to Single Level Page Tables

Mutli-Level page tables



Benefits and Drawbacks?

- Pointer chasing down each level \Rightarrow More memory accesses
- + Address spaces are *sparse* \Rightarrow Only instantiate page tables you need

What do those do?

- Map *physical* addresses to virtual ones
 - Why? Physical address space is much smaller and you don't need a table per address space
 - If you don't have a table per address space, how can you still have multiple?
- ⇒ e.g. Linked List of entries per physical frame

Any drawbacks?


- We mostly care about *the other direction*, i.e. virtual ⇒ physical
- That requires iteration :(

Inverted Page Tables

How can we speed them up?

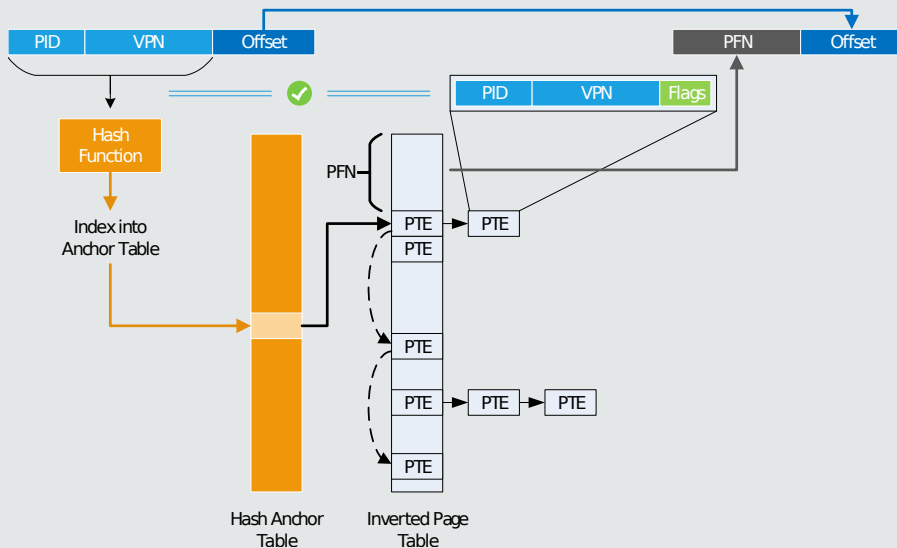
After having attended *Algorithmen I* we all know: Hacktables are $O(1)$

Hashed inverted page tables



Inverted Page Tables

Hashed inverted page tables



Why do pages have a valid/present bit?

What is it for?

- The page might not be present. Why? Swapped out, not zeroed yet, ...

What happens on access if it is not set?

- Page fault!
- Handle it and do sth. sensible (or crash the process...)

TLB



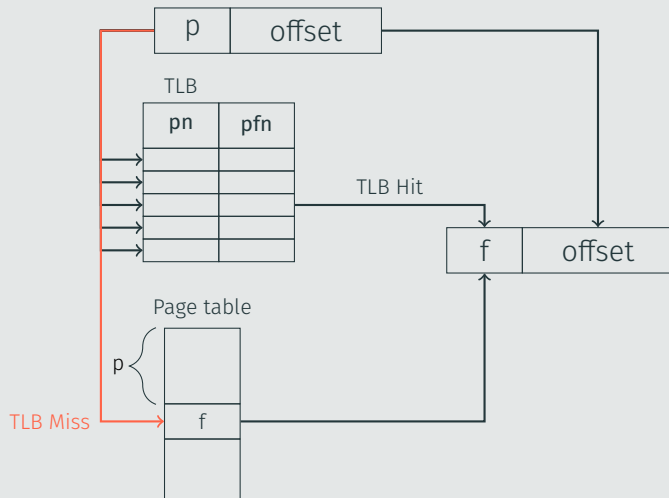
What does that refer to?

Translation Lookaside Buffer. What's that for?

- Lookup from virtual to physical address can be *slow*
- You need to translate addresses *all the time*

⇒ There is no problem you can't solve with another caching layer
(except having too many caching layers) [Nearly the Fundamental theorem of software engineering](#)

TLB layout



What is the difference between software and hardware walked Page Tables?

The TLB is *the same*, what differs is the behaviour on a *cache miss*.

Hardware walked

- Hardware looks at your page table
 - Hardware *resolves* the physical address using it
 - On failure a page fault is raised
- ⇒ What does that imply? Page table layout is **fixed**

Software walked

- Transfers control to TLB miss handler
- Kernel routine walks the page table and tries to find a mapping
- Loads that mapping into the TLB and can choose which entry to evict!
- If there is none ⇒ Jump to page fault handler

What are benefits / drawbacks of software walked TLBs?

- + Free to choose page table layout (fit your algorithm)
- + Free to choose which TLB entries to evict \Rightarrow Use optimized strategy
- Greater overhead

What information does the TLB store?

- Physical Frame Number, Virtual Page Number
- Valid / Present Bit
- Modified bit, permissions, ...

When is a TLB miss or page fault raised?

Due to the TLB:

- Software walked: No entry found in TLB \Rightarrow TLB miss exception
- Hardware walked: No mapping in page table \Rightarrow Page fault

And other problems: *Invalid access to a mapped page*

- Software walked: Raise some kind of TLB fault if the page is e.g. read-only
- Hardware walked: Page fault raised, page fault handler has to find out what happend

Page Fault Handling

What is Demand-Paging and Pre-Paging?

Demand-Paging:

- Load pages on-demand, right when they are needed

Pre-Paging:

- Loaded Pages speculatively in batches, even *before* you need them

Why would you (not?) use Demand-Paging?


- + Only loads needed data \Rightarrow Less memory wasted
- Generates lots of page faults before working set is in memory

Why would you (not?) use Pre-Paging?

- + Might reduce number of page faults
- Loads more than needed \Rightarrow Wasteful
- More I/O \Rightarrow Slower?
- + HDDs a lot faster when reading chunks

Different kind of page faults

Not all pages are created equal. Do you have any idea what types of page faults typically exist?



Different kind of page faults

Why are pages to generic memory zero filled? It could leak other processes' memory otherwise (called a [Covert Channel](#)).

And there is one other kind of page fault...

...The page was stolen by the OS and swapped out!

Also supported on some systems: *Purgable memory*. Stolen from [Apple](#) and also implemented [in SerenityOS in this video](#).

What kind of information does the page fault handler need?

- Access flags: Can the user perform the operation on this page?
- Where to find the most recent version (different for zero filled, file backed, etc.)

How could you implement Copy-on-Write memory?

- Mark memory as read-only on fork
- Add an additional **CoW** flag: When a page fault is raised check it, copy the page and clear the **CoW** and **ro** flag

Page replacement

The pager in some systems tries to keep „spare pages“. Why?

- OS needs free (pre-zeroed) frames to assign to processes
- What happens when there are none and a page fault occurs?

⇒ Needs to write dirty pages back to disk

⇒ Sloow

If you need to swap out a page, what pages do you search for a victim?

Local page replacement algorithms:

- Only look through the frames *of that process*

Global page replacement algorithms:

- Look through *all* frames, even that of other processes

What are the pro and cons of local algorithms?

- + Guaranteed number of pages per application
- + Potentially faster
- Guaranteed number of pages per application \Rightarrow Can not adapt as easily to changing demands
- Difficult selection of optimal number of frames per application (see point above)

Every process should get an equal number of pages. Do you use a Global or Local Algorithm?

Local

What is the working set?

- The set of pages that a process accessed in the last Δ page references

What is Thrashing?



- Not enough frames to fit the working set
- ⇒ Pages will be stored to disk and reloaded very often


What are those?

We need to find a victim page to replace with a new one.

Common strategies

- **FIFO:** First page to be loaded is the first to be unloaded
- **LRU:** Least recently used page is evicted
- **Optimal:** Evict the page that is used the *furthest into the future*. Feasible?
Used as a *benchmark*
- **Clock:** Uff, let's talk about that on its own page

Clock



Where did that R come from?

- Referenced (and modified) bits are set by the CPU when accessing or writing to the page
- Referenced bits are periodically cleared by the kernel using timer interrupts

Do it

- Clock: Ordered by Load time ASC, Head is on Frame 3
- Reference order: 4, 0, 0, 0, 2, 4, 2, 1, 0, 3, 2

Frame	Virtual page	Load time	Access time	Referenced	Modified
0	2	60	161	0	1
1	1	130	160	0	0
2	0	26	162	1	0
3	3	20	163	1	1

How well does LRU work for the Stack, the Heap and the Code segment?

- Stack: Beautifully. The older it is the longer it will take to be reached again
- Code: Mostly, loops are mostly linear and it follows certain patterns
- Heap: Well, the heap has more random access patterns. So not that well, probably

Pointer Arithmetic

Are numbers!

```
1  int  hello      = 20;      // no worries
2  int* hPointer   = &hello;  // no worries
3  int* pointer    = 0x200;   // no worries
```

And you can convert between them

```
1      int  hello          = 20;
2      int* aPointer       = &hello;
3      intptr_t asInt      = (intptr_t) aPointer;
4      int* backToPointer  = (int*) asInt;
```

And what do operations do?

```
1  int array[5] = {0, 1, 2, 3, 4};  
2  printf("Current value: %d", *array); // 0  
3  array++;  
4  printf("Current value: %d", *array); // 1
```

Wait, weren't integers more than one byte?

- + and - on a pointer decrease it by *the size of its type*

```
1  int* pointer;  
2  pointer++; // is the same as  
3  intptr_t asInt = (intptr_t) pointer;  
4  asInt += sizeof(int);  
5  pointer = (int*) asInt;
```


Multiplication? Division?

- Probably a compiler error
- Sounds a bit useless

We can also cast them!

```
1  int hello[5] = {0};
2  char* start = (char*) hello;
3  *start = 1;
4  printf("First value: %d\n", hello[0]); // 1!
5  // Why is it 1 and not 2147483648?
```

Some interesting things are writable

DoubleForker

DoubleForker!

Modifying Stackframes in GDB

Capture the flag!

```
1  #include <stdio.h>
2  #include <string.h>
3
4  int readIt() {
5      char buffer[10];
6      int result = 1;
7
8      printf("Enter the password: ");
9      gets(buffer);
10
11     if(strcmp(buffer, "Secret") == 0) {
12         result = 0;
13     }
14     return result;
15 }
16
17 void flag() {
18     printf("FLAG!\n");
19 }
20
21 int main() {
22     if (readIt() == 0) {
23         flag();
24     } else {
25         printf("Nope!\n");
26     }
27     return 0;
28 }
```

Modifying Stackframes in GDB

Useful Resources

- [Stackframe layout \(for overflowing\)](#)
- [Input non printable characters in GDB](#)
- Run it in GDB, Compile with **-fno-stack-protector**
- `lscpu` to find Endianness of your computer (probably Little Endian)

Spoiler Example Solution

Created on my computer while running in GDB. Addresses may vary.

```
r <<< $(python2 -c 'print
    "A"*14 # Padding to overflow to base pointer
    + "\x00\x00\x7f\xff\xff\xff\xdd\x02"[:-1] # overwrite BP
    + "\x00\x00\x55\x55\x55\x55\x51\xbd"[:-1] # overwrite ret addr
    ')
```

Modifying Stackframes in Real Life

Why that won't work 1: `gets`

- We needed to inject some null bytes to pad our addresses to 64 bit.
- `gets` just reads until EOF/newline and doesn't care about size or zeroes

⇒ *Really dangerous*

⇒ Removed from the C standard (but it still probably compiles)

Why that won't work 2: **Stack-Canary**

- The compiler inserts a random number between the local variables and the return address
- If it is changed the program will die with a nice message:
***** stack smashing detected ***: terminated**
- Disabled with `-fno-stack-protector`

Why that won't work 3: ASLR

- We injected *absolute addresses* to jump to (or set the base pointer to)
 - Modern Operating Systems employ **A**ddress **S**pace **L**ayout **R**andomization
- ⇒ Your code is loaded at a random address and absolute addresses don't work
- GDB disables ASLR for the loaded program so we ran it in there

Functions are...

Code! And where is code stored?

In memory. So let's point to it.

```
1 #include <stdio.h>
2
3 void sayHi(char* name) {
4     printf("Hello, %s!\n", name);
5 }
6
7 int main() {
8     (&sayHi)("John");
9     return 0;
10 }
```


How can we declare a variable?

```
1   ??? myFunction = &sayHi;
2   // We basically copy the declaration
3   void sayHi(char* name) = &sayHi; // Nearly!
4   void (sayHi)(char* name) = &sayHi; // Closer!
5   void (* myName)(char* name) = &sayHi;
6   //^^^  ^  ^^^^^^  ^^^^^^^^^^^
7   // | Pointer    | Parameters delimited with ,
8   // Return Type |
9   //              Var name
```

How can we declare a variable?

```
1 #include <stdio.h>
2
3 void greetMultiple(char* name, char* other) {
4     printf("Hello, %s and %s!\n", name, other);
5 }
6
7 int main() {
8     void (*myFunction)(char* name, char*) = &greetMultiple;
9     myFunction("Peter", "Jane");
10    return 0;
11 }
```

Aufgabe

1. Schreibe eine Insertion-Sort die ein int array, einen compare-Funktionspointer und was man sonst so braucht entgegennimmt
2. Schreibe eine compare-Funktion, die Integer absteigend sortiert und nutze sie
3. Schreibe eine compare-Funktion, die Integer basierend auf der Anzahl an 1-Bits sortiert

```
1  int main() {  
2      int array[] = {5, 3, 2, 4, 1};  
3      int length = sizeof(array) / sizeof(int);  
4  
5      insertionSort(array, length, compareDesc);  
6  
7      for(int i = 0; i < length; i++) {  
8          printf("%d - %d\n", i, array[i]);  
9      }
```

Caches



How many do you typically have in your computer?

Quite a few! Typically L1, L2, L3 <maybe more> and your RAM

And why do you have multiple?

- Fast caches are expensive (and size limited) \Rightarrow Small
- Multiple levels of slower but larger caches arranged in a *cache hierarchy*

Why do caches even work?

Two main principles:

- Temporal locality: Memory that was accessed will be accessed again soon
- Spatial locality: If address X was accessed, $X \pm 1$ will likely be accessed soon
 \Rightarrow Don't just cache single words but entire *lines*. How large? Probably ≥ 64 Bytes for *reasons*

Cache lines

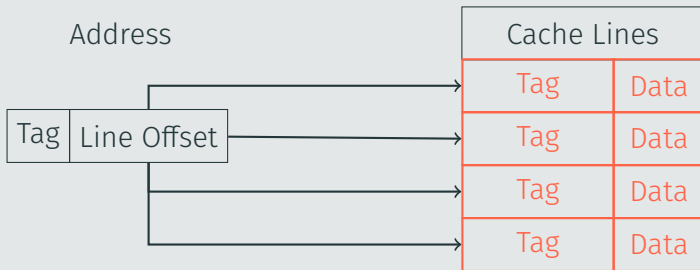
Why is a cache line useful? Can't we just make N memory accesses? Reading a chunk from DRAM is *much* more efficient than reading them one after another

What types do you know?

- Fully Associative
- Set-Associative
- Direct Mapped

Cache Types - Fully Associative

What is that?




When can cache misses happen here?

- Never accessed that address before. Called a **Compulsory Miss**
- Cache was full and it was evicted. Called a **Capacity Miss**

Cache Types - Direct Mapped


What is that?



When can cache misses happen here?

- Never accessed that address before. Called a **Compulsory Miss**
- Cache was full and it was evicted. Called a **Capacity Miss**

What is that?




Conflict misses

There was space, but it was mapped to the same slot!

Cache Types - Set associative

What is that?



Insertion

Find the correct *set* using the index, then treat each set as a *Fully Associative* cache.

Conflict misses?

There was space, but it was mapped to the same *set*!

You have a Cache and real memory behind

What do you do when you write to ...


... an address in the cache?

- Write Through: Update the cache *and the main memory*
- Write Back: Update only the cache. Update main memory *on eviction*

... an address *not* in the cache?


- Write-allocate: First load it into the cache. Then see above
- Write-to-memory: Don't load into cache, modify in memory

At which position do you put the cache?



Placing The Cache

At which position do you put the cache?



What kind of addresses does this use for Index/Tag?

Virtually Indexed, Virtually Tagged (VIVT)

What kind of addresses does this use for Index/Tag?


Virtually Indexed, Virtually Tagged (VIVT)

Benefits? Drawbacks?

- + Fast, no address translation
- Ambiguity Problem (The same virtual address might point to different physical addresses over time)
- Alias Problem (Multiple virtual addresses might point to the same physical address)

Placing The Cache

At which position do you put the cache?



What kind of addresses does this use for Index/Tag?

Virtually Indexed, Physically Tagged (VIPT)

What kind of addresses does this use for Index/Tag?


Virtually Indexed, Physically Tagged (VIPT)

Benefits? Drawbacks?

- + Still relatively fast: Can lookup the index while the MMU works
- + Ambiguity solved, as it is tagged by physical address (iff the entire PFN is used as a tag!)
- Alias Problem (sometimes!)

Placing The Cache

At which position do you put the cache?



What kind of addresses does this use for Index/Tag?

Physically Indexed, Physically Tagged (PIPT)

What kind of addresses does this use for Index/Tag?

Physically Indexed, Physically Tagged (PIPT)

Benefits? Drawbacks?

- + Transparent to CPU: no alias, no ambiguity
- + Coherency can be implemented in hardware
- Allocation conflicts

Allocation conflicts

- How well do you think *contiguous* virtual pages fit into the cache?
 - Maybe not so well. The cache operates on *physical* addresses
- ⇒ Sequential virtual pages might be mapped to *conflicting* physical ones!

Inter-Process-Communication

How can different (concurrent) activities interact?

- Shared memory (implicitly as they share an address space and explicitly via a shared memory area)
- OS facilities (e.g. messages, pipes, Signals, sockets)
- High level abstractions (files, database entries)

What mechanism do you need for IPC in an imperfect world?

Timeouts! You don't want to wait for buggy programs or poor dead ones :(

Asynchronous send Operations require a buffer. Where do you put that?

- Actually, *why do they need a buffer?* Sent but not yet received messages
- They can be in the *receiver's* address space, the *sender's* AS or in the kernel

Buffering in the Receiver's Address Space

- Is that a good idea? What happens when you have many clients?
- ⇒ Does not scale well. You either allow every sender to allocate memory for you (*ouch*) or you might run out with many clients

Buffering in the Kernel

- Is that a good idea? What happens when you have many clients?
- ⇒ Does not scale well. You either allow every sender to allocate kernel memory (*ouch*) or you might run out with many clients

Buffering in the Sender's Address Space

- Sounds good? Yea, we are running out of options. But it's mostly alright!
- + We only need to store it once: The sender has it in a buffer somewhere anyways
- + Scales better, as each sender keeps their messages
- ± We need to tell the client when it can reclaim the buffer

You are a very popular process that receives and handles many messages.

- For simplicities sake you chose *Synchronous IPC*. What problem might occur with many clients?
- You spent your whole life *waiting for timeouts to expire*
- How could you solve that with a new syscall? How does **send-and-receive**, which sends and instantly receives help?
- The server can assume you are using it and set a **zero** timeout. After all, if you are using that syscall you *will* be waiting.

How can you emulate asynchronous IPC using synchronous IPC?

- Use a Proxy thread
 1. Copy message to proxy thread
 2. Proxy threads sends synchronously and might block until recipient calls receive
- + Allows async I/O
- How many messages can you send? Yea, one per thread...

How can you emulate Synchronous IPC using asynchronous IPC?

What about

```
1  do {  
2    res = async_send(message);  
3  } while(res != MESSAGE_SENT);
```

Will wait until it was *sent*, not until it was *received*.

⇒ Implement a very simple protocol involving multiple messages

⇒ Send **ACK** message on receiver side, wait for **ACK** to be received.

And receive?

Just loop until `async_receive` receives a message (that is not an **ACK**)

Race Conditions

The Parallel Wizard Strikes Again

```
1 current = get_balance();  
2 current += delta; // delta ∈ {−50, 100}  
3 set_balance(current);
```


What happens if this code is executed in *parallel* for the two values of `delta`?

Unrolled


```
1 current = get_balance();  
2 int tmp = current;  
3 current = tmp + delta; // delta ∈ {−50, 100}  
4 set_balance(current);
```

Race Conditions

Thread execution



And if you run it concurrently on a single-core system?




What is the value of `tally` at the end?

```
1 #include <stdio.h>
2
3 int tally;
4 void total(int N) {
5     for(int i = 0; i < N; i++)
6         tally += 1;
7 }
8
9 int main() {
10     tally = 0;
11
12     #pragma omp parallel for
13     for(int i = 0; i < 2; i++)
14         total(50);
15
16     printf("%d\n" , tally);
17     return 0;
18 }
```


Having Fun With It...

This!



And what happens when we use $N, N > 0$ threads instead of 2?

The range is now $[2, 50 \cdot N]$. We can still have the same problem.

...And Adding ULTs!

What would happen if we used ULTs?

Can ULTs be pre-empted? With cooperatively-scheduled (the common case) no!
⇒ Everything will be fine, *unless* the threads **yield** *during the increment*

In that spirit: Does this work (with 1:1 threads)?

```
1 void total(int N) {  
2     for(int i = 0; i < N; ++ i) {  
3         tally += 1;  
4         sched_yield();  
5     }  
6 }
```

Nope. It might be less likely that it is interrupted during the increment (Why? it never uses its timeslice), but it is still possible (e.g. due to a hardware interrupt)

Critical Sections

Does this work?

```
1 void foo() { // called in parallel
2   if(random() < 0.5) {
3     sleep(10);
4   }
5   a += 10;
6 }
```

Nope. Multiple threads can call `a += 10` at the same time.

Property 1: **Mutual Exclusion**

Synchronizing Properly - Notes on Terminology

Some boring definitions

```
1 void foo() { // called in parallel
2
3     // This is the code before the critical section
4     // ==> Entry section
5
6     // Here common data is accessed (e.g. shared variable).
7     // Only one thread might be in here at a time.
8     // ==> Critical Section
9     a += 10;
10
11    // This is the code after the critical section
12    // ==> Exit section
13 }
```

Everything else is the Remainder section

Does this work?

```
1 void foo() { // called in parallel
2   // The late bird catches the worm
3   while(me != selectLastWaiting()) {
4     sleep(10);
5   }
6   a += 10;
7 }
```

How long could a thread wait?

Forever :(

Property 2: **Bounded Waiting**

Does this work?

```
1 void foo() { // called in parallel
2   while(waitingThreadCount > 1) {
3     sleep(10);
4   }
5   a += 10;
6 }
```

Threads *outside* the critical section prevent threads from *entering* it

⇒ There's no progress!

Property 3: **Progress**

And the last one isn't really a property of a correct solution

Property 4: **Performance**

Synchronizing Properly - Core Properties

Mutual Exclusion

Only *one thread* can enter the critical section at once.

Progress

Threads in the *remainder* section do not prevent threads from *entering* the critical section.

Bounded Waiting

There is an upper bound on *how many* different threads can enter the CS while a thread is waiting.

Important: This is not a *time* bound!

Performance

The time overhead of the synchronization primitive is low (for low / medium / high contention).

Synchronizing Properly - Core Properties

Does this code fulfill all properties?

```
1  /* aligned to cache lines */
2  volatile int next = 0;
3  volatile int executing = 0;
4
5  lock_acquire() {
6      /* atomic version of "ticket = next; next += 1;" */
7      int ticket = fetch_and_add(next, 1);
8      /* busy wait until the counter matches */
9      while (ticket != executing ) {}
10 }
11
12 lock_release() {
13     executing++;
14 }
```

Does this code fulfill all properties?

- Mutual Exclusion: Yes (though if you have 2^{32} threads waiting it doesn't)
- Bounded Waiting: Yes. Eventually my number is the next one!
- Progress: Yes, threads in the remainder section do not hinder any thread from entering the critical section. Only threads in the *entry section* can do that temporarily.

Fixing The Bank

Remember the bank example?

```
1  current = get_balance();  
2  current += delta; // delta ∈ {−50, 100}  
3  set_balance(current);
```

How could you fix that?

Synchronize it!

```
1  lock(L);  
2  current = get_balance();  
3  current += delta; // delta ∈ {−50, 100}  
4  set_balance(current);  
5  unlock(L);
```

Synchronization Primitives

There are different kinds of synchronization primitives

Which ones do you know?

Spinlock

- `lock` / `unlock`
- Busy-waiting and atomic instructions (e.g. compare-and-set)
- Recommended for *short* critical sections as it wastes CPU time
- Preemption wastes more resources (threads can't make progress)
- \Rightarrow Mostly used in the kernel without interrupts

Semaphore

- `wait(sem)` (also called *acquire*) / `signal(sem)` (Also called *release/post*)
- Has an internal counter that is decremented (wait) or incremented (signal)
- Blocks if you try to decrement it below 0
- Can be used to implement bounded buffers

Mutex (Binary Semaphore)

- `lock(m)`, `unlock(m)`
- Or a Semaphore with values 0 and 1

Synchronization Primitives

Condition Variables

```
1 void consume() {
2     lock(l);
3     while(queue.size == 0) {
4         unlock(l);
5         sleep(); lock(l);
6     }
7     queue.poll(); unlock(l); signal();
8 }
9 void produce() {
10    lock(l);
11    while(queue.size == MAX_SIZE) {
12        unlock(l);
13        sleep(); lock(l);
14    }
15    queue.add(X); unlock(l); signal();
16 }
```

This code can incorrectly sleep a consumer/producer. How? *Lost wakeup problem*

Synchronization Primitives

Condition Variables

```
1 void consume() {
2     lock(l);
3     while(queue.size == 0) {
4         // unlocks and sleeps atomically. Relocks when waking up
5         wait(cond_filled, l);
6     }
7     queue.poll(); signal(cond_empty); unlock(l);
8 }
9 void produce() {
10    lock(l);
11    while(queue.size == MAX_SIZE) {
12        // unlocks and sleeps atomically. Relocks when waking up
13        wait(cond_empty, l);
14    }
15    queue.add(X); signal(cond_filled); unlock(l);
16 }
```

Now no wakeup is lost :)

Kernel Synchronization

How could you achieve mutual exclusion on Single-Core systems?

Masking interrupts! Why? Only one thread can run at a time, disabling interrupts prevents preemption (and other interrupt handlers)

Does this work on Multi-Core systems?

Nope! Masking only affects the current CPU.

Additionally, another core could be in the same routine and access the same data

How would you solve that problem in the kernel?

Real locks

The Big Kernel Lock™

Have a big lock *for the whole kernel*. Implications? The kernel effectively serialises access

⇒ Can't make use of your processors if you have many syscalls

This removes the implementation of the big kernel lock, at last. A lot of people have worked on this in the past, so the credit for this patch should be with everyone who participated in the hunt. ([Commit message](#))

Mutual Exclusion In The Kernel

Fine Grained Locking

- + Only lock areas that are *relevant* for the *current* task
- ⇒ Have many small locks
- Complex and error prone

Remember Spinlocks and Interrupt handlers?

Without disabling interrupts there is a problem: *Lockholder Preemption*

1. Thread enters spinlock
2. Thread gets pre-empted by an interrupt handler
3. Interrupt handler needs the same lock ⇒ Can never acquire it!

⇒ You might still need to disable interrupts for those


Deadlocks



What is that? Do you know an example?

- Several processes or activities can not make progress, as they are waiting for resources held by each other
- Examples: 4-way intersection, *Dining Philosophers*

The Problem



Philosophers (P) want to eat, but to do that they need two Chopsticks (Chop)!

How can this deadlock?

Why did that happen? What fateful circumstances lead to this starvation?

The Four Horsemen of the Apocalypse *Coffman Conditions*

Mutual Exclusion

Resources can not be shared between processes

Hold and wait

A process already holding resources can acquire more

No Preemption

Resources can not be forcibly taken away from processes

Circular Wait

There exists a set of Processes $P_0, P_1, \dots P_n$ where P_0 is waiting for a resource held by P_1 . P_1 is waiting for a resource held by P_2 , ...and P_n is waiting for a resource held by P_1

Note

These conditions *are not independent!* (e.g. Circular Wait \Rightarrow Hold And Wait)

Finding a deadlock

Code

```
1  Spinlock s1,s2, s3 = FREE;
2  int counter = 0;
3  void Thread1() {
4      if(counter == 0) {
5          lock(s1);
6          counter++;
7          unlock(s1);
8      }
9      lock(s2);
10     lock(s3);
11     // update some more data
12     unlock(s3);
13     unlock(s2);
14 }

15 void Thread2() {
16     lock(s3);
17     counter++;
18     // update some data
19     if(counter == 2) {
20         lock(s2);
21         // update some more data
22         unlock(s2);
23     }
24     lock(s1);
25     // update even more data
26     unlock(s3);
27     unlock(s1);
28 }
```

Deadlock Prevention

Make a deadlock *impossible*! How? Break ≥ 1 of the four necessary conditions

Deadlock Avoidance

- Deadlocks are still possible
- The resource allocator knows what resources are used by the processes
- The resource allocator denies requests that *might* lead to a deadlock

How can you negate *Mutual Exclusion*?

Sometimes: Spooling

Like a Printer

- You send a job
- It is executed

⇒ Only the executor has access to the resource

Negate *Hold And Wait*

Allocate resources atomically: All you will need or nothing

⇒ Once you have resources, you can no longer request new ones

How can you negate *No Preemption*?

Allow Preemption! Normally done by *multiplexing* resources (how RAM or CPU time is handled).


Not always possible

Negate *Circular Wait*

Order resources and only allocate in the *same* order, everywhere.

Commonly used (and also in the current exercise (**not anymore** :() :))

The Problem





What kind of vertices and edges are in this graph?

How can you detect a deadlock in there?

Resource Allocation Graphs

Some examples



Is there a deadlock in one of the graphs?

Yes, in the left. Right has a cycle *but no deadlock*.

Cycle \equiv Deadlock only holds if you have *one* instance of each resource

Deadlock Empire

<https://deadlockempire.github.io/>

Solid-State Drives

NAND based flash memory

Rejoice, TI might be useful once

How long do writes/reads take normally?

- Reading a page: 25 μ s
- Writing a page: 250 μ s
- Erasing a block: 2ms

What happens when you just write to a random page?

Speeding things up

What could you change so writing pages is faster?

- Keep around spare *erased* pages

⇒ You do not pay the erase penalty!

- When do you create / reserve / erase those spare pages? Probably in the background. Any problems? Might get exhausted if you write too much data in a short timeframe or the disk is full!

Deleting files

What happens when you delete a file? What effect does that have on the SSD performance? The block is not freed \Rightarrow Can't be used as an erased empty page

What can the OS do to combat that?

The **trim** command

Can be issued by the OS to tell the SSD firmware what pages can be safely erased.

RAID



What is that?


A Redundant Array of Independent/Inexpensive Disks

Why would you use that?

- Probably cheaper than a SLED (Single Large Expensive Disk)
- Might be more resilient
- Might be faster


Great, you now have multiple disks. How do you store your files on them?

- „I like to live dangerously“ - RAID Level 0
- Mirroring: RAID Level 1
- Historic variants: RAID Level 2 and 3
- Block striping and parity: RAID Level 4
- Block striping and *distributed* parity: RAID Level 5




Benefits / Drawbacks?

- + Extremely fast (parallel reads and writes)
- + Can use full capacity
- If a single disk fails your files are toast




Benefits / Drawbacks?

- + You can lose all but one disk without losing data
- + Parallel reads possible
- Writes slower as they need to write to all disks
- Size equals the size of a single disk



What is that?

- Have $\log_2(N)$ parity disk
- Stripe data at the *bit* level
- Use a hamming code of proper size
- Spin the disks in lockstep (so you read all bits of your word at once)




What is that?

- Have a dedicated parity disk
- Stripe data at the *byte* level
- Spin the disks in lockstep (so you read all bytes of your word at once)

Benefits / Drawbacks

- You can lose a disk and restore it using the parity




What is that?

- Have a dedicated parity disk
- Stripe data at the *block* level

Benefits / Drawbacks

- + You can lose a disk and restore it using the parity
- + Good read performance



What is that?

- Stripe data at the *block* level
- Distribute parity across your disks

Benefits / Drawbacks

- + You can lose a disk and restore it using the parity
- + Good read performance

Compare SLED and RAID (Level 0, 1, 4, 5)

Each RAID uses 4 disks for actual data storage.

How many disks do you need?

- SLED: 1
- RAID 0: 4
- RAID 1: 8
- RAID 4: 5
- RAID 5: 5

You want to modify one byte of data. How many blocks do you need to read/write?

- SLED: 1 read + 1 write
- RAID 0: 1 read + 1 write
- RAID 1: 1 read + 2 write (1 data + 1 mirror)
- RAID 4: 2 read (data + old parity) + 2 write (data + new parity)
- RAID 5: 2 read (data + old parity) + 2 write (data + new parity)

You are using RAID


- You accidentally delete a file. **GONE**
- You accidentally overwrite a file. **GONE**
- Some data gets corrupted on one disk. **GONE** (*probably*)
- [The poor intern connects to the production database.](#) (Or [here](#)) **GONE**
- A crypto-locker takes out your computer. Believe it or not, ~~JAIL~~ **GONE**

So what do we learn?

RAID IS NO SUBSTITUTE FOR A BACKUP

Files And Directories

Paths (Linux)



What are the two basic access methods (patterns) for reading a file?

- **Sequential Access**

Accessed in order, reading sequential bytes. Writes append at the end.

- **Random Access**

Reading and writing at arbitrary positions, programmer needs to specify where to write/read

Which Library Functions and System Calls do we have to access files?

Opening the file

`fopen` / `open`. Provide mode as argument (read/write, where, create, etc.)

Closing a file

`fclose` / `close`

Reading from a file

`fread` / `read`

Writing to a file

`fwrite` / `write`

Flushing dirty buffers

The library functions sometimes buffer to reduce the amount of syscalls.

`fflush` flushes those buffers.

Which system calls / library functions move the „current position pointer“?

- Linux: `fseek` / `lseek`
- Windows: `SetFilePointerEx` / `SetFilePointer`

What happens when you move the cursor behind the end of the file?

It creates a hole filled with zeros! Such a file is called a *sparse* file and some file systems might not store empty regions.

How could you implement random access without a dedicated file pointer?

Add an offset to the `read` and `write` system call.

- + Save a system call
- Caller need to keep track of the current offset for sequential reads

`mmap` the file (works quite well, might cause page faults. Probably faster for large files than read/write calls)

What system calls do you need to list files in a Linux directory?

1. `opendir`
2. `readdir`: Returns a `dirent` with a *relative path*
3. `closedir`

Open Files


Kernel Data Structures For Open Files

What structures does the kernel use for open files?

PCBs with local open file tables

Global open file table


vnode table



Disk Space Allocation

Contiguous Allocation

What is that?




Challenges

- Sizing your block (internal fragmentation, growth)
- External fragmentation

Chained Allocation

What's that?




Benefits? Drawbacks?

- + No longer need a contiguous chunk
- Only sequential access
- A single corrupted pointer is very bad news

Linked List Allocation With FAT

What is that?





0		
1	7	⇐ File A
2	-1	
3	5	
4	-1	
5	2	
6	8	⇐ File B
7	3	
8	4	

Benefits? Drawbacks?

- + Hopefully fits in RAM \Rightarrow Iteration fast
- + Even if it doesn't fit: Less disk seeks
- Size depends on size of hard disk (one entry per Block)
- Might be too large to be cached in RAM (on large disk)

What is that?



Benefits? Drawbacks?


- + Hopefully fits in RAM \Rightarrow Lookups fast
- + Random access possible (just fetch the i th block)
- File size limited by how many references fit into one block
- Wasteful for small files

How could you store *Huge* files?

- Increase the Block size
- + Indirection! Make the index block link to other index blocks (like pagetables)
You can also mix that: First N pointers point to data blocks, next to indirect blocks, next to double-indirect blocks, etc.

Indexed Allocations With Inodes

What does an inode look like?



Indexed Allocations With Inodes

What is the maximum file size?

Assume that disk blocks are 8 KiB in size and a pointer to a disk block is 4 bytes long. An inode contains 12 pointers to direct blocks, and one pointer to a single, double, and triple indirect block, respectively.

Calculating the amount of pointers to blocks

Pointers per Block: $8 \text{ KiB} / 4 \text{ bytes} = 2^{13} / 2^2 = 2^{11} = 2048$

Direct pointers	12 pointers to blocks
One single-indirect block	2048 pointers to blocks
One double-indirect block	$2048 \cdot 2048$ pointers to blocks
One triple-indirect block	$2048 \cdot 2048 \cdot 2048$ pointers to blocks

Final result

$(12 + 2048 + 2048^2 + 2048^3) \cdot 8 \text{ KiB} \approx 64 \text{ TiB}$

File System Implementation

Hard links

Pointer **to the same inode**.

⇒ *Everything* is identical: Size, content, access time, mode, ...

Can therefore *not* cross file system boundaries.

Symlinks

Are (mostly) **normal files** containing **a filepath** with their own access time, mode, When programs access the file the OS transparently:

1. Reads the contents of the symlink file
2. Resolves the file path it read
3. Performs the operation on that path instead

Can cross file system boundaries or point to non-existent files or change what they point to if the file is moved, ...

Having Fun With Paths I

Hard and symlink renames

```
1 echo "Hello" > test.txt
2 ln test.txt hardlink.txt
3 ln -s test.txt symlink.txt
4
5 mv test.txt renamed_test.txt
6 // Is the hardlink / symlink broken?
```

Reaching for the stars

Does this work (on my machine ;)?

```
1 cd /tmp/
2 echo "Hello" > test.txt
3 ln test.txt ~/test_link.txt
```

No, as **/tmp** is mounted as tmpfs and / as ext4.

Hard and symlink renames

```
1 echo "Hello" > test.txt
2 ln test.txt hardlink.txt
3 ln -s test.txt symlink.txt
4
5 cp test.txt renamed_test.txt
6 rm test.txt
7 // Is the hardlink / symlink broken?
8
9 echo "Hello" > test.txt
10 // Is the hardlink / symlink still broken?
```

If you delete a hardlink, how can the FS know when it can delete the file?

If the refcount stored in the inode is zero

How can directories be implemented? What information is stored in them?

As a normal file with variable-length entries consisting of

- The filename
- The inode that represents that file

Which of the following data are typically stored in an inode?

- | | |
|---------------------------------|---|
| 1. filename | 8. name/location of hardlinks |
| 2. name of containing directory | 9. access rights |
| 3. file size | 10. timestamps (last access/modify) |
| 4. file type | 11. file contents |
| 5. number of symlinks to file | 12. ordered list of blocks occupied by file |
| 6. name/location of symlinks | |
| 7. number of hardlinks | |

Virtual File System

What is the purpose of the VFS layer in an OS?

Abstraction. Provide a high-level API that works no matter what file system you use (XFS, ext4, NTFS, SMB, IliadFUSE, ...).

What are some drawbacks / problems?

- Lowest common denominator: Might hide special FS features
- ⇒ Allow accessing underlying FS (e.g. via `ioctl`)
- ⇒ Lose compatibility with other file systems

What happens when you mount a filesystem?

E.g. `mount /dev/sda1 /mnt`.

Makes files from that file system accessible. It will make the given path (e.g. `/mnt`) the *root* of the new file system. Any access to `/tmp/*` is stripped of the `/tmp/` prefix and then searched in the mounted file system.

FUSE

File System in **US**erspace.

- Allows users to write their own file system without writing kernel code
- Runs as an unprivileged user process
- Is really awesome! You suddenly can use all your normal tools on whatever the FUSE filesystem exposes!

File System Cache


Is data persisted after a call to `write`? If not, why and when?

- Data is probably buffered by libc first
- Even after that, it doesn't directly get written to disk. Enter...The *File System Cache*!
- Why? Writing to disk every time is *reeeaaally slow*, so it is probably cached in memory and flushed in chunks

The standard question: Why is a cache even helpful?

- Temporal and spatial locality
- When exactly is it flushed to disk then? Is that even important to know/control?
- Using `flush` (for buffers) and `fsync` (for the page cache) and after some time by a daemon

Let's talk about the page cache



Sizing the page cache

How large would you make it? Do you give it a fixed size or let it grow?

Fix Easy to implement

Fix Can give hard guarantees (real-time systems anyone?)

Fix Can't adapt to different workloads

When do you load data in the page cache?

Read-ahead: Read more data in the cache than requested

- + Good sequential performance
- + Improved throughput if files are sequential
- Wasted time and memory if not needed

How could you implement `mmap` with the file system cache?

Just map the file system cache in the process's memory (and handle faults!)

⇒ Acts as a shared memory segment

You synchronize accesses yourself :(

`read()` and `write()` and the file system cache

- `read()`: Copy data from cache to your application buffers
- `write()`: Copy data from your application buffers to the file system cache


File system synchronizes access!

Modern File Systems

Crashing Is Fun


Your computer crashes (e.g. your power fails). What horrible death does your file system die?

A fun game for the whole family and your „why would I need backups“ crowd




The inode has a refcount of 1 but is no longer referenced!

Where is Data?



Both files might point to this block!

Where is Data?



Block still referenced but marked free

Or multiple directory entries with the same name, as another problem

What might be the outcome?

```
1 // create a file
2 echo "Hey" > a.txt
3 // And a second
4 echo "Hello" > b.txt
5 // CRASH
```

- Both files exist
- No files exist
- Only **a** exists
- Only **b** exists !!

Reordering is *allowed* (unless you take precautions)

How could you detect those inconsistencies?

Use the **fsck** (file system consistency check) program. And what could that do?

- Check whether all data blocks are referenced by *exactly one* inode
- Do all directories contain valid entries (and just one with a given name)
- Is every directory / file connected to the directory tree
- Is the refcount consistent with the number of hardlinks? What do you do if not?

$refcount > 0$ but no links? \Rightarrow Attach to **lost+found**

Update the refcount to be consistent

Can this fix your application data?

No! It just tries to keep the filesystem internally consistent, it won't find corrupted data blocks

Journal File System - Recovery

General principle

Walk over the journal and execute any outstanding entries.

Let's crash



What to do

The happy path, everything's nice ⇒ We didn't write anything!

⇒ Operation **failed** ⇒ Invalid checksum

⇒ Skip entry

⇒ Operation **failed** ⇒ Non-terminated journal entry

⇒ Retry and complete operation

⇒ Operation **successful** ⇒ Consistent state

Physical vs logical logging

What data could you log in the journal?

- **Logical logging:** Store a high level entry (like: „rename a to b“)
- **Physical logging:** Store the file system blocks that will be modified

What might be problems with the journal presented before?

You write *every block twice*! How could you make that less painful? Only journal *metadata*.

How could you design a safe system that only needs to write once?

A Log Structured File System (Copy on Write). What is the core idea there?

Only write to unused blocks.

Then you can *atomically* update indexing datastructures to point to the new blocks. If it crashes, you either have *all* of the new state or *exactly* the old state.