

Betriebssysteme

Tutorium 3

Peter Bohner

14. November 2024

ITEC - Operating Systems Group

Scheduling basics

What is a Scheduler? Why do we even need it?

What is a Scheduler? Why do we even need it?

- Maps processes to resources

What is a Scheduler? Why do we even need it?

- Maps processes to resources
- Ideally: Every process gets what it needs some day

What is a Scheduler? Why do we even need it?

- Maps processes to resources
- Ideally: Every process gets what it needs some day

What Schedulers do you know?

What is a Scheduler? Why do we even need it?

- Maps processes to resources
- Ideally: Every process gets what it needs some day

What Schedulers do you know?

- CPU-Scheduler: The classic

What is a Scheduler? Why do we even need it?

- Maps processes to resources
- Ideally: Every process gets what it needs some day

What Schedulers do you know?

- CPU-Scheduler: The classic
- Disk-Scheduler: Why have one?

What is a Scheduler? Why do we even need it?

- Maps processes to resources
- Ideally: Every process gets what it needs some day

What Schedulers do you know?

- CPU-Scheduler: The classic
- Disk-Scheduler: Why have one? Multiplexing but also efficiency!

What is a Scheduler? Why do we even need it?

- Maps processes to resources
- Ideally: Every process gets what it needs some day

What Schedulers do you know?

- CPU-Scheduler: The classic
- Disk-Scheduler: Why have one? Multiplexing but also efficiency!
- Network I/O:

What is a Scheduler? Why do we even need it?

- Maps processes to resources
- Ideally: Every process gets what it needs some day

What Schedulers do you know?

- CPU-Scheduler: The classic
- Disk-Scheduler: Why have one? Multiplexing but also efficiency!
- Network I/O: When to send packets, which packets to drop, QoL,...

What are the differences? When are they used?

What are the differences? When are they used?

- LTS: Decide which processes to put in the *run queue*

What are the differences? When are they used?

- LTS: Decide which processes to put in the *run queue*
- STS: Decide which process runs on the *CPU*

What are the differences? When are they used?

- LTS: Decide which processes to put in the *run queue*
 - STS: Decide which process runs on the *CPU*
 - MTS: Temporarily removes processes from main memory (and e.g. writes them out to disk)
- ⇒ Reduce degree of multiprogramming, make room in memory (and a few other reasons)



Process states

new

ready

Process states

new

ready

running

Process states

new

ready


running

waiting


Process states




Process states



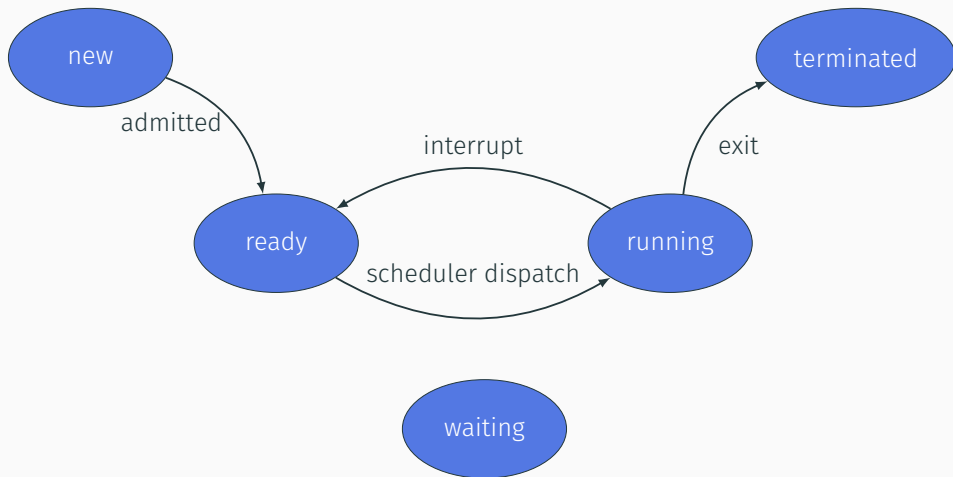
Process states




Process states



Process states



Process states



„I/O or event wait“? When does a process move from ready to waiting?

- Network / Disk I/O

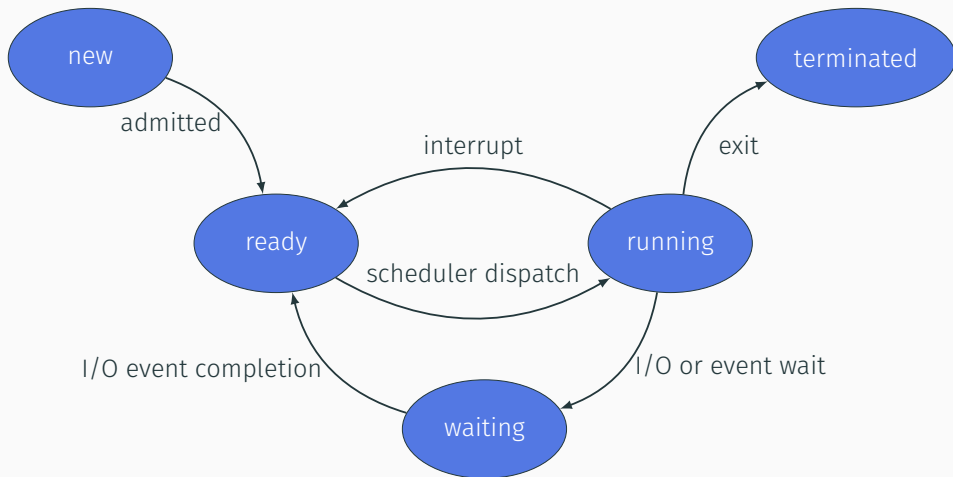
„I/O or event wait“? When does a process move from ready to waiting?

- Network / Disk I/O
- Mutex or other inter-process synchronisation

„I/O or event wait“? When does a process move from ready to waiting?

- Network / Disk I/O
- Mutex or other inter-process synchronisation
- Sleepyness

Process states



What makes a good Scheduler good?

Let's play scheduler!

What makes a good Scheduler good?

Let's play scheduler!

Some metrics

- Processor utilization: Percentage of working time
- Throughput: How many jobs do you finish?
- Turnaround time: Wallclock-time from submission to finish
- Waiting time: How long did it spend in the ready queue
- Response time: Time between submission of a request and first response (e.g. key press to echo on screen)

What does your hardware need to support to allow non-cooperative scheduling?

What does your hardware need to support to allow non-cooperative scheduling?

Timer Interrupts! Waiting for a cosmic ray to hit, a network package to arrive, a system call or any other random interrupt gets old fast :)

Scheduling - When to interrupt

Any guesses for how long a timeslice usually is?

Scheduling - When to interrupt

Any guesses for how long a timeslice usually is?

2ms - 200ms

- On windows it depends on the configuration (favor foreground / background processes)

Scheduling - When to interrupt

Any guesses for how long a timeslice usually is?

2ms - 200ms

- On windows it depends on the configuration (favor foreground / background processes)

Which setting has the longer timeslice?

Scheduling - When to interrupt

Any guesses for how long a timeslice usually is?

2ms - 200ms

- On windows it depends on the configuration (favor foreground / background processes)

Which setting has the longer timeslice?

20ms to 30ms for foreground, 180ms to 200ms for background

Any guesses for how long a timeslice usually is?

2ms - 200ms

- On windows it depends on the configuration (favor foreground / background processes)

Which setting has the longer timeslice?

20ms to 30ms for foreground, 180ms to 200ms for background

- Linux's „Completely Fair Scheduler“ adjusts them dynamically based on the priority, number of processes, ...

Scheduling - When to interrupt

Any guesses for how long a timeslice usually is?

2ms - 200ms

- On windows it depends on the configuration (favor foreground / background processes)

Which setting has the longer timeslice?

20ms to 30ms for foreground, 180ms to 200ms for background

- Linux's „Completely Fair Scheduler“ adjusts them dynamically based on the priority, number of processes, ...

Benefits of shorter/longer timeslices?

Scheduling - When to interrupt

Any guesses for how long a timeslice usually is?

2ms - 200ms

- On windows it depends on the configuration (favor foreground / background processes)

Which setting has the longer timeslice?

20ms to 30ms for foreground, 180ms to 200ms for background

- Linux's „Completely Fair Scheduler“ adjusts them dynamically based on the priority, number of processes, ...

Benefits of shorter/longer timeslices?

- Short: High interactivity, higher overhead

Scheduling - When to interrupt

Any guesses for how long a timeslice usually is?

2ms - 200ms

- On windows it depends on the configuration (favor foreground / background processes)


Which setting has the longer timeslice?

20ms to 30ms for foreground, 180ms to 200ms for background

- Linux's „Completely Fair Scheduler“ adjusts them dynamically based on the priority, number of processes, ...

Benefits of shorter/longer timeslices?

- Short: High interactivity, higher overhead
- Long: Lower interactivity, smaller overhead



XKCD 1542 - Scheduling Conflict

FRAGEN?

Bis nächste Woche :)