

Betriebssysteme

Tutorium 4

Peter Bohner

22. November 2024

ITEC - Operating Systems Group

Pitfalls - How would you implement this?

Pitfalls - How would you implement this?

- You would need to have future knowledge to figure out the job length!

Pitfalls - How would you implement this?

- You would need to have future knowledge to figure out the job length!
How do you solve this?

Pitfalls - How would you implement this?

- You would need to have future knowledge to figure out the job length!
How do you solve this?
- **Predict** the *future* based on *past* behaviour

Pitfalls - How would you implement this?

- You would need to have future knowledge to figure out the job length!
How do you solve this?
- **Predict** the *future* based on *past* behaviour
- Does this work?

Pitfalls - How would you implement this?

- You would need to have future knowledge to figure out the job length!
How do you solve this?
- **Predict** the *future* based on *past* behaviour
- Does this work?
„*THERE IS AS YET INSUFFICIENT DATA FOR A MEANINGFUL ANSWER.*“
~ Isaac Asimov, „*The Last Question*“ (Comic)

Pitfalls - How would you implement this?

- You would need to have future knowledge to figure out the job length!
How do you solve this?
- **Predict** the *future* based on *past* behaviour
- Does this work?

„THERE IS AS YET INSUFFICIENT DATA FOR A MEANINGFUL ANSWER.“

~ Isaac Asimov, *„The Last Question“* (Comic)

You need some balanced initial value. Not *that* big of a deal with preemption though. Why?

Pitfalls - How would you implement this?

- You would need to have future knowledge to figure out the job length!
How do you solve this?
- **Predict** the *future* based on *past* behaviour
- Does this work?

„THERE IS AS YET INSUFFICIENT DATA FOR A MEANINGFUL ANSWER.“

~ Isaac Asimov, *„The Last Question“* (Comic)

You need some balanced initial value. Not *that* big of a deal with preemption though. Why?


Interrupt the process after the estimated time is over.

What is priority scheduling? Why would you use it?

What is priority scheduling? Why would you use it?

- Each process is assigned a priority
- The process with the highest priority is chosen

Multi-Level Feedback Queues



How it works

- All processes start in the highest queue
 - When they use up their timeslice and are preempted, they descend
 - If they block before, they stay in the level (optionally: Are moved up)
- ⇒ I/O bound processes rise to the top and react quickly, CPU bound processes get longer timeslices but less often

Multi-Level Feedback Queues



How it works

- All processes start in the highest queue
 - When they use up their timeslice and are preempted, they descend
 - If they block before, they stay in the level (optionally: Are moved up)
- ⇒ I/O bound processes rise to the top and react quickly, CPU bound processes get longer timeslices but less often

Multi-Level Feedback Queues

Short time slice

		C	D
--	--	---	---

B didn't use it fully

	B		
--	---	--	--

A used its full timeslice

A			
---	--	--	--

How it works

- All processes start in the highest queue
 - When they use up their timeslice and are preempted, they descend
 - If they block before, they stay in the level (optionally: Are moved up)
- ⇒ I/O bound processes rise to the top and react quickly, CPU bound processes get longer timeslices but less often

A Flawless Scheduling Algorithm?

Consider the waiting time

A Flawless Scheduling Algorithm?

Consider the waiting time

- A few I/O-bound jobs could saturate the CPU!
- ⇒ The lower-level processes starve

A Flawless Scheduling Algorithm?

Consider the waiting time

- A few I/O-bound jobs could saturate the CPU!
- ⇒ The lower-level processes starve

How could you fix this?

A Flawless Scheduling Algorithm?

Consider the waiting time

- A few I/O-bound jobs could saturate the CPU!
- ⇒ The lower-level processes starve

How could you fix this?

- E.g. reset the whole thing after a given interval, so all start in the highest level again

A Flawless Scheduling Algorithm?

Consider the waiting time

- A few I/O-bound jobs could saturate the CPU!
- ⇒ The lower-level processes starve

How could you fix this?

- E.g. reset the whole thing after a given interval, so all start in the highest level again
- "Boost" processes that waited for a long time

A Flawless Scheduling Algorithm?

Consider the waiting time

- A few I/O-bound jobs could saturate the CPU!
- ⇒ The lower-level processes starve

How could you fix this?

- E.g. reset the whole thing after a given interval, so all start in the highest level again
- "Boost" processes that waited for a long time

What metrics does it optimize?

- Utilization? Turnaround time? Throughput? Waiting time? Response time?

A Flawless Scheduling Algorithm?

Consider the waiting time

- A few I/O-bound jobs could saturate the CPU!
- ⇒ The lower-level processes starve

How could you fix this?

- E.g. reset the whole thing after a given interval, so all start in the highest level again
- "Boost" processes that waited for a long time

What metrics does it optimize?

- Utilization? Turnaround time? Throughput? Waiting time? Response time?
- Prefer I/O bound, prefer short jobs, group the rest based on their needs

Process switching

What does the kernel need to do when switching processes?

- Adjust Instruction, Base and Frame Pointer – okay

What does the kernel need to do when switching processes?

- Adjust Instruction, Base and Frame Pointer – okay

The process shouldn't notice it was interrupted! What else do you need to adjust?

What does the kernel need to do when switching processes?

- Adjust Instruction, Base and Frame Pointer – okay

The process shouldn't notice it was interrupted! What else do you need to adjust?

- General purpose register

What does the kernel need to do when switching processes?

- Adjust Instruction, Base and Frame Pointer – okay

The process shouldn't notice it was interrupted! What else do you need to adjust?

- General purpose register
- Address space

What does the kernel need to do when switching processes?

- Adjust Instruction, Base and Frame Pointer – okay

The process shouldn't notice it was interrupted! What else do you need to adjust?

- General purpose register
- Address space
- And a bit of housekeeping:

What does the kernel need to do when switching processes?

- Adjust Instruction, Base and Frame Pointer – okay

The process shouldn't notice it was interrupted! What else do you need to adjust?

- General purpose register
- Address space
- And a bit of housekeeping: Open files, scheduling priorities, ...

⇒ These things make up the PCB - The

What does the kernel need to do when switching processes?

- Adjust Instruction, Base and Frame Pointer – okay

The process shouldn't notice it was interrupted! What else do you need to adjust?

- General purpose register
- Address space
- And a bit of housekeeping: Open files, scheduling priorities, ...

⇒ These things make up the PCB - The Process Control Block

What does the kernel need to do when switching processes?

- Adjust Instruction, Base and Frame Pointer – okay

The process shouldn't notice it was interrupted! What else do you need to adjust?

- General purpose register
- Address space
- And a bit of housekeeping: Open files, scheduling priorities, ...

⇒ These things make up the PCB - The Process Control Block

Where are the PCBs stored? In user or kernel space?

What does the kernel need to do when switching processes?

- Adjust Instruction, Base and Frame Pointer – okay

The process shouldn't notice it was interrupted! What else do you need to adjust?

- General purpose register
- Address space
- And a bit of housekeeping: Open files, scheduling priorities, ...

⇒ These things make up the PCB - The Process Control Block

Where are the PCBs stored? In user or kernel space?

- Kernel space! Users shouldn't be able to modify them

Are the PCBs always valid?

Are the PCBs always valid?

No! Some parts (registers, PC, etc) only when the process is not running. Why?

Are the PCBs always valid?

No! Some parts (registers, PC, etc) only when the process is not running. Why?
They are saved when it is switched out.


In pictures

P1 ——— Living life

P2


Process switching

In pictures




Process switching

In pictures



Process switching

In pictures



As text

1. Transition to the kernel (How?

As text

1. Transition to the kernel (How? Interrupt, Exception, Syscall)

As text

1. Transition to the kernel (How? Interrupt, Exception, Syscall)
2. Save the context of the process. What was that again?

As text

1. Transition to the kernel (How? Interrupt, Exception, Syscall)
2. Save the context of the process. What was that again? Registers, Program counter, Stack pointer, etc.

As text

1. Transition to the kernel (How? Interrupt, Exception, Syscall)
2. Save the context of the process. What was that again? Registers, Program counter, Stack pointer, etc.
3. Where do you save it to? A per-process kernel stack typically (as it is often moved there automatically) and then move it to the PCB.

As text

1. Transition to the kernel (How? Interrupt, Exception, Syscall)
2. Save the context of the process. What was that again? Registers, Program counter, Stack pointer, etc.
3. Where do you save it to? A per-process kernel stack typically (as it is often moved there automatically) and then move it to the PCB.
4. Restore the context of the next process

As text

1. Transition to the kernel (How? Interrupt, Exception, Syscall)
2. Save the context of the process. What was that again? Registers, Program counter, Stack pointer, etc.
3. Where do you save it to? A per-process kernel stack typically (as it is often moved there automatically) and then move it to the PCB.
4. Restore the context of the next process
5. Leave kernel mode and transfer control to the PC of the next process

Threads



What are processes, address spaces and threads? How do they relate to each other?

What are processes, address spaces and threads? How do they relate to each other?

- A thread is an *entity of execution*, the personification of control flow

What are processes, address spaces and threads? How do they relate to each other?

- A thread is an *entity of execution*, the personification of control flow
- A thread lives in an address space, i.e. all the addresses that it can access and the data that is stored there

What are processes, address spaces and threads? How do they relate to each other?

- A thread is an *entity of execution*, the personification of control flow
- A thread lives in an address space, i.e. all the addresses that it can access and the data that is stored there
- Thread + Address Space = Process

Spawn a few threads using pthreads!

Write a small program that creates five threads using the pthread library. Each thread should print its number (e.g., *Hello, I am 4*) and the main program should wait for each thread to exit.

Thread models — One To One

One To One

new Thread()

new Thread()

Process

A diagram illustrating the One To One thread model. It features a large light gray rectangular area. Inside this area, on the right side, is a smaller square box with a black border, labeled "Process". To the left of this box, there are two lines of text, each in an italicized font: "new Thread()" and "new Thread()".


Problems?

Thread models — One To One

One To One

new Thread()


new Thread()



Problems?

Thread models — One To One


One To One



Problems?


Thread models — One To One

One To One



Thread models — One To One

One To One



Problems?

Problems and benefits of *One To One*?

Problems and benefits of *One To One*?

- + Scales with core count

Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff

Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff
- + Blocking does not affect other threads

Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff
- + Blocking does not affect other threads
- + Can piggy-back onto the OS scheduler

Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff
- + Blocking does not affect other threads
- + Can piggy-back onto the OS scheduler
- **Must** piggy-back onto the OS scheduler

Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff
- + Blocking does not affect other threads
- + Can piggy-back onto the OS scheduler
- **Must** piggy-back onto the OS scheduler
- Relatively high overhead due to context switches

Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff
- + Blocking does not affect other threads
- + Can piggy-back onto the OS scheduler
- **Must** piggy-back onto the OS scheduler
- Relatively high overhead due to context switches
- Relatively high overhead *when creating one*

Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff
- + Blocking does not affect other threads
- + Can piggy-back onto the OS scheduler
- **Must** piggy-back onto the OS scheduler
- Relatively high overhead due to context switches
- Relatively high overhead *when creating one!!*

Many to One

new Thread()

User Stack

Process



Kernel Stack

Many to One

new Thread()

User Stack

new Thread()

User Stack

Process



Kernel Stack

Many to One

new Thread()

User Stack

new Thread()

User Stack

new Thread()

User Stack


Process



Kernel Stack


Thread models

Many to One




Thread models

Many to One



Thread models

Many to One



Problems and benefits of *Many To One*?

Do they improve anything?

Problems and benefits of *Many To One*?

Do they improve anything?

- + Scales with core count?

Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core

Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- + Conceptually easy — the OS does the hard stuff?

Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much

Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- + **Blocking does not affect other threads?**

Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads

Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- + Can piggy-back onto the OS scheduler?

Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler

Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler
- **Must** piggy-back onto the OS scheduler?

Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler

Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- **Relatively high overhead due to context switches?**

Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- + Low overhead during context switches

Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- + Low overhead during context switches
- **Relatively high overhead *when creating one*?**

Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- + Low overhead during context switches
- + Low overhead when creating one

Thread models - Many To One

Do you know a programming language / runtime using that?

E.g. *nodejs* using its „event loop“

A small excursion - Structured Programming

Control flow should fall into one of four patterns:

- Sequence: One block is executed after another

Thread models - Many To One

Do you know a programming language / runtime using that?

E.g. *nodejs* using its „event loop“

A small excursion - Structured Programming

Control flow should fall into one of four patterns:

- Sequence: One block is executed after another
- Selection: One or more are executed (i.e. an *if*-statement)

Thread models - Many To One

Do you know a programming language / runtime using that?

E.g. *nodejs* using its „event loop“

A small excursion - Structured Programming

Control flow should fall into one of four patterns:

- Sequence: One block is executed after another
- Selection: One or more are executed (i.e. an *if*-statement)
- Iteration: A block is executed more than once (i.e. a loop)

Thread models - Many To One

Do you know a programming language / runtime using that?

E.g. *nodejs* using its „event loop“

A small excursion - Structured Programming

Control flow should fall into one of four patterns:

- Sequence: One block is executed after another
- Selection: One or more are executed (i.e. an *if*-statement)
- Iteration: A block is executed more than once (i.e. a loop)
- Recursion: A block calls itself until an exit condition is met (i.e. recursion!)

Do you know any keyword in C which *doesn't* quite adhere to that but can instead totally spaghetti your control flow?

Thread models - Many To One

Do you know a programming language / runtime using that?

E.g. *nodejs* using its „event loop“

A small excursion - Structured Programming

Control flow should fall into one of four patterns:

- Sequence: One block is executed after another
- Selection: One or more are executed (i.e. an *if*-statement)
- Iteration: A block is executed more than once (i.e. a loop)
- Recursion: A block calls itself until an exit condition is met (i.e. recursion!)

Do you know any keyword in C which *doesn't* quite adhere to that but can instead totally spaghetti your control flow? *goto*

Thread models - Many To One

Do you know a programming language / runtime using that?

E.g. *nodejs* using its „event loop“

A small excursion - Structured Programming

Control flow should fall into one of four patterns:

- Sequence: One block is executed after another
- Selection: One or more are executed (i.e. an *if*-statement)
- Iteration: A block is executed more than once (i.e. a loop)
- Recursion: A block calls itself until an exit condition is met (i.e. recursion!)

Do you know any keyword in C which *doesn't* quite adhere to that but can instead totally spaghetti your control flow? *goto*

Famous paper by a proponent of Structured Programming:

„*Go To Statement Considered Harmful*“ by Edsger W. Dijkstra

And what about threads?

- Can outlive the methods they were spawned in

And what about threads?

- Can outlive the methods they were spawned in
- Can use variables and fields after they went out of scope in a method

And what about threads?

- Can outlive the methods they were spawned in
- Can use variables and fields after they went out of scope in a method
- Can split up or transfer their control flow arbitrarily

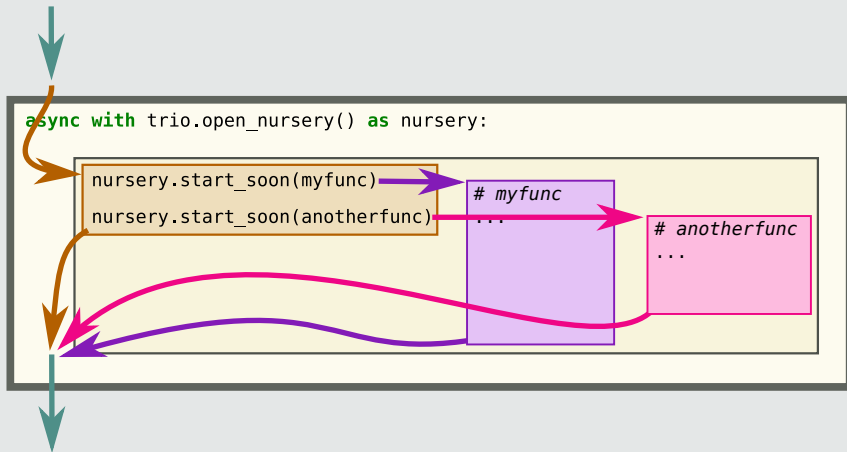
And what about threads?

- Can outlive the methods they were spawned in
- Can use variables and fields after they went out of scope in a method
- Can split up or transfer their control flow arbitrarily

So that might sound familiar...

Thread models - Many To One

Structured Concurrency



Taken from vorpus.org

Nice, but what does this have to do with ULTs?

Nice, but what does this have to do with ULTs?

- Spawning lots of threads for small operations *is too slow otherwise*

Further reading:

[Notes on Structured Concurrency](#)

[ULTs and Structured concurrency in Java - Project Loom](#)

Thread models - Many To Many

Many To Many

new Thread()

User Stack

Process



Kernel Stack

Kernel Stack

Thread models - Many To Many

Many To Many

new Thread()

User Stack

new Thread()

User Stack

Process



Kernel Stack

Kernel Stack

Thread models - Many To Many

Many To Many

new Thread()

User Stack

new Thread()

User Stack

new Thread()

User Stack

Process




Kernel Stack

Kernel Stack


Thread models - Many To Many

Many To Many




Thread models - Many To Many

Many To Many



Thread models - Many To Many

Many To Many



Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count?

Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count

Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff?

Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much

Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + **Blocking does not affect other threads?**

Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (*upcalls*) and does *not* pause the whole kernel level thread

Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (*upcalls*) and does *not* pause the whole kernel level thread
- + Can piggy-back onto the OS scheduler?

Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (*upcalls*) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler

Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (*upcalls*) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler
- **Must** piggy-back onto the OS scheduler?

Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (*upcalls*) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler
- + Can implement its *own* scheduler

Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (*upcalls*) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- Relatively high overhead due to context switches?

Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (*upcalls*) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- + Low overhead during context switches (unless you need to interact with the kernel threads, e.g. to schedule between them)

Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (*upcalls*) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- + Low overhead during context switches (unless you need to interact with the kernel threads, e.g. to schedule between them)
- Relatively high overhead *when creating one?*

Problems and benefits of *Many To Many*?


Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (*upcalls*) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- + Low overhead during context switches (unless you need to interact with the kernel threads, e.g. to schedule between them)
- + Low overhead when creating one

Thread models - Many To Many

Problems the second

You have all attended SWT 1! So let's have a look.




And preemption is now possible, which might complicate user code.

Thread models - Many To Many

Problems the second


You have all attended SWT 1! So let's have a look.



Thread models - Many To Many

Problems the second


You have all attended SWT 1! So let's have a look.



Thread models - Many To Many

Problems the second

You have all attended SWT 1! So let's have a look.



And preemption is now possible, which might complicate user code.

What events can trigger a context switch?

What events can trigger a context switch?

- Voluntary:

What events can trigger a context switch?

- Voluntary: yield, blocking system call

What events can trigger a context switch?

- Voluntary: yield, blocking system call
- Involuntary:
 - interrupts, exceptions, syscalls

What events can trigger a context switch?

- Voluntary: yield, blocking system call
- Involuntary:
 - interrupts, exceptions, syscalls
 - end of time-slice

What events can trigger a context switch?

- Voluntary: yield, blocking system call
- Involuntary:
 - interrupts, exceptions, syscalls
 - end of time-slice
 - high priority thread becoming ready

What events can trigger a context switch?

What events can trigger a context switch?

- Most libraries only support *cooperative scheduling*

What events can trigger a context switch?

- Most libraries only support *cooperative scheduling*
- Why is switching with preemption, interrupts, blocking system calls hard?

What events can trigger a context switch?

- Most libraries only support *cooperative scheduling*
- Why is switching with preemption, interrupts, blocking system calls hard?
Kernel is not aware of the ULTs and will return where it came from — *but not call out to the scheduler and carry on*
- The benefit of platforms: How can **Java** (using Project Loom) or **Node.js** switch on most of the above?

What events can trigger a context switch?

- Most libraries only support *cooperative scheduling*
- Why is switching with preemption, interrupts, blocking system calls hard?
Kernel is not aware of the ULTs and will return where it came from — *but not call out to the scheduler and carry on*
- The benefit of platforms: How can **Java** (using Project Loom) or **Node.js** switch on most of the above?
You can not execute syscalls directly, but need to call library methods!
Suspension points can be inserted there.

„Jobs are either I/O-bound or compute-bound. In neither case would user-level threads be a win. Why would one go for pure user-level threads at all?“


„Jobs are either I/O-bound or compute-bound. In neither case would user-level threads be a win. Why would one go for pure user-level threads at all?“

- Program structure (e.g. Structured Concurrency, channels or just easier pipelines)


„Jobs are either I/O-bound or compute-bound. In neither case would user-level threads be a win. Why would one go for pure user-level threads at all?“

- Program structure (e.g. Structured Concurrency, channels or just easier pipelines)
- The same or higher I/O throughput if on an abstracted platform

Should a fork do this?



Should a fork do this?



Who is aware of the fork?

- The thread that executed the fork and the child

Who is aware of the fork?

- The thread that executed the fork and the child
- Who else?

Who is aware of the fork?

- The thread that executed the fork and the child
- Who else? Nobody!

Who is aware of the fork?

- The thread that executed the fork and the child
- Who else? Nobody!

Can you foresee any problems?

- If you copy the thread it might do weird things

Who is aware of the fork?

- The thread that executed the fork and the child
- Who else? Nobody!

Can you foresee any problems?

- If you copy the thread it might do weird things
- Locks are memory and will be copied as well! What if a thread holds a critical lock when it is forked?

Who is aware of the fork?

- The thread that executed the fork and the child
- Who else? Nobody!

Can you foresee any problems?

- If you copy the thread it might do weird things
- Locks are memory and will be copied as well! What if a thread holds a critical lock when it is forked?

Do you need thread duplication for our shell example?

Who is aware of the fork?

- The thread that executed the fork and the child
- Who else? Nobody!

Can you foresee any problems?

- If you copy the thread it might do weird things
- Locks are memory and will be copied as well! What if a thread holds a critical lock when it is forked?

Do you need thread duplication for our shell example?

No, we `exec` anyways.

Threads — Forking

Who is aware of the fork?

- The thread that executed the fork and the child
- Who else? Nobody!

Can you foresee any problems?

- If you copy the thread it might do weird things
- Locks are memory and will be copied as well! What if a thread holds a critical lock when it is forked?

Do you need thread duplication for our shell example?

No, we `exec` anyways.

Summary

fork is not as simple as it once was. Is it still a good abstraction?

What is the difference?

What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread

What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

Why would you need kernel *mode* threads?

What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

Why would you need kernel *mode* threads?

- Sometimes the OS needs to do some housekeeping that doesn't belong in a system call

What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

Why would you need kernel *mode* threads?

- Sometimes the OS needs to do some housekeeping that doesn't belong in a system call
- Free some pages, swap something in and out of memory

What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

Why would you need kernel *mode* threads?

- Sometimes the OS needs to do some housekeeping that doesn't belong in a system call
- Free some pages, swap something in and out of memory
- Flush pages from the disk cache to the hard disk

What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

Why would you need kernel *mode* threads?

- Sometimes the OS needs to do some housekeeping that doesn't belong in a system call
- Free some pages, swap something in and out of memory
- Flush pages from the disk cache to the hard disk
- Perform VFS Checkpointing



What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

Why would you need kernel *mode* threads?

- Sometimes the OS needs to do some housekeeping that doesn't belong in a system call
- Free some pages, swap something in and out of memory
- Flush pages from the disk cache to the hard disk
- Perform VFS Checkpointing
- ...

F R A G E N ?



XKCD 292 - Goto

Bis nächste Woche :)