

# Betriebssysteme

## 5. Tutorium - Threading, Segmentation, MM

---

Peter Bohner

28. November 2024

ITEC - Operating Systems Group

# Threads

---

What are processes, address spaces and threads? How do they relate to each other?

What are processes, address spaces and threads? How do they relate to each other?

- A thread is an *entity of execution*, the personification of control flow

What are processes, address spaces and threads? How do they relate to each other?

- A thread is an *entity of execution*, the personification of control flow
- A thread lives in an address space, i.e. all the addresses that it can access and the data that is stored there

What are processes, address spaces and threads? How do they relate to each other?

- A thread is an *entity of execution*, the personification of control flow
- A thread lives in an address space, i.e. all the addresses that it can access and the data that is stored there
- Thread + Address Space = Process

## Spawn a few threads using pthreads!

Write a small program that creates five threads using the pthread library. Each thread should print its number (e.g., **Hello, I am 4**) and the main program should wait for each thread to exit.

## Thread models — One To One

### One To One

`new Thread()`

`new Thread()`

Process

A diagram illustrating the One To One thread model. A large light gray rectangle represents the operating system or hardware. Inside this rectangle, on the left, are two instances of the `new Thread()` method. On the right, there is a smaller, outlined rectangle labeled "Process". This visualizes that each thread created in this model corresponds to a separate process.

### Problems?

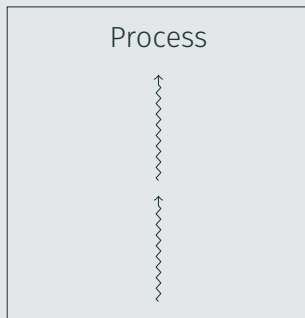


## Thread models — One To One

### One To One

`new Thread()`


`new Thread()`



### Problems?

# Thread models — One To One


## One To One



## Problems?


# Thread models — One To One

## One To One



# Thread models — One To One

## One To One



## Problems?

Problems and benefits of *One To One*?

### Problems and benefits of *One To One*?

- + Scales with core count

## Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff

### Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff
- + Blocking does not affect other threads



### Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff
- + Blocking does not affect other threads
- + Can piggy-back onto the OS scheduler

## Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff
- + Blocking does not affect other threads
- + Can piggy-back onto the OS scheduler
- **Must** piggy-back onto the OS scheduler

## Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff
- + Blocking does not affect other threads
- + Can piggy-back onto the OS scheduler
- **Must** piggy-back onto the OS scheduler
- Relatively high overhead due to context switches

## Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff
- + Blocking does not affect other threads
- + Can piggy-back onto the OS scheduler
- **Must** piggy-back onto the OS scheduler
- Relatively high overhead due to context switches
- Relatively high overhead *when creating one*

## Problems and benefits of *One To One*?

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff
- + Blocking does not affect other threads
- + Can piggy-back onto the OS scheduler
- **Must** piggy-back onto the OS scheduler
- Relatively high overhead due to context switches
- Relatively high overhead *when creating one!!*

## Many to One

`new Thread()`

User Stack

Process



Kernel Stack

# Thread models

## Many to One

`new Thread()`

User Stack

`new Thread()`

User Stack

Process



Kernel Stack


## Many to One






# Thread models

## Many to One




# Thread models

## Many to One



# Thread models

## Many to One



Problems and benefits of *Many To One*?

Do they improve anything?

### Problems and benefits of *Many To One*?

Do they improve anything?

- + Scales with core count?

### Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core

### Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- + Conceptually easy — the OS does the hard stuff?

### Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much



## Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- + **Blocking does not affect other threads?**

## Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads

### Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- + Can piggy-back onto the OS scheduler?

### Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler

### Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler
- **Must** piggy-back onto the OS scheduler?

## Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler

## Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- **Relatively high overhead due to context switches?**

## Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- + Low overhead during context switches



## Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- + Low overhead during context switches
- Relatively high overhead *when creating one?*

## Problems and benefits of *Many To One*?

Do they improve anything?

- Can only use one core
- Harder to implement — the OS doesn't help you much
- Blocking *does* affect other threads
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- + Low overhead during context switches
- + Low overhead when creating one

## Thread models - Many To One

Do you know a programming language / runtime using that?

E.g. `nodejs` using its „event loop“

### A small excursion - Structured Programming

Control flow should fall into one of four patterns:

- Sequence: One block is executed after another

## Thread models - Many To One

Do you know a programming language / runtime using that?

E.g. `nodejs` using its „event loop“

### A small excursion - Structured Programming

Control flow should fall into one of four patterns:

- Sequence: One block is executed after another
- Selection: One or more are executed (i.e. an `if`-statement)

## Thread models - Many To One

Do you know a programming language / runtime using that?

E.g. `nodejs` using its „event loop“

### A small excursion - Structured Programming

Control flow should fall into one of four patterns:

- Sequence: One block is executed after another
- Selection: One or more are executed (i.e. an `if`-statement)
- Iteration: A block is executed more than once (i.e. a loop)

## Thread models - Many To One

Do you know a programming language / runtime using that?

E.g. `node.js` using its „event loop“

### A small excursion - Structured Programming

Control flow should fall into one of four patterns:

- Sequence: One block is executed after another
- Selection: One or more are executed (i.e. an `if`-statement)
- Iteration: A block is executed more than once (i.e. a loop)
- Recursion: A block calls itself until an exit condition is met (i.e. recursion!)

Do you know any keyword in C which *doesn't* quite adhere to that but can instead totally spaghetti your control flow?

## Thread models - Many To One

Do you know a programming language / runtime using that?

E.g. `node.js` using its „event loop“

### A small excursion - Structured Programming

Control flow should fall into one of four patterns:

- Sequence: One block is executed after another
- Selection: One or more are executed (i.e. an `if`-statement)
- Iteration: A block is executed more than once (i.e. a loop)
- Recursion: A block calls itself until an exit condition is met (i.e. recursion!)

Do you know any keyword in C which *doesn't* quite adhere to that but can instead totally spaghetti your control flow? `goto`

# Thread models - Many To One

Do you know a programming language / runtime using that?

E.g. `node.js` using its „event loop“

## A small excursion - Structured Programming

Control flow should fall into one of four patterns:

- Sequence: One block is executed after another
- Selection: One or more are executed (i.e. an `if`-statement)
- Iteration: A block is executed more than once (i.e. a loop)
- Recursion: A block calls itself until an exit condition is met (i.e. recursion!)

Do you know any keyword in C which *doesn't* quite adhere to that but can instead totally spaghetti your control flow? `goto`

Famous paper by a proponent of Structured Programming:

„*Go To Statement Considered Harmful*“ by Edsger W. Dijkstra



### And what about threads?

- Can outlive the methods they were spawned in

### And what about threads?

- Can outlive the methods they were spawned in
- Can use variables and fields after they went out of scope in a method

### And what about threads?

- Can outlive the methods they were spawned in
- Can use variables and fields after they went out of scope in a method
- Can split up or transfer their control flow arbitrarily

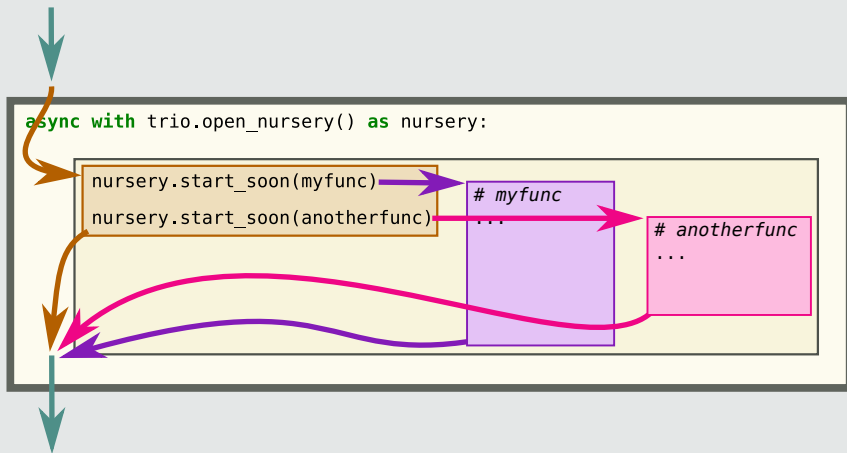
### And what about threads?

- Can outlive the methods they were spawned in
- Can use variables and fields after they went out of scope in a method
- Can split up or transfer their control flow arbitrarily

So that might sound familiar...

# Thread models - Many To One

## Structured Concurrency



Taken from [vorpus.org](http://vorpus.org)

Nice, but what does this have to do with ULTs?

Nice, but what does this have to do with ULTs?

- Spawning lots of threads for small operations *is too slow otherwise*

Further reading:

[Notes on Structured Concurrency](#)

[ULTs and Structured concurrency in Java - Project Loom](#)

# Thread models - Many To Many

## Many To Many

`new Thread()`

User Stack

Process





Kernel Stack

Kernel Stack





# Thread models - Many To Many

## Many To Many




# Thread models - Many To Many

## Many To Many




# Thread models - Many To Many

## Many To Many




# Thread models - Many To Many

## Many To Many



# Thread models - Many To Many

## Many To Many



### Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

### Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count?

### Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count



### Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- + Conceptually easy — the OS does the hard stuff?

## Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much

## Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + **Blocking does not affect other threads?**

## Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (**upcalls**) and does *not* pause the whole kernel level thread

### Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (**upcalls**) and does *not* pause the whole kernel level thread
- + Can piggy-back onto the OS scheduler?

### Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (**upcalls**) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler

## Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (**upcalls**) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler
- **Must** piggy-back onto the OS scheduler?

## Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (**upcalls**) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler



### Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (**upcalls**) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- Relatively high overhead due to context switches?

## Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (**upcalls**) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- + Low overhead during context switches (unless you need to interact with the kernel threads, e.g. to schedule between them)

## Problems and benefits of *Many To Many*?

Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (**upcalls**) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- + Low overhead during context switches (unless you need to interact with the kernel threads, e.g. to schedule between them)
- Relatively high overhead *when creating one?*

## Problems and benefits of *Many To Many*?


Important: The kernel *knows about the user level scheduler*

- + Scales with core count
- Harder to implement — the OS doesn't help you much
- + Blocking does not affect other threads as the kernel informs the user scheduler (**upcalls**) and does *not* pause the whole kernel level thread
- Can *not* piggy-back onto the OS scheduler
- + Can implement its **own** scheduler
- + Low overhead during context switches (unless you need to interact with the kernel threads, e.g. to schedule between them)
- + Low overhead when creating one

# Thread models - Many To Many

## Problems the second

You have all attended SWT 1! So let's have a look.




And preemption is now possible, which might complicate user code.

# Thread models - Many To Many

## Problems the second


You have all attended SWT 1! So let's have a look.



# Thread models - Many To Many

## Problems the second


You have all attended SWT 1! So let's have a look.



# Thread models - Many To Many

## Problems the second

You have all attended SWT 1! So let's have a look.



And preemption is now possible, which might complicate user code.



What events can trigger a context switch?

What events can trigger a context switch?

- Voluntary:

What events can trigger a context switch?

- Voluntary: yield, blocking system call

### What events can trigger a context switch?

- Voluntary: yield, blocking system call
- Involuntary:
  - interrupts, exceptions, syscalls

### What events can trigger a context switch?

- Voluntary: yield, blocking system call
- Involuntary:
  - interrupts, exceptions, syscalls
  - end of time-slice

### What events can trigger a context switch?

- Voluntary: yield, blocking system call
- Involuntary:
  - interrupts, exceptions, syscalls
  - end of time-slice
  - high priority thread becoming ready

What events can trigger a context switch?

What events can trigger a context switch?

- Most libraries only support *cooperative scheduling*



### What events can trigger a context switch?

- Most libraries only support *cooperative scheduling*
- Why is switching with preemption, interrupts, blocking system calls hard?

### What events can trigger a context switch?

- Most libraries only support *cooperative scheduling*
- Why is switching with preemption, interrupts, blocking system calls hard?  
Kernel is not aware of the ULTs and will return where it came from — *but not call out to the scheduler and carry on*
- The benefit of platforms: How can **Java** (using Project Loom) or **Node.js** switch on most of the above?

### What events can trigger a context switch?

- Most libraries only support *cooperative scheduling*
- Why is switching with preemption, interrupts, blocking system calls hard?  
Kernel is not aware of the ULTs and will return where it came from — *but not call out to the scheduler and carry on*
- The benefit of platforms: How can **Java** (using Project Loom) or **Node.js** switch on most of the above?  
You can not execute syscalls directly, but need to call library methods!  
Suspension points can be inserted there.

„Jobs are either I/O-bound or compute-bound. In neither case would user-level threads be a win. Why would one go for pure user-level threads at all?“


„Jobs are either I/O-bound or compute-bound. In neither case would user-level threads be a win. Why would one go for pure user-level threads at all?“

- Program structure (e.g. Structured Concurrency, channels or just easier pipelines)


„Jobs are either I/O-bound or compute-bound. In neither case would user-level threads be a win. Why would one go for pure user-level threads at all?“

- Program structure (e.g. Structured Concurrency, channels or just easier pipelines)
- The same or higher I/O throughput if on an abstracted platform

Should a fork do this?



Should a fork do this?





### Who is aware of the fork?

- The thread that executed the fork and the child

### Who is aware of the fork?

- The thread that executed the fork and the child
- Who else?

### Who is aware of the fork?

- The thread that executed the fork and the child
- Who else? Nobody!

### Who is aware of the fork?

- The thread that executed the fork and the child
- Who else? Nobody!

### Can you foresee any problems?

- If you copy the thread it might do weird things

### Who is aware of the fork?

- The thread that executed the fork and the child
- Who else? Nobody!

### Can you foresee any problems?

- If you copy the thread it might do weird things
- Locks are memory and will be copied as well! What if a thread holds a critical lock when it is forked?

### Who is aware of the fork?

- The thread that executed the fork and the child
- Who else? Nobody!

### Can you foresee any problems?

- If you copy the thread it might do weird things
- Locks are memory and will be copied as well! What if a thread holds a critical lock when it is forked?

### Do you need thread duplication for our shell example?

### Who is aware of the fork?

- The thread that executed the fork and the child
- Who else? Nobody!

### Can you foresee any problems?

- If you copy the thread it might do weird things
- Locks are memory and will be copied as well! What if a thread holds a critical lock when it is forked?

### Do you need thread duplication for our shell example?

No, we `exec` anyways.

# Threads — Forking

## Who is aware of the fork?

- The thread that executed the fork and the child
- Who else? Nobody!

## Can you foresee any problems?

- If you copy the thread it might do weird things
- Locks are memory and will be copied as well! What if a thread holds a critical lock when it is forked?

## Do you need thread duplication for our shell example?

No, we `exec` anyways.

## Summary

`fork` is not as simple as it once was. Is it still a good abstraction?



What is the difference?

### What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread

### What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

### What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

### Why would you need kernel *mode* threads?

### What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

### Why would you need kernel *mode* threads?

- Sometimes the OS needs to do some housekeeping that doesn't belong in a system call

### What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

### Why would you need kernel *mode* threads?

- Sometimes the OS needs to do some housekeeping that doesn't belong in a system call
- Free some pages, swap something in and out of memory

### What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

### Why would you need kernel *mode* threads?

- Sometimes the OS needs to do some housekeeping that doesn't belong in a system call
- Free some pages, swap something in and out of memory
- Flush pages from the disk cache to the hard disk

### What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

### Why would you need kernel *mode* threads?

- Sometimes the OS needs to do some housekeeping that doesn't belong in a system call
- Free some pages, swap something in and out of memory
- Flush pages from the disk cache to the hard disk
- Perform VFS Checkpointing



### What is the difference?

- Kernel Level Thread: The kernel knows about and manages the thread
- Kernel Mode Thread: The thread runs in *kernel mode*

### Why would you need kernel *mode* threads?

- Sometimes the OS needs to do some housekeeping that doesn't belong in a system call
- Free some pages, swap something in and out of memory
- Flush pages from the disk cache to the hard disk
- Perform VFS Checkpointing
- ...

## Stackframes

---

Let's look at some code

```
1  int foo(int a, int b) {  
2      int local = a + b;  
3      int anotherLocal = 'a';  
4      return local + anotherLocal;  
5  }
```


What is this? What does it contain for `int foo(int, int)`?

What is this? What does it contain for `int foo(int,int)?`

Return Address
<code>int local;</code>
<code>int anotherLocal;</code>

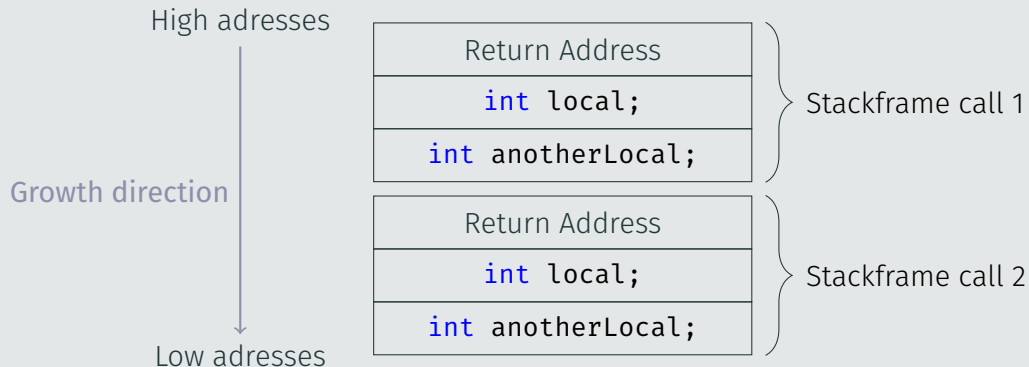
# Stackframes

What is this? What does it contain for `int foo(int,int)`?



# Stackframes

What is this? What does it contain for `int foo(int,int)?`



How are arguments passed?



How are arguments passed?

Registers!

How are arguments passed?

Registers! What happens for

```
1 void foo(int a, int b, int c, int d, int e, int f, int g, int h);
```

How are arguments passed?

Registers! What happens for

```
1 void foo(int a, int b, int c, int d, int e, int f, int g, int h);
```

We do not have enough registers!

Fixing...

# Stackframe

How are arguments passed?

Registers! What happens for

```
1 void foo(int a, int b, int c, int d, int e, int f, int g, int h);
```

We do not have enough registers!

Fixing...

int h
int g
Return Address
int local;
int anotherLocal;

## Popping frames

How could you release the space? I.e. figure out by how much to decrement the stackpointer?

## Popping frames

How could you release the space? I.e. figure out by how much to decrement the stackpointer?

The compiler surely knows how much space the local variables take up, right?

## Popping frames

How could you release the space? I.e. figure out by how much to decrement the stackpointer?

The compiler surely knows how much space the local variables take up, right?

No! C has e.g. *variable-length arrays* (VLA): `new int[someVariable]`.

⇒ Store the address at the start of the frame!

## Popping frames

How could you release the space? I.e. figure out by how much to decrement the stackpointer?

The compiler surely knows how much space the local variables take up, right?


No! C has e.g. *variable-length arrays* (VLA): `new int[someVariable]`.

⇒ Store the address at the start of the frame! Enter: `rbp`, the base pointer.



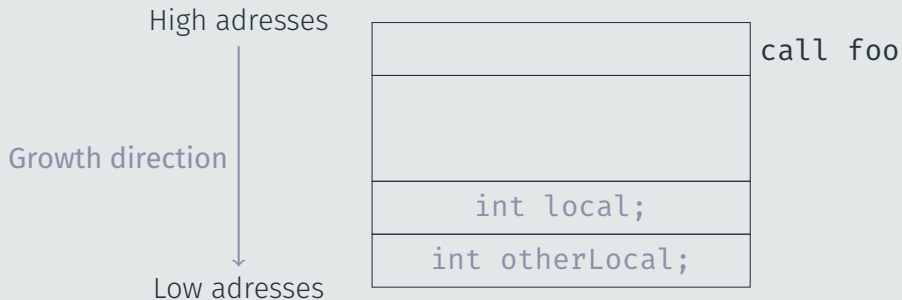
## Base pointer and the function (Pro|Epi)log

### Base pointer and the function (Pro|Epi)log



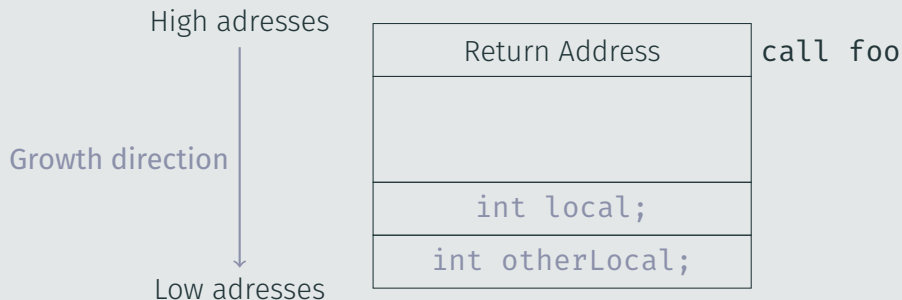
## Base pointer and the function (Pro|Epi)log

### Base pointer and the function (Pro|Epi)log



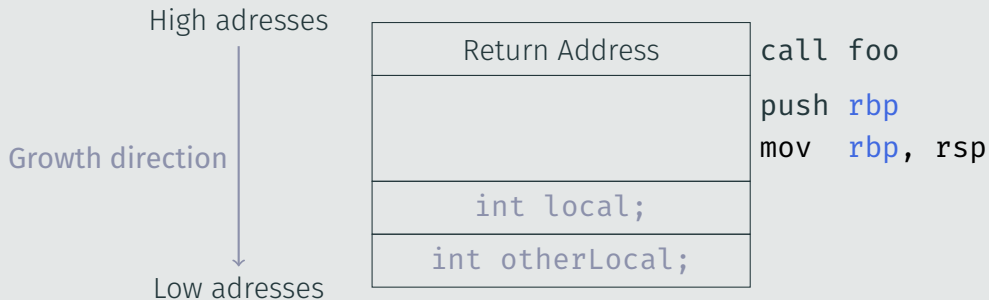
## Base pointer and the function (Pro|Epi)log

### Base pointer and the function (Pro|Epi)log



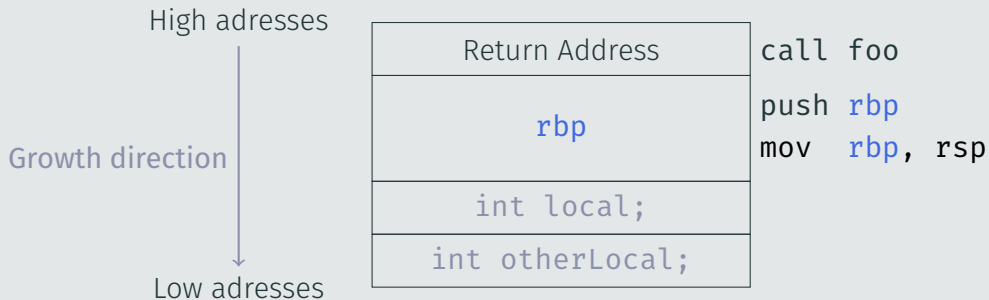
# Base pointer and the function (Pro|Epi)log

## Base pointer and the function (Pro|Epi)log



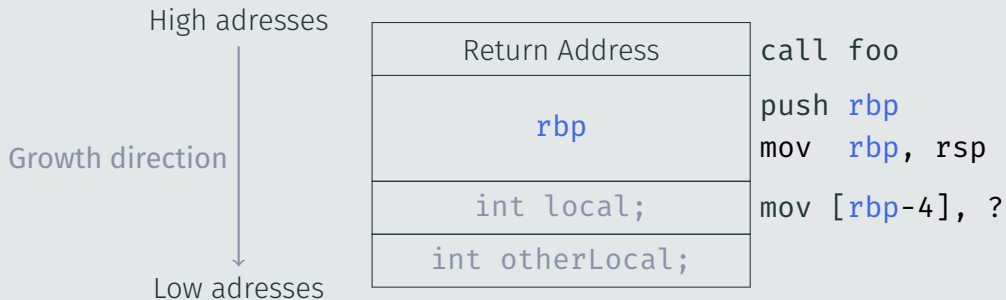
## Base pointer and the function (Pro|Epi)log

### Base pointer and the function (Pro|Epi)log



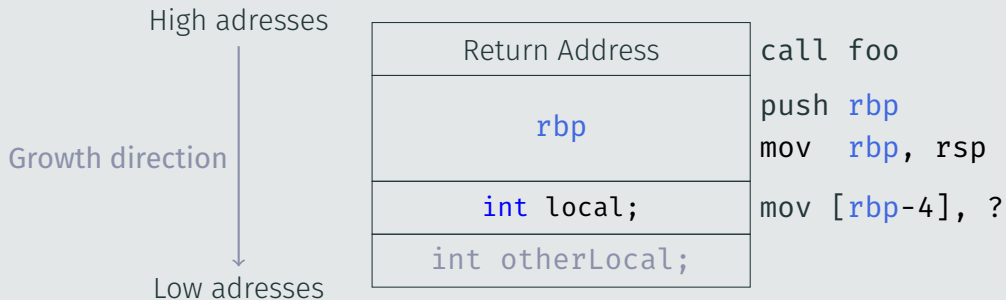
# Base pointer and the function (Pro|Epi)log

## Base pointer and the function (Pro|Epi)log



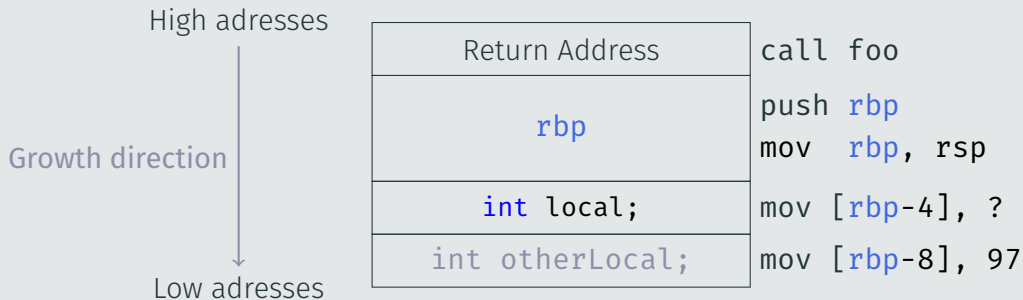
# Base pointer and the function (Pro|Epi)log

## Base pointer and the function (Pro|Epi)log



# Base pointer and the function (Pro|Epi)log

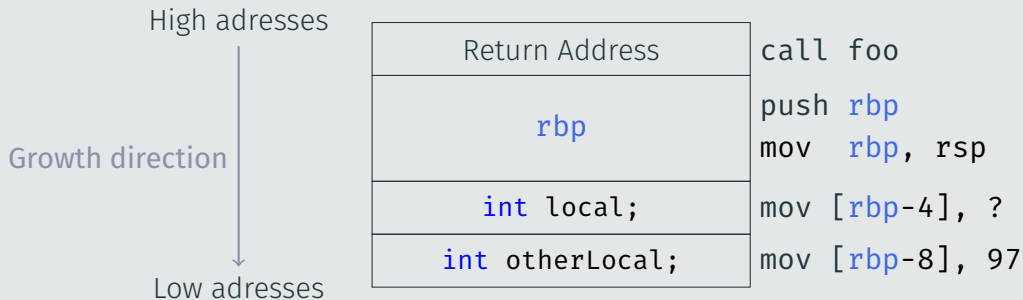
## Base pointer and the function (Pro|Epi)log





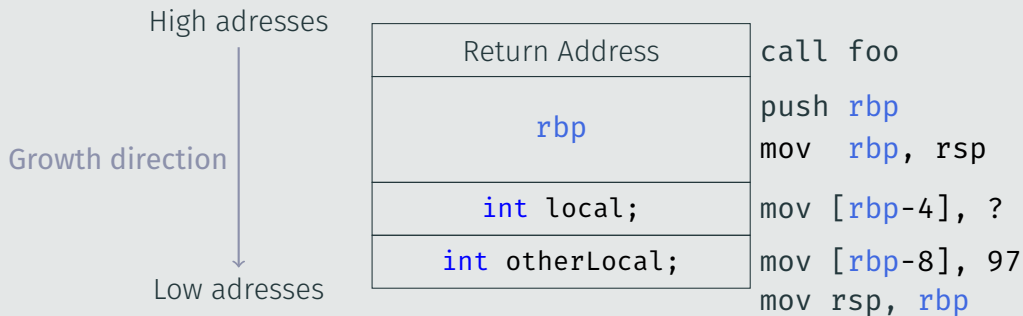
## Base pointer and the function (Pro|Epi)log

### Base pointer and the function (Pro|Epi)log



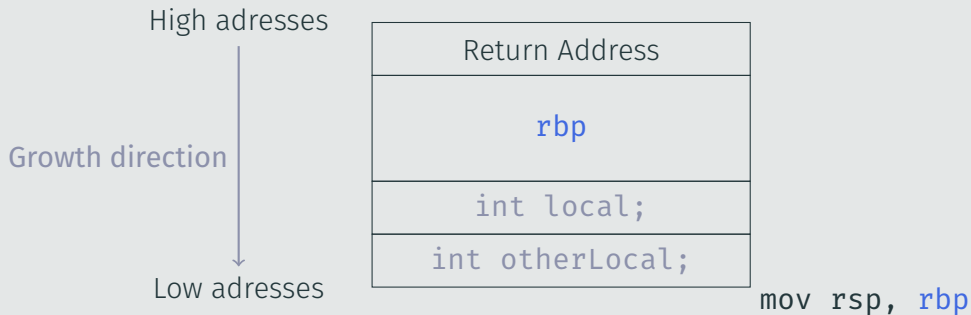
# Base pointer and the function (Pro|Epi)log

## Base pointer and the function (Pro|Epi)log



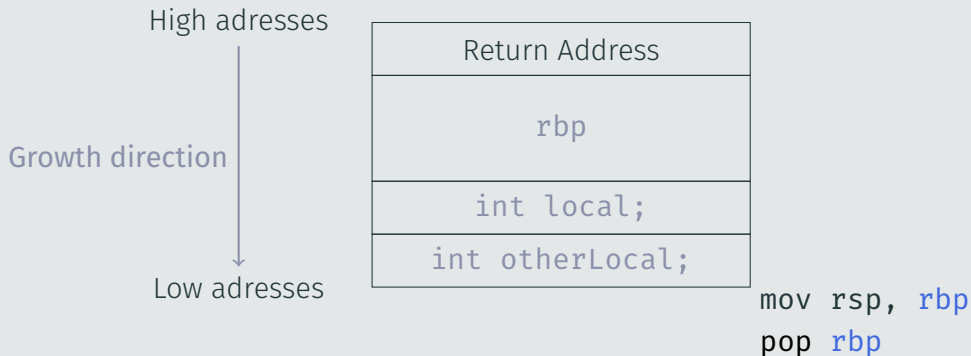
## Base pointer and the function (Pro|Epi)log

### Base pointer and the function (Pro|Epi)log



# Base pointer and the function (Pro|Epi)log

## Base pointer and the function (Pro|Epi)log



# Base pointer and the function (Pro|Epi)log

## Base pointer and the function (Pro|Epi)log

High addresses

Growth direction

Low addresses

Return Address
rbp
int local;
int otherLocal;

```
mov rsp, rbp
pop rbp
ret
```


# Base pointer and the function (Pro|Epi)log

## Base pointer and the function (Pro|Epi)log

High addresses

Growth direction

Low addresses



```
mov rsp, rbp
pop rbp
ret
```

# Segmentation

---

Where have you seen that word before while sadly staring at your screen?





Where have you seen that word before while sadly staring at your screen?

> Segmentation fault (core dumped)

## A few example segments

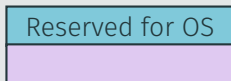
0xFFFFFFFF



0x00000000

## A few example segments


0xFFFFFFFF



0x00000000

## A few example segments


0xFFFFFFFF



0x00000000

## A few example segments


0xFFFFFFFF



0x00000000

## A few example segments

0xFFFFFFFF



0x00000000

## A few example segments

0xFFFFFFFF




0x00000000

# Segmentation

## A few example segments

0xFFFFFFFF




0x00000000



# Segmentation

## A few example segments

0xFFFFFFFF



0x00000000

# A Virtual Address


What does it look like?

42	0xCAFE
----	--------

Segment Number	Base	Limit
...	...	...
42	0xB105	0xF00D
...	...	...

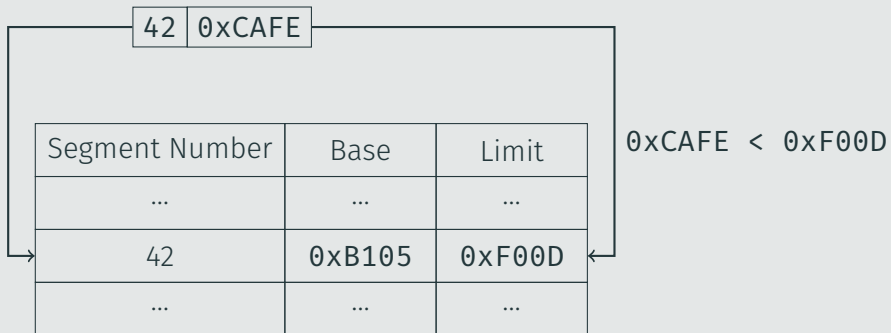
# A Virtual Address

What does it look like?



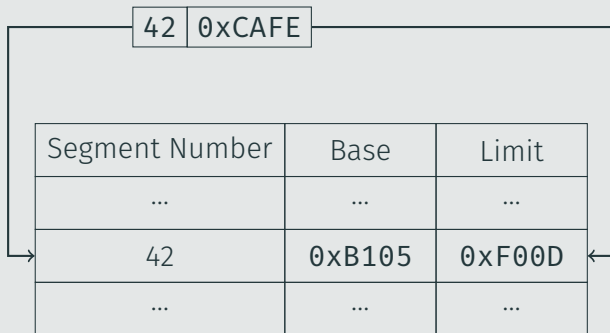
# A Virtual Address

What does it look like?



# A Virtual Address


What does it look like?



0xCAFE < 0xF00D

# A Virtual Address

What does it look like?



$$0xB105 + 0xCAFE = 0x17C03$$

## And let's try it

### Segments

Segment Number	Base	Limit
0	0xdead	0x00ef
1	0xf154	0x013a
2	0x0000	0x0000
3	0x0000	0x3fff

### Your task

Virtual Address	Segment Number	Offset	Valid?	Physical Address
0x2020	3	0x3999		
		0x0204	yes	
			yes	0xf15f

## And let's try it

### Solution

Virtual Address	Segment Number	Offset	Valid?	Physical Address
0xf999	3	0x3999	yes	0x3999
0x2020	0	0x2020	no	Offset outside limit
0xc204	3	0x0204	yes	0x0204
0x400b	1	0x000b	yes	0xf15f



# Memory Management Basics

---

And once again: What is the difference?

And once again: What is the difference?

- We assume no 1:1 mapping (i.e. we have virtual memory)

And once again: What is the difference?

- We assume no 1:1 mapping (i.e. we have virtual memory)
- *All program addresses are virtual*

## And once again: What is the difference?

- We assume no 1:1 mapping (i.e. we have virtual memory)
- *All program addresses are virtual*
- Mapped to *physical* addresses as needed by the memory management unit

What is *internal* fragmentation?


What is *internal* fragmentation?



Internal, i.e. *within* a block

# Fragmentation


What is *internal* fragmentation?



Internal, i.e. *within* a block



What is *internal* fragmentation?



Internal, i.e. *within* a block


What is *external* fragmentation?

What is *external* fragmentation?




External, i.e. due to *external factors* (different time-to-free)

What is *external* fragmentation?



External, i.e. due to *external factors* (different time-to-free)

What is *external* fragmentation?



External, i.e. due to *external factors* (different time-to-free)

Can you have *both* types at the same time?

Can you have *both* types at the same time?

Yes!

Can you have *both* types at the same time?

Yes!

- Allocate in *chunks* by e.g. rounding up to  $2^x$



Can you have *both* types at the same time?

Yes!

- Allocate in *chunks* by e.g. rounding up to  $2^x$
- Have different lifetimes

Can you have *both* types at the same time?


Yes!

- Allocate in *chunks* by e.g. rounding up to  $2^x$
- Have different lifetimes

⇒ Wasteful allocations scattered throughout RAM

# Fragmentation

What do we do now? This sounds bad!



What do we do now? This sounds bad!



What do we do now? This sounds bad!



# Fragmentation

What do we do now? This sounds bad!



This is called **Compaction**

# Fragmentation

What do we do now? This sounds bad!



This is called **Compaction**

Compaction - Is that even possible?

- C uses direct pointers

# Fragmentation

What do we do now? This sounds bad!



This is called **Compaction**

Compaction - Is that even possible?

- C uses direct pointers
- ⇒ They are all garbage now!



# Fragmentation

What do we do now? This sounds bad!



This is called **Compaction**

**Compaction - Is that even possible?**

- C uses direct pointers
- ⇒ They are all garbage now!
- Works just fine in languages with indirections (e.g. garbage collection)

# Fragmentation

What do we do now? This sounds bad!



This is called **Compaction**

**Compaction - Is that even possible?**

- C uses direct pointers
- ⇒ They are all garbage now!
- Works just fine in languages with indirections (e.g. garbage collection)
  - Also works for segments in physical memory! How?

# Fragmentation

What do we do now? This sounds bad!



This is called **Compaction**

**Compaction - Is that even possible?**

- C uses direct pointers
- ⇒ They are all garbage now!
- Works just fine in languages with indirections (e.g. garbage collection)
  - Also works for segments in physical memory! How? Update base addresses in MMU

## Memory allocation policies

---

Which strategies for finding free blocks do you know?

First Fit,

Which strategies for finding free blocks do you know?

First Fit, Best Fit,

## Some common strategies

Which strategies for finding free blocks do you know?

First Fit, Best Fit, Worst Fit

## Some common strategies

Which strategies for finding free blocks do you know?

First Fit, Best Fit, Worst Fit

**First Fit**

Pick the first block that is large enough



## Some common strategies

Which strategies for finding free blocks do you know?

First Fit, Best Fit, Worst Fit

**First Fit**

Pick the first block that is large enough

**Best Fit**

Pick the *smallest* block that fits

## Some common strategies

Which strategies for finding free blocks do you know?

First Fit, Best Fit, Worst Fit

**First Fit**

Pick the first block that is large enough

**Best Fit**

Pick the *smallest* block that fits


**Worst Fit**

Pick the *largest* block that fits

## Let's try them

### Best fit


Allocate 500, 1200, and 200, fail if not possible.



# Let's try them

## Best fit

Allocate 500, 1200, and 200, fail if not possible.





And compact it to fit the next one!


# Let's try them

## Best fit

Allocate 500, 1200, and 200, fail if not possible.



And compact it to fit the next one!





XKCD 138 - Pointers

F R A G E N?

Bis nächste Woche :)