

Betriebssysteme

7. Tutorium - Paging

Peter Bohner

13. Dezember 2024

ITEC - Operating Systems Group

Page Fault Handling

What is Demand-Paging and Pre-Paging?

Demand-Paging:

What is Demand-Paging and Pre-Paging?

Demand-Paging:

- Load pages on-demand, right when they are needed

What is Demand-Paging and Pre-Paging?

Demand-Paging:

- Load pages on-demand, right when they are needed

Pre-Paging:

- Loaded Pages speculatively in batches, even *before* you need them

What is Demand-Paging and Pre-Paging?

Demand-Paging:

- Load pages on-demand, right when they are needed

Pre-Paging:

- Loaded Pages speculatively in batches, even *before* you need them

Why would you (not?) use Demand-Paging?

What is Demand-Paging and Pre-Paging?

Demand-Paging:

- Load pages on-demand, right when they are needed

Pre-Paging:

- Loaded Pages speculatively in batches, even *before* you need them

Why would you (not?) use Demand-Paging?

- + Only loads needed data \Rightarrow Less memory wasted

What is Demand-Paging and Pre-Paging?

Demand-Paging:

- Load pages on-demand, right when they are needed

Pre-Paging:

- Loaded Pages speculatively in batches, even *before* you need them

Why would you (not?) use Demand-Paging?

- + Only loads needed data \Rightarrow Less memory wasted
- Generates lots of page faults before working set is in memory

Why would you (not?) use Pre-Paging?

Why would you (not?) use Pre-Paging?

- + Might reduce number of page faults

Why would you (not?) use Pre-Paging?

- + Might reduce number of page faults
- Loads more than needed \Rightarrow Wasteful

Why would you (not?) use Pre-Paging?

- + Might reduce number of page faults
- Loads more than needed \Rightarrow Wasteful
- More I/O \Rightarrow Slower?

Why would you (not?) use Pre-Paging?

- + Might reduce number of page faults
- Loads more than needed \Rightarrow Wasteful
- More I/O \Rightarrow Slower?
- + HDDs a lot faster when reading chunks

Different kind of page faults

Not all pages are created equal. Do you have any idea what types of page faults typically exist?

On-Demand Paging

Different kind of page faults

Not all pages are created equal. Do you have any idea what types of page faults typically exist?



Kernel Space

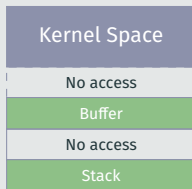
The diagram shows a light gray rectangular area representing memory space. Inside this area, there is a solid purple rectangle labeled 'Kernel Space'. Below the purple rectangle is a dashed-line rectangle labeled 'No access'.

No access

On-Demand Paging

Different kind of page faults

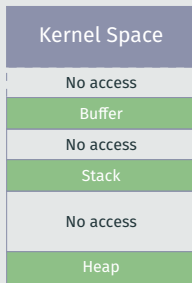
Not all pages are created equal. Do you have any idea what types of page faults typically exist?



On-Demand Paging

Different kind of page faults

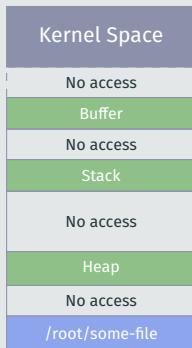
Not all pages are created equal. Do you have any idea what types of page faults typically exist?



On-Demand Paging

Different kind of page faults

Not all pages are created equal. Do you have any idea what types of page faults typically exist?



On-Demand Paging

Different kind of page faults

Not all pages are created equal. Do you have any idea what types of page faults typically exist?

Kernel Space
No access
Buffer
No access
Stack
No access
Heap
No access
/root/some-file
No access

On-Demand Paging

Different kind of page faults

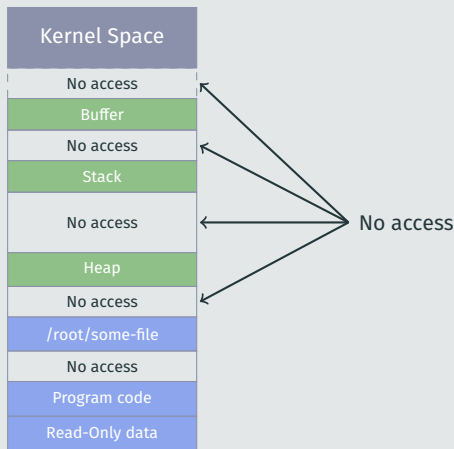
Not all pages are created equal. Do you have any idea what types of page faults typically exist?

Kernel Space
No access
Buffer
No access
Stack
No access
Heap
No access
/root/some-file
No access
Program code
Read-Only data

On-Demand Paging

Different kind of page faults

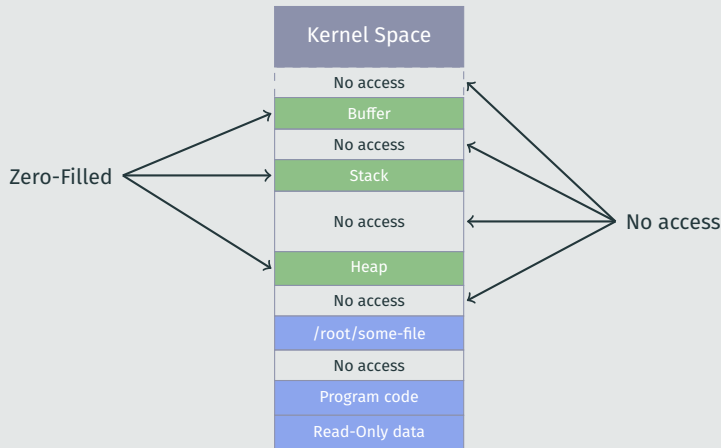
Not all pages are created equal. Do you have any idea what types of page faults typically exist?



On-Demand Paging

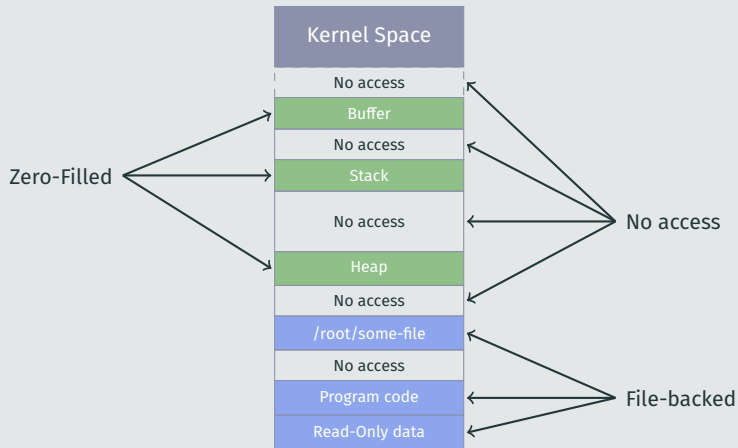
Different kind of page faults

Not all pages are created equal. Do you have any idea what types of page faults typically exist?



Different kind of page faults

Not all pages are created equal. Do you have any idea what types of page faults typically exist?



Different kind of page faults

Why are pages to generic memory zero filled?

Different kind of page faults

Why are pages to generic memory zero filled? It could leak other processes' memory otherwise (called a [Covert Channel](#)).

Different kind of page faults

Why are pages to generic memory zero filled? It could leak other processes' memory otherwise (called a [Covert Channel](#)).

And there is one other kind of page fault...

Different kind of page faults

Why are pages to generic memory zero filled? It could leak other processes' memory otherwise (called a [Covert Channel](#)).

And there is one other kind of page fault...

...The page was stolen by the OS and swapped out!

Different kind of page faults

Why are pages to generic memory zero filled? It could leak other processes' memory otherwise (called a [Covert Channel](#)).

And there is one other kind of page fault...

...The page was stolen by the OS and swapped out!

Also supported on some systems: *Purgable memory*. Stolen from [Apple](#) and also implemented [in SerenityOS in this video](#).

What kind of information does the page fault handler need?

What kind of information does the page fault handler need?

- Access flags: Can the user perform the operation on this page?

What kind of information does the page fault handler need?

- Access flags: Can the user perform the operation on this page?
- Where to find the most recent version (different for zero filled, file backed, etc.)

How could you implement Copy-on-Write memory?

How could you implement Copy-on-Write memory?

- Mark memory as read-only on fork
- Add an additional **CoW** flag: When a page fault is raised check it, copy the page and clear the **CoW** and **ro** flag

Page replacement

The pager in some systems tries to keep „spare pages“. Why?

The pager in some systems tries to keep „spare pages“. Why?

- OS needs free (pre-zeroed) frames to assign to processes

The pager in some systems tries to keep „spare pages“. Why?

- OS needs free (pre-zeroed) frames to assign to processes
- What happens when there are none and a page fault occurs?

The pager in some systems tries to keep „spare pages“. Why?

- OS needs free (pre-zeroed) frames to assign to processes
- What happens when there are none and a page fault occurs?

⇒ Needs to write dirty pages back to disk

⇒ Sloow

If you need to swap out a page, what pages do you search for a victim?

If you need to swap out a page, what pages do you search for a victim?

Local page replacement algorithms:

- Only look through the frames *of that process*

Global page replacement algorithms:

- Look through *all* frames, even that of other processes

What are the pro and cons of local algorithms?

What are the pro and cons of local algorithms?

- + Guaranteed number of pages per application

What are the pro and cons of local algorithms?

- + Guaranteed number of pages per application
- + Potentially faster

What are the pro and cons of local algorithms?

- + Guaranteed number of pages per application
- + Potentially faster
- Guaranteed number of pages per application \Rightarrow Can not adapt as easily to changing demands

What are the pro and cons of local algorithms?

- + Guaranteed number of pages per application
- + Potentially faster
- Guaranteed number of pages per application \Rightarrow Can not adapt as easily to changing demands
- Difficult selection of optimal number of frames per application (see point above)

What are the pro and cons of local algorithms?

- + Guaranteed number of pages per application
- + Potentially faster
- Guaranteed number of pages per application \Rightarrow Can not adapt as easily to changing demands
- Difficult selection of optimal number of frames per application (see point above)

Every process should get an equal number of pages. Do you use a Global or Local Algorithm?

What are the pro and cons of local algorithms?

- + Guaranteed number of pages per application
- + Potentially faster
- Guaranteed number of pages per application \Rightarrow Can not adapt as easily to changing demands
- Difficult selection of optimal number of frames per application (see point above)

Every process should get an equal number of pages. Do you use a Global or Local Algorithm?

Local

What is the working set?

- The set of pages that a process accessed in the last Δ page references

What is the working set?

- The set of pages that a process accessed in the last Δ page references

What is Thrashing?

Stack Page

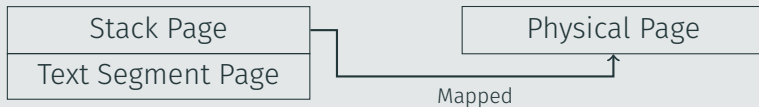
Text Segment Page

Physical Page

What is the working set?

- The set of pages that a process accessed in the last Δ page references

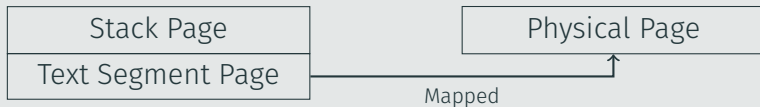
What is Thrashing?



What is the working set?

- The set of pages that a process accessed in the last Δ page references

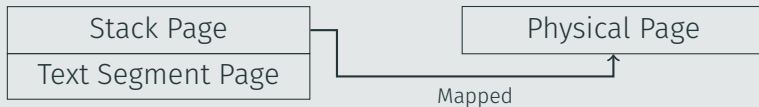
What is Thrashing?



What is the working set?

- The set of pages that a process accessed in the last Δ page references

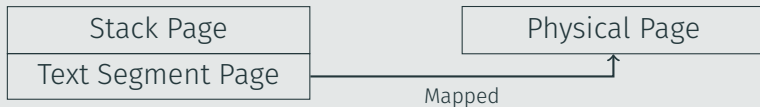
What is Thrashing?



What is the working set?

- The set of pages that a process accessed in the last Δ page references

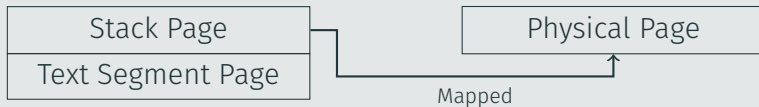
What is Thrashing?



What is the working set?

- The set of pages that a process accessed in the last Δ page references

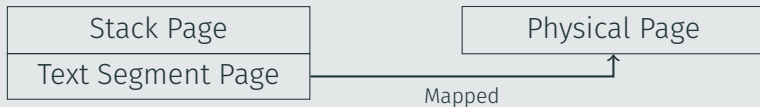
What is Thrashing?



What is the working set?

- The set of pages that a process accessed in the last Δ page references

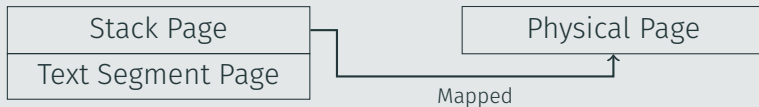
What is Thrashing?



What is the working set?

- The set of pages that a process accessed in the last Δ page references

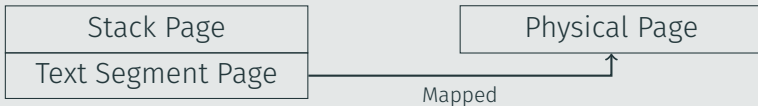
What is Thrashing?



What is the working set?

- The set of pages that a process accessed in the last Δ page references

What is Thrashing?



- Not enough frames to fit the working set
- ⇒ Pages will be stored to disk and reloaded very often

What are those?

We need to find a victim page to replace with a new one.

What are those?

We need to find a victim page to replace with a new one.

Common strategies

- FIFO:

What are those?

We need to find a victim page to replace with a new one.

Common strategies

- **FIFO:** First page to be loaded is the first to be unloaded
- **LRU:**

What are those?

We need to find a victim page to replace with a new one.

Common strategies

- **FIFO:** First page to be loaded is the first to be unloaded
- **LRU:** Least recently used page is evicted
- **Optimal:**

What are those?

We need to find a victim page to replace with a new one.

Common strategies

- **FIFO:** First page to be loaded is the first to be unloaded
- **LRU:** Least recently used page is evicted
- **Optimal:** Evict the page that is used the *furthest into the future*. Feasible?

What are those?

We need to find a victim page to replace with a new one.

Common strategies

- **FIFO:** First page to be loaded is the first to be unloaded
- **LRU:** Least recently used page is evicted
- **Optimal:** Evict the page that is used the *furthest into the future*. Feasible?
Used as a *benchmark*
- **Clock:**

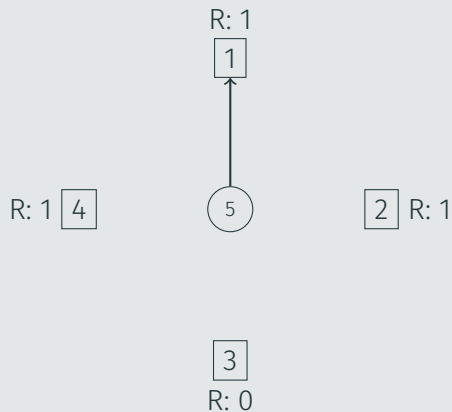
What are those?

We need to find a victim page to replace with a new one.

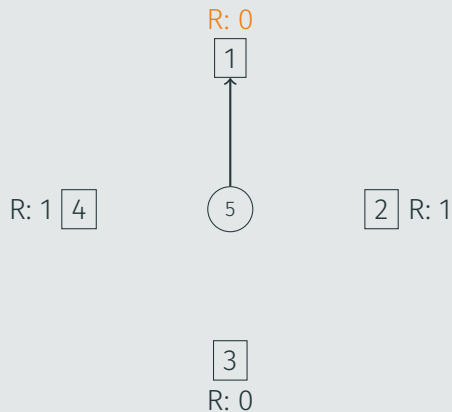
Common strategies

- **FIFO:** First page to be loaded is the first to be unloaded
- **LRU:** Least recently used page is evicted
- **Optimal:** Evict the page that is used the *furthest into the future*. Feasible?
Used as a *benchmark*
- **Clock:** Uff, let's talk about that on its own page

Clock

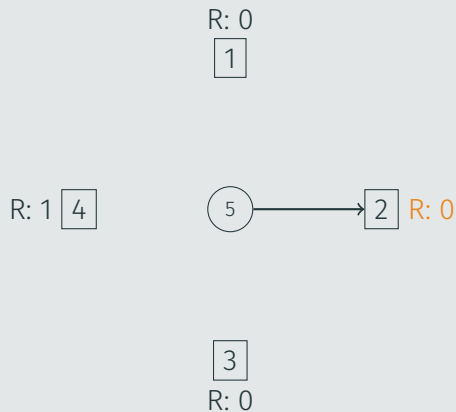


Clock

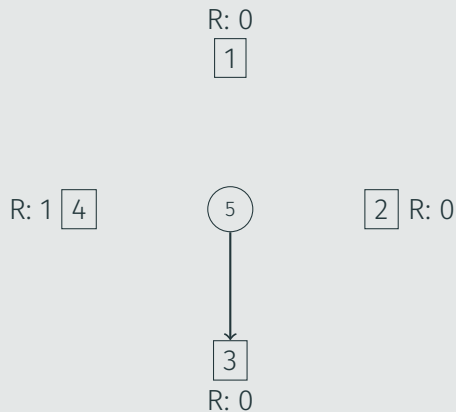




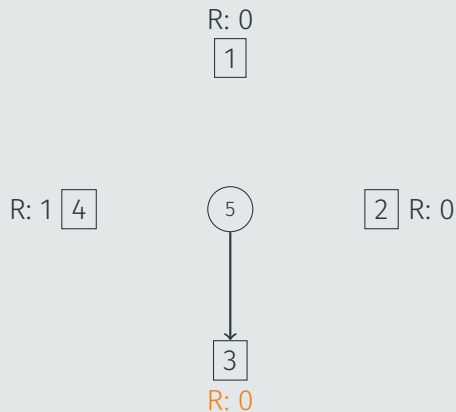
Clock



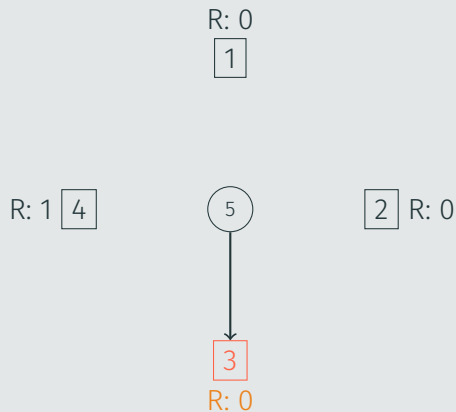
Clock



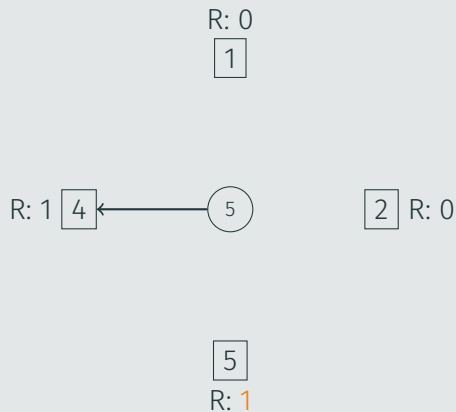
Clock



Clock



Clock



Where did that R come from?

Where did that R come from?

- Referenced (and modified) bits are set by the CPU when accessing or writing to the page

Where did that R come from?

- Referenced (and modified) bits are set by the CPU when accessing or writing to the page
- Referenced bits are periodically cleared by the kernel using timer interrupts

Do it

- Clock: Ordered by Load time ASC, Head is on Frame 3
- Reference order: 4, 0, 0, 0, 2, 4, 2, 1, 0, 3, 2

Frame	Virtual page	Load time	Access time	Referenced	Modified
0	2	60	161	0	1
1	1	130	160	0	0
2	0	26	162	1	0
3	3	20	163	1	1

How well does LRU work for the Stack, the Heap and the Code segment?

How well does LRU work for the Stack, the Heap and the Code segment?

- Stack: Beautifully. The older it is the longer it will take to be reached again

How well does LRU work for the Stack, the Heap and the Code segment?

- Stack: Beautifully. The older it is the longer it will take to be reached again
- Code: Mostly, loops are mostly linear and it follows certain patterns

How well does LRU work for the Stack, the Heap and the Code segment?

- Stack: Beautifully. The older it is the longer it will take to be reached again
- Code: Mostly, loops are mostly linear and it follows certain patterns
- Heap: Well, the heap has more random access patterns. So not that well, probably



XKCD 313 - Insomnia

F R A G E N?

Bis nächste Woche :)