

Betriebssysteme

12. Tutorium - Files and Directories

Peter Bohner

30. Januar 2025

ITEC - Operating Systems Group

Solid-State Drives

NAND based flash memory

Rejoice, TI might be useful once

How long do writes/reads take normally?

NAND based flash memory

Rejoice, TI might be useful once

How long do writes/reads take normally?

- Reading a page: 25 μ s

NAND based flash memory

Rejoice, TI might be useful once

How long do writes/reads take normally?

- Reading a page: 25 μ s
- Writing a page: 250 μ s

NAND based flash memory

Rejoice, TI might be useful once

How long do writes/reads take normally?

- Reading a page: 25 μ s
- Writing a page: 250 μ s
- Erasing a block:

NAND based flash memory

Rejoice, TI might be useful once

How long do writes/reads take normally?

- Reading a page: 25 μ s
- Writing a page: 250 μ s
- Erasing a block: 2ms

What happens when you just write to a random page?

NAND based flash memory

Rejoice, TI might be useful once

How long do writes/reads take normally?

- Reading a page: 25 μ s
- Writing a page: 250 μ s
- Erasing a block: 2ms

What happens when you just write to a random page?

Speeding things up

What could you change so writing pages is faster?

NAND based flash memory

Rejoice, TI might be useful once

How long do writes/reads take normally?

- Reading a page: 25 μ s
- Writing a page: 250 μ s
- Erasing a block: 2ms

What happens when you just write to a random page?

Speeding things up

What could you change so writing pages is faster?

- Keep around spare *erased* pages
- ⇒ You do not pay the erase penalty!
- When do you create / reserve / erase those spare pages?

NAND based flash memory

Rejoice, TI might be useful once

How long do writes/reads take normally?

- Reading a page: 25μs
- Writing a page: 250μs
- Erasing a block: 2ms

What happens when you just write to a random page?

Speeding things up

What could you change so writing pages is faster?

- Keep around spare *erased* pages

⇒ You do not pay the erase penalty!

- When do you create / reserve / erase those spare pages? Probably in the background. Any problems?

NAND based flash memory

Rejoice, TI might be useful once

How long do writes/reads take normally?

- Reading a page: 25 μ s
- Writing a page: 250 μ s
- Erasing a block: 2ms

What happens when you just write to a random page?

Speeding things up

What could you change so writing pages is faster?

- Keep around spare *erased* pages

⇒ You do not pay the erase penalty!

- When do you create / reserve / erase those spare pages? Probably in the background. Any problems? Might get exhausted if you write too much data in a short timeframe or the disk is full!

Deleting files

What happens when you delete a file? What effect does that have on the SSD performance?

Deleting files

What happens when you delete a file? What effect does that have on the SSD performance? The block is not freed \Rightarrow Can't be used as an erased empty page

What can the OS do to combat that?

Deleting files


What happens when you delete a file? What effect does that have on the SSD performance? The block is not freed \Rightarrow Can't be used as an erased empty page

What can the OS do to combat that?

The **trim** command

Can be issued by the OS to tell the SSD firmware what pages can be safely erased.

RAID



What is that?

A Redundant Array of Independent/Inexpensive Disks

What is that?

A Redundant Array of Independent/Inexpensive Disks

Why would you use that?

What is that?

A Redundant Array of Independent/Inexpensive Disks

Why would you use that?

- Probably cheaper than a SLED (Single Large Expensive Disk)

What is that?

A Redundant Array of Independent/Inexpensive Disks

Why would you use that?

- Probably cheaper than a SLED (Single Large Expensive Disk)
- Might be more resilient

What is that?


A Redundant Array of Independent/Inexpensive Disks

Why would you use that?


- Probably cheaper than a SLED (Single Large Expensive Disk)
- Might be more resilient
- Might be faster

Great, you now have multiple disks. How do you store your files on them?

- „I like to live dangerously“ - RAID Level 0
- Mirroring: RAID Level 1
- Historic variants: RAID Level 2 and 3
- Block striping and parity: RAID Level 4
- Block striping and *distributed* parity: RAID Level 5




Benefits / Drawbacks?




Benefits / Drawbacks?

- + Extremely fast (parallel reads and writes)




Benefits / Drawbacks?

- + Extremely fast (parallel reads and writes)
- + Can use full capacity




Benefits / Drawbacks?

- + Extremely fast (parallel reads and writes)
- + Can use full capacity
- If a single disk fails your files are toast




Benefits / Drawbacks?




Benefits / Drawbacks?

- + You can lose all but one disk without losing data




Benefits / Drawbacks?

- + You can lose all but one disk without losing data
- + Parallel reads possible




Benefits / Drawbacks?

- + You can lose all but one disk without losing data
- + Parallel reads possible
- Writes slower as they need to write to all disks




Benefits / Drawbacks?

- + You can lose all but one disk without losing data
- + Parallel reads possible
- Writes slower as they need to write to all disks
- Size equals the size of a single disk




What is that?

- Have $\log_2(N)$ parity disk




What is that?

- Have $\log_2(N)$ parity disk
- Stripe data at the *bit* level




What is that?

- Have $\log_2(N)$ parity disk
- Stripe data at the *bit* level
- Use a hamming code of proper size




What is that?

- Have $\log_2(N)$ parity disk
- Stripe data at the *bit* level
- Use a hamming code of proper size
- Spin the disks in lockstep (so you read all bits of your word at once)




Benefits / Drawbacks




Benefits / Drawbacks

- + External error checking




Benefits / Drawbacks

- + External error checking
- Really slow




Benefits / Drawbacks

- + External error checking
- Really slow
- Not that useful as disks have internal error checking by now




Benefits / Drawbacks

- + External error checking
- Really slow
- Not that useful as disks have internal error checking by now
- Spins in lockstep \Rightarrow Can only service one request at a time




What is that?

- Have a dedicated parity disk




What is that?

- Have a dedicated parity disk
- Stripe data at the *byte* level




What is that?

- Have a dedicated parity disk
- Stripe data at the *byte* level
- Spin the disks in lockstep (so you read all bytes of your word at once)




Benefits / Drawbacks




Benefits / Drawbacks

- + You can lose a disk and restore it using the parity




Benefits / Drawbacks

- + You can lose a disk and restore it using the parity
- Slow when reading/writing small files at random locations




Benefits / Drawbacks

- + You can lose a disk and restore it using the parity
- Slow when reading/writing small files at random locations
- Spins in lockstep \Rightarrow Can only service one request at a time




Benefits / Drawbacks

- + You can lose a disk and restore it using the parity
- Slow when reading/writing small files at random locations
- Spins in lockstep \Rightarrow Can only service one request at a time
- Every write and read hits the same single parity disk




What is that?

- Have a dedicated parity disk




What is that?

- Have a dedicated parity disk
- Stripe data at the *block* level




Benefits / Drawbacks




Benefits / Drawbacks

- + You can lose a disk and restore it using the parity




Benefits / Drawbacks

- + You can lose a disk and restore it using the parity
- + Good read performance




Benefits / Drawbacks

- + You can lose a disk and restore it using the parity
- + Good read performance
- Every write and read hits the same single parity disk \Rightarrow Bottleneck, prone to failure




Benefits / Drawbacks

- + You can lose a disk and restore it using the parity
- + Good read performance
- Every write and read hits the same single parity disk \Rightarrow Bottleneck, prone to failure
- Slow writes (write to same parity disk)




What is that?

- Stripe data at the *block* level




What is that?

- Stripe data at the *block* level
- Distribute parity across your disks




Benefits / Drawbacks




Benefits / Drawbacks

- + You can lose a disk and restore it using the parity




Benefits / Drawbacks

- + You can lose a disk and restore it using the parity
- + Good read performance



Benefits / Drawbacks

- + You can lose a disk and restore it using the parity
- + Good read performance
- + Okay write performance



Benefits / Drawbacks

- + You can lose a disk and restore it using the parity
- + Good read performance
- + Okay write performance
- Still slower than RAID 0 or a SLED

Compare SLED and RAID (Level 0, 1, 4, 5)

Each RAID uses 4 disks for actual data storage.

How many disks do you need?

- SLED: 1

Compare SLED and RAID (Level 0, 1, 4, 5)

Each RAID uses 4 disks for actual data storage.

How many disks do you need?

- SLED: 1
- RAID 0: 4

Compare SLED and RAID (Level 0, 1, 4, 5)

Each RAID uses 4 disks for actual data storage.

How many disks do you need?

- SLED: 1
- RAID 0: 4
- RAID 1: 8

Compare SLED and RAID (Level 0, 1, 4, 5)

Each RAID uses 4 disks for actual data storage.

How many disks do you need?

- SLED: 1
- RAID 0: 4
- RAID 1: 8
- RAID 4: 5

Compare SLED and RAID (Level 0, 1, 4, 5)

Each RAID uses 4 disks for actual data storage.

How many disks do you need?

- SLED: 1
- RAID 0: 4
- RAID 1: 8
- RAID 4: 5
- RAID 5: 5

You want to modify one byte of data. How many blocks do you need to read/write?

- SLED: 1 read + 1 write

You want to modify one byte of data. How many blocks do you need to read/write?

- SLED: 1 read + 1 write
- RAID 0: 1 read + 1 write

You want to modify one byte of data. How many blocks do you need to read/write?

- SLED: 1 read + 1 write
- RAID 0: 1 read + 1 write
- RAID 1: 1 read + 2 write (1 data + 1 mirror)

You want to modify one byte of data. How many blocks do you need to read/write?

- SLED: 1 read + 1 write
- RAID 0: 1 read + 1 write
- RAID 1: 1 read + 2 write (1 data + 1 mirror)
- RAID 4: 2 read (data + old parity) + 2 write (data + new parity)

You want to modify one byte of data. How many blocks do you need to read/write?

- SLED: 1 read + 1 write
- RAID 0: 1 read + 1 write
- RAID 1: 1 read + 2 write (1 data + 1 mirror)
- RAID 4: 2 read (data + old parity) + 2 write (data + new parity)
- RAID 5: 2 read (data + old parity) + 2 write (data + new parity)

You are using RAID

- You accidentally delete a file.

You are using RAID

- You accidentally delete a file. *GONE*
- You accidentally overwrite a file.

You are using RAID

- You accidentally delete a file. ***GONE***
- You accidentally overwrite a file. ***GONE***
- Some data gets corrupted on one disk.

You are using RAID

- You accidentally delete a file. **GONE**
- You accidentally overwrite a file. **GONE**
- Some data gets corrupted on one disk. **GONE** (*probably*)
- [The poor intern connects to the production database.](#) (Or [here](#))

You are using RAID

- You accidentally delete a file. **GONE**
- You accidentally overwrite a file. **GONE**
- Some data gets corrupted on one disk. **GONE** (*probably*)
- [The poor intern connects to the production database.](#) (Or [here](#)) **GONE**
- A crypto-locker takes out your computer.

You are using RAID

- You accidentally delete a file. **GONE**
- You accidentally overwrite a file. **GONE**
- Some data gets corrupted on one disk. **GONE** (*probably*)
- [The poor intern connects to the production database.](#) (Or [here](#)) **GONE**
- A crypto-locker takes out your computer. Believe it or not, ~~JAIL~~ **GONE**

You are using RAID

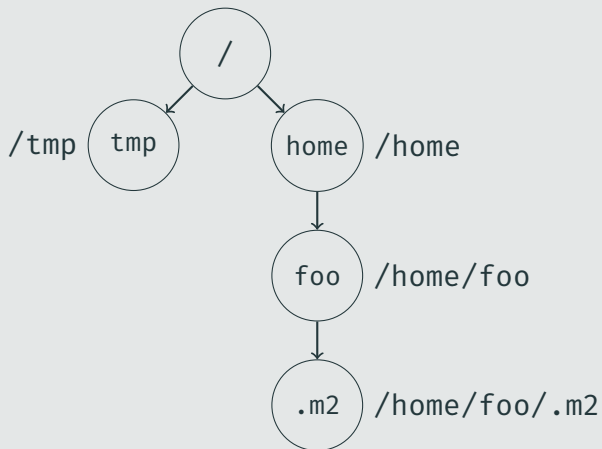
- You accidentally delete a file. **GONE**
- You accidentally overwrite a file. **GONE**
- Some data gets corrupted on one disk. **GONE** (*probably*)
- [The poor intern connects to the production database.](#) (Or [here](#)) **GONE**
- A crypto-locker takes out your computer. Believe it or not, ~~JAIL~~ **GONE**

So what do we learn?

RAID IS NO SUBSTITUTE FOR A BACKUP

Files And Directories

Paths (Linux)



What are the two basic access methods (patterns) for reading a file?

- **Sequential Access**

Accessed in order, reading sequential bytes. Writes append at the end.

- **Random Access**

Reading and writing at arbitrary positions, programmer needs to specify where to write/read

Which Library Functions and System Calls do we have to access files?

Opening the file

Which Library Functions and System Calls do we have to access files?

Opening the file

`fopen` / `open`. Provide mode as argument (read/write, where, create, etc.)

Which Library Functions and System Calls do we have to access files?

Opening the file

`fopen` / `open`. Provide mode as argument (read/write, where, create, etc.)

Closing a file

Which Library Functions and System Calls do we have to access files?

Opening the file

`fopen` / `open`. Provide mode as argument (read/write, where, create, etc.)

Closing a file

`fclose` / `close`

Which Library Functions and System Calls do we have to access files?

Opening the file

`fopen` / `open`. Provide mode as argument (read/write, where, create, etc.)

Closing a file

`fclose` / `close`

Reading from a file

Which Library Functions and System Calls do we have to access files?

Opening the file

`fopen` / `open`. Provide mode as argument (read/write, where, create, etc.)

Closing a file

`fclose` / `close`

Reading from a file

`fread` / `read`

Which Library Functions and System Calls do we have to access files?

Opening the file

`fopen` / `open`. Provide mode as argument (read/write, where, create, etc.)

Closing a file

`fclose` / `close`

Reading from a file

`fread` / `read`

Writing to a file

Which Library Functions and System Calls do we have to access files?

Opening the file

`fopen` / `open`. Provide mode as argument (read/write, where, create, etc.)

Closing a file

`fclose` / `close`

Reading from a file

`fread` / `read`

Writing to a file

`fwrite` / `write`

Which Library Functions and System Calls do we have to access files?

Opening the file

`fopen` / `open`. Provide mode as argument (read/write, where, create, etc.)

Closing a file

`fclose` / `close`

Reading from a file

`fread` / `read`

Writing to a file

`fwrite` / `write`

Flushing dirty buffers

Which Library Functions and System Calls do we have to access files?

Opening the file

`fopen` / `open`. Provide mode as argument (read/write, where, create, etc.)

Closing a file

`fclose` / `close`

Reading from a file

`fread` / `read`

Writing to a file

`fwrite` / `write`

Flushing dirty buffers

The library functions sometimes buffer to reduce the amount of syscalls.

`fflush` flushes those buffers.

Which system calls / library functions move the „current position pointer“?

Which system calls / library functions move the „current position pointer“?

- Linux: `fseek` / `lseek`
- Windows: `SetFilePointerEx` / `SetFilePointer`

Which system calls / library functions move the „current position pointer“?

- Linux: `fseek` / `lseek`
- Windows: `SetFilePointerEx` / `SetFilePointer`

What happens when you move the cursor behind the end of the file?

Which system calls / library functions move the „current position pointer“?

- Linux: `fseek` / `lseek`
- Windows: `SetFilePointerEx` / `SetFilePointer`

What happens when you move the cursor behind the end of the file?

It creates a hole filled with zeros!

Which system calls / library functions move the „current position pointer“?

- Linux: `fseek` / `lseek`
- Windows: `SetFilePointerEx` / `SetFilePointer`

What happens when you move the cursor behind the end of the file?

It creates a hole filled with zeros! Such a file is called a *sparse* file and some file systems might not store empty regions.

How could you implement random access without a dedicated file pointer?

Add an offset to the `read` and `write` system call.

How could you implement random access without a dedicated file pointer?

Add an offset to the `read` and `write` system call.

- + Save a system call

How could you implement random access without a dedicated file pointer?

Add an offset to the `read` and `write` system call.

- + Save a system call
- Caller need to keep track of the current offset for sequential reads

How could you implement random access without a dedicated file pointer?

Add an offset to the `read` and `write` system call.

- + Save a system call
- Caller need to keep track of the current offset for sequential reads

`mmap` the file (works quite well, might cause page faults. Probably faster for large files than read/write calls)

What system calls do you need to list files in a Linux directory?

What system calls do you need to list files in a Linux directory?

1. `opendir`

What system calls do you need to list files in a Linux directory?

1. `opendir`
2. `readdir`:

What system calls do you need to list files in a Linux directory?

1. `opendir`
2. `readdir`: Returns a `dirent` with a *relative path*

What system calls do you need to list files in a Linux directory?

1. `opendir`
2. `readdir`: Returns a `dirent` with a *relative path*
3. `closedir`

Open Files

What structures does the kernel use for open files?

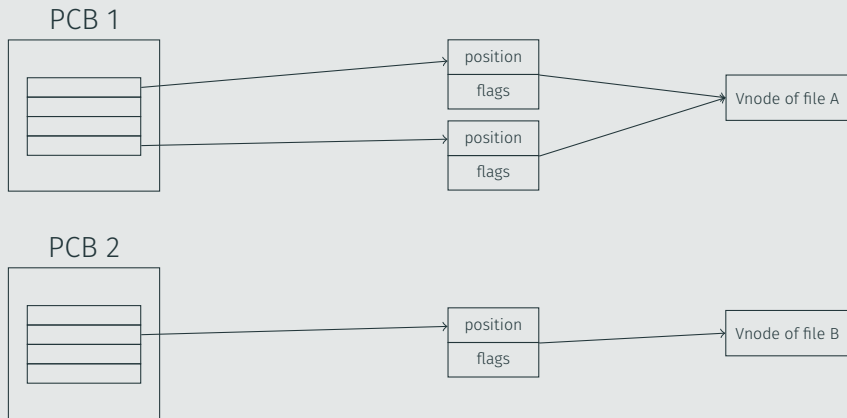
Kernel Data Structures For Open Files

What structures does the kernel use for open files?

PCBs with local open file tables

Global open file table

vnode table





XKCD 1084 - Server Problem

FRAGEN?



<https://forms.gle/9CwJSKidKibubran9>

Bis nächste Woche