

Betriebssysteme

13. Tutorium - Files Systems

Peter Bohner

6. Februar 2025


ITEC - Operating Systems Group

Disk Space Allocation

What is that?


Contiguous Allocation

What is that?




Contiguous Allocation

What is that?




Contiguous Allocation

What is that?



Contiguous Allocation

What is that?



Challenges

- Sizing your block (internal fragmentation, growth)
- External fragmentation

What's that?

Chained Allocation

What's that?


File A

File B




Chained Allocation

What's that?




Chained Allocation

What's that?




Chained Allocation

What's that?




Chained Allocation

What's that?



Chained Allocation


What's that?



Benefits? Drawbacks?

Chained Allocation

What's that?




Benefits? Drawbacks?

- + No longer need a contiguous chunk

Chained Allocation

What's that?




Benefits? Drawbacks?

- + No longer need a contiguous chunk
- Only sequential access

Chained Allocation

What's that?




Benefits? Drawbacks?

- + No longer need a contiguous chunk
- Only sequential access
- A single corrupted pointer is very bad news


Linked List Allocation With FAT

What is that?




Linked List Allocation With FAT

What is that?



Linked List Allocation With FAT


What is that?



0		
1	7	⇐ File A
2		
3		
4		
5		
6		
7		
8		

Linked List Allocation With FAT


What is that?



0		
1	7	⇐ File A
2		
3		
4		
5		
6		
7	3	
8		

Linked List Allocation With FAT


What is that?



0		
1	7	⇐ File A
2		
3	5	
4		
5		
6		
7	3	
8		

Linked List Allocation With FAT


What is that?



0		
1	7	⇐ File A
2		
3	5	
4		
5	2	
6		
7	3	
8		

Linked List Allocation With FAT


What is that?



0		
1	7	⇐ File A
2	-1	
3	5	
4		
5	2	
6		
7	3	
8		

Linked List Allocation With FAT


What is that?



0		
1	7	⇐ File A
2	-1	
3	5	
4		
5	2	
6		
7	3	
8		

Linked List Allocation With FAT


What is that?



0		
1	7	⇐ File A
2	-1	
3	5	
4		
5	2	
6	8	⇐ File B
7	3	
8		

Linked List Allocation With FAT


What is that?



0		
1	7	⇐ File A
2	-1	
3	5	
4		
5	2	
6	8	⇐ File B
7	3	
8	4	

Linked List Allocation With FAT

What is that?



0		
1	7	⇐ File A
2	-1	
3	5	
4	-1	
5	2	
6	8	⇐ File B
7	3	
8	4	

Benefits? Drawbacks?

Benefits? Drawbacks?

- + Hopefully fits in RAM \Rightarrow Iteration fast

Benefits? Drawbacks?

- + Hopefully fits in RAM \Rightarrow Iteration fast
- + Even if it doesn't fit: Less disk seeks

Benefits? Drawbacks?

- + Hopefully fits in RAM \Rightarrow Iteration fast
- + Even if it doesn't fit: Less disk seeks
- Size depends on size of hard disk (one entry per Block)

Benefits? Drawbacks?


- + Hopefully fits in RAM \Rightarrow Iteration fast
- + Even if it doesn't fit: Less disk seeks
- Size depends on size of hard disk (one entry per Block)
- Might be too large to be cached in RAM (on large disk)

What is that?

0	1	2
3	4	5
6	7	8



Indexed Allocation

What is that?




Indexed Allocation

What is that?



Indexed Allocation

What is that?





0	7
1	3
2	5
3	2

} Index Block 1

Indexed Allocation

What is that?



Benefits? Drawbacks?

Benefits? Drawbacks?

- + Hopefully fits in RAM \Rightarrow Lookups fast

Benefits? Drawbacks?

- + Hopefully fits in RAM \Rightarrow Lookups fast
- + Random access possible (just fetch the i th block)

Benefits? Drawbacks?

- + Hopefully fits in RAM \Rightarrow Lookups fast
- + Random access possible (just fetch the i th block)
- File size limited by how many references fit into one block

Benefits? Drawbacks?

- + Hopefully fits in RAM \Rightarrow Lookups fast
- + Random access possible (just fetch the *i*th block)
- File size limited by how many references fit into one block
- Wasteful for small files

Benefits? Drawbacks?

- + Hopefully fits in RAM \Rightarrow Lookups fast
- + Random access possible (just fetch the i th block)
- File size limited by how many references fit into one block
- Wasteful for small files

How could you store *Huge* files?

Benefits? Drawbacks?

- + Hopefully fits in RAM \Rightarrow Lookups fast
- + Random access possible (just fetch the *i*th block)
- File size limited by how many references fit into one block
- Wasteful for small files

How could you store *Huge* files?

- Increase the Block size

Benefits? Drawbacks?

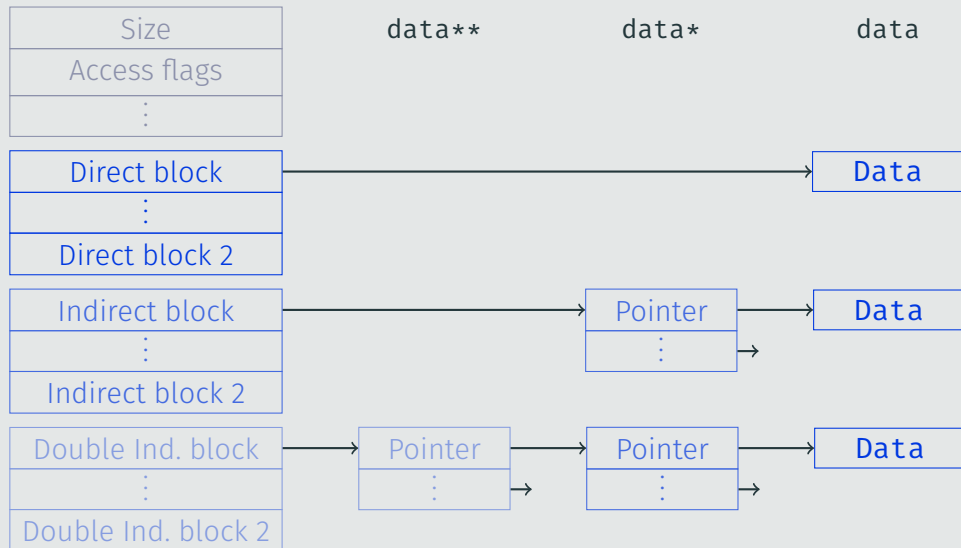
- + Hopefully fits in RAM \Rightarrow Lookups fast
- + Random access possible (just fetch the i th block)
- File size limited by how many references fit into one block
- Wasteful for small files

How could you store *Huge* files?

- Increase the Block size
- + Indirection! Make the index block link to other index blocks (like pagetables)
You can also mix that: First N pointers point to data blocks, next to indirect blocks, next to double-indirect blocks, etc.

Indexed Allocations With Inodes

What does an inode look like?



Indexed Allocations With Inodes

What is the maximum file size?

Assume that disk blocks are 8 KiB in size and a pointer to a disk block is 4 bytes long. An inode contains 12 pointers to direct blocks, and one pointer to a single, double, and triple indirect block, respectively.

Indexed Allocations With Inodes

What is the maximum file size?

Assume that disk blocks are 8 KiB in size and a pointer to a disk block is 4 bytes long. An inode contains 12 pointers to direct blocks, and one pointer to a single, double, and triple indirect block, respectively.

Calculating the amount of pointers to blocks

Pointers per Block: $8 \text{ KiB} / 4 \text{ bytes} = 2^{13} / 2^2 = 2^{11} = 2048$

Direct pointers

Indexed Allocations With Inodes

What is the maximum file size?

Assume that disk blocks are 8 KiB in size and a pointer to a disk block is 4 bytes long. An inode contains 12 pointers to direct blocks, and one pointer to a single, double, and triple indirect block, respectively.

Calculating the amount of pointers to blocks

Pointers per Block: $8 \text{ KiB} / 4 \text{ bytes} = 2^{13} / 2^2 = 2^{11} = 2048$

Direct pointers 12 pointers to blocks

One single-indirect block

Indexed Allocations With Inodes

What is the maximum file size?

Assume that disk blocks are 8 KiB in size and a pointer to a disk block is 4 bytes long. An inode contains 12 pointers to direct blocks, and one pointer to a single, double, and triple indirect block, respectively.

Calculating the amount of pointers to blocks

Pointers per Block: $8 \text{ KiB} / 4 \text{ bytes} = 2^{13} / 2^2 = 2^{11} = 2048$

Direct pointers 12 pointers to blocks

One single-indirect block 2048 pointers to blocks

One double-indirect block

Indexed Allocations With Inodes

What is the maximum file size?

Assume that disk blocks are 8 KiB in size and a pointer to a disk block is 4 bytes long. An inode contains 12 pointers to direct blocks, and one pointer to a single, double, and triple indirect block, respectively.

Calculating the amount of pointers to blocks

Pointers per Block: $8 \text{ KiB} / 4 \text{ bytes} = 2^{13} / 2^2 = 2^{11} = 2048$

Direct pointers	12 pointers to blocks
-----------------	-----------------------

One single-indirect block	2048 pointers to blocks
---------------------------	-------------------------

One double-indirect block	$2048 \cdot 2048$ pointers to blocks
---------------------------	--------------------------------------

One triple-indirect block	
---------------------------	--

Indexed Allocations With Inodes

What is the maximum file size?

Assume that disk blocks are 8 KiB in size and a pointer to a disk block is 4 bytes long. An inode contains 12 pointers to direct blocks, and one pointer to a single, double, and triple indirect block, respectively.

Calculating the amount of pointers to blocks

Pointers per Block: $8 \text{ KiB} / 4 \text{ bytes} = 2^{13} / 2^2 = 2^{11} = 2048$

Direct pointers	12 pointers to blocks
One single-indirect block	2048 pointers to blocks
One double-indirect block	$2048 \cdot 2048$ pointers to blocks
One triple-indirect block	$2048 \cdot 2048 \cdot 2048$ pointers to blocks

Indexed Allocations With Inodes

What is the maximum file size?

Assume that disk blocks are 8 KiB in size and a pointer to a disk block is 4 bytes long. An inode contains 12 pointers to direct blocks, and one pointer to a single, double, and triple indirect block, respectively.

Calculating the amount of pointers to blocks

Pointers per Block: $8 \text{ KiB} / 4 \text{ bytes} = 2^{13} / 2^2 = 2^{11} = 2048$

Direct pointers	12 pointers to blocks
One single-indirect block	2048 pointers to blocks
One double-indirect block	$2048 \cdot 2048$ pointers to blocks
One triple-indirect block	$2048 \cdot 2048 \cdot 2048$ pointers to blocks

Final result

$$(12 + 2048 + 2048^2 + 2048^3) \cdot 8\text{KiB} \approx 64\text{TiB}$$

File System Implementation

Hard links

Hard links

Pointer **to the same inode**.

⇒ *Everything* is identical: Size, content, access time, mode, ...

Can therefore *not* cross file system boundaries.

Hard links

Pointer **to the same inode**.

⇒ *Everything* is identical: Size, content, access time, mode, ...

Can therefore *not* cross file system boundaries.

Symlinks

File System Implementation

Hard links

Pointer **to the same inode**.

⇒ *Everything* is identical: Size, content, access time, mode, ...

Can therefore *not* cross file system boundaries.

Symlinks

Are (mostly) **normal files** containing **a filepath** with their own access time, mode, When programs access the file the OS transparently:

1. Reads the contents of the symlink file
2. Resolves the file path it read
3. Performs the operation on that path instead

Can cross file system boundaries or point to non-existent files or change what they point to if the file is moved, ...

Hard and symlink renames

```
1 echo "Hello" > test.txt
2 ln test.txt hardlink.txt
3 ln -s test.txt symlink.txt
4
5 mv test.txt renamed_test.txt
6 // Is the hardlink / symlink broken?
```

Having Fun With Paths I

Hard and symlink renames

```
1 echo "Hello" > test.txt
2 ln test.txt hardlink.txt
3 ln -s test.txt symlink.txt
4
5 mv test.txt renamed_test.txt
6 // Is the hardlink / symlink broken?
```

Reaching for the stars

Does this work (on my machine ;)?

```
1 cd /tmp/
2 echo "Hello" > test.txt
3 ln test.txt ~/test_link.txt
```

Having Fun With Paths I

Hard and symlink renames

```
1 echo "Hello" > test.txt
2 ln test.txt hardlink.txt
3 ln -s test.txt symlink.txt
4
5 mv test.txt renamed_test.txt
6 // Is the hardlink / symlink broken?
```

Reaching for the stars

Does this work (on my machine ;)?

```
1 cd /tmp/
2 echo "Hello" > test.txt
3 ln test.txt ~/test_link.txt
```

No, as **/tmp** is mounted as tmpfs and / as ext4.

Hard and symlink renames

```
1 echo "Hello" > test.txt
2 ln test.txt hardlink.txt
3 ln -s test.txt symlink.txt
4
5 cp test.txt renamed_test.txt
6 rm test.txt
7 // Is the hardlink / symlink broken?
8
```

Hard and symlink renames

```
1  echo "Hello" > test.txt
2  ln test.txt hardlink.txt
3  ln -s test.txt symlink.txt
4
5  cp test.txt renamed_test.txt
6  rm test.txt
7  // Is the hardlink / symlink broken?
8
9  echo "Hello" > test.txt
10 // Is the hardlink / symlink still broken?
```


Hard and symlink renames

```
1  echo "Hello" > test.txt
2  ln test.txt hardlink.txt
3  ln -s test.txt symlink.txt
4
5  cp test.txt renamed_test.txt
6  rm test.txt
7  // Is the hardlink / symlink broken?
8
9  echo "Hello" > test.txt
10 // Is the hardlink / symlink still broken?
```

If you delete a hardlink, how can the FS know when it can delete the file?

If you delete a hardlink, how can the FS know when it can delete the file?

If the refcount stored in the inode is zero

If you delete a hardlink, how can the FS know when it can delete the file?

If the refcount stored in the inode is zero

How can directories be implemented? What information is stored in them?

If you delete a hardlink, how can the FS know when it can delete the file?

If the refcount stored in the inode is zero

How can directories be implemented? What information is stored in them?

As a normal file with variable-length entries consisting of

- The filename
- The inode that represents that file

Which of the following data are typically stored in an inode?

- | | |
|---------------------------------|---|
| 1. filename | 8. name/location of hardlinks |
| 2. name of containing directory | 9. access rights |
| 3. file size | 10. timestamps (last access/modify) |
| 4. file type | 11. file contents |
| 5. number of symlinks to file | 12. ordered list of blocks occupied by file |
| 6. name/location of symlinks | |
| 7. number of hardlinks | |

Which of the following data are typically stored in an inode?

- | | |
|---------------------------------|---|
| 1. filename | 8. name/location of hardlinks |
| 2. name of containing directory | 9. access rights |
| 3. file size | 10. timestamps (last access/modify) |
| 4. file type | 11. file contents |
| 5. number of symlinks to file | 12. ordered list of blocks occupied by file |
| 6. name/location of symlinks | |
| 7. number of hardlinks | |

Which of the following data are typically stored in an inode?

- | | |
|---------------------------------|---|
| 1. filename | 8. name/location of hardlinks |
| 2. name of containing directory | 9. access rights |
| 3. file size | 10. timestamps (last access/modify) |
| 4. file type | 11. file contents |
| 5. number of symlinks to file | 12. ordered list of blocks occupied by file |
| 6. name/location of symlinks | |
| 7. number of hardlinks | |

Which of the following data are typically stored in an inode?

- | | |
|---------------------------------|---|
| 1. filename | 8. name/location of hardlinks |
| 2. name of containing directory | 9. access rights |
| 3. file size | 10. timestamps (last access/modify) |
| 4. file type | 11. file contents |
| 5. number of symlinks to file | 12. ordered list of blocks occupied by file |
| 6. name/location of symlinks | |
| 7. number of hardlinks | |

Which of the following data are typically stored in an inode?

- | | |
|---------------------------------|---|
| 1. filename | 8. name/location of hardlinks |
| 2. name of containing directory | 9. access rights |
| 3. file size | 10. timestamps (last access/modify) |
| 4. file type | 11. file contents |
| 5. number of symlinks to file | 12. ordered list of blocks occupied by file |
| 6. name/location of symlinks | |
| 7. number of hardlinks | |

Which of the following data are typically stored in an inode?

- | | |
|---------------------------------|---|
| 1. filename | 8. name/location of hardlinks |
| 2. name of containing directory | 9. access rights |
| 3. file size | 10. timestamps (last access/modify) |
| 4. file type | 11. file contents |
| 5. number of symlinks to file | 12. ordered list of blocks occupied by file |
| 6. name/location of symlinks | |
| 7. number of hardlinks | |

Which of the following data are typically stored in an inode?

- | | |
|---------------------------------|---|
| 1. filename | 8. name/location of hardlinks |
| 2. name of containing directory | 9. access rights |
| 3. file size | 10. timestamps (last access/modify) |
| 4. file type | 11. file contents |
| 5. number of symlinks to file | 12. ordered list of blocks occupied by file |
| 6. name/location of symlinks | |
| 7. number of hardlinks | |

Which of the following data are typically stored in an inode?

- | | |
|---------------------------------|---|
| 1. filename | 8. name/location of hardlinks |
| 2. name of containing directory | 9. access rights |
| 3. file size | 10. timestamps (last access/modify) |
| 4. file type | 11. file contents |
| 5. number of symlinks to file | 12. ordered list of blocks occupied by file |
| 6. name/location of symlinks | |
| 7. number of hardlinks | |

Which of the following data are typically stored in an inode?

- | | |
|---------------------------------|---|
| 1. filename | 8. name/location of hardlinks |
| 2. name of containing directory | 9. access rights |
| 3. file size | 10. timestamps (last access/modify) |
| 4. file type | 11. file contents |
| 5. number of symlinks to file | 12. ordered list of blocks occupied by file |
| 6. name/location of symlinks | |
| 7. number of hardlinks | |

Which of the following data are typically stored in an inode?

- | | |
|---------------------------------|---|
| 1. filename | 8. name/location of hardlinks |
| 2. name of containing directory | 9. access rights |
| 3. file size | 10. timestamps (last access/modify) |
| 4. file type | 11. file contents |
| 5. number of symlinks to file | 12. ordered list of blocks occupied by file |
| 6. name/location of symlinks | |
| 7. number of hardlinks | |

Which of the following data are typically stored in an inode?

- | | |
|---------------------------------|---|
| 1. filename | 8. name/location of hardlinks |
| 2. name of containing directory | 9. access rights |
| 3. file size | 10. timestamps (last access/modify) |
| 4. file type | 11. file contents |
| 5. number of symlinks to file | 12. ordered list of blocks occupied by file |
| 6. name/location of symlinks | |
| 7. number of hardlinks | |

Which of the following data are typically stored in an inode?

- | | |
|---------------------------------|---|
| 1. filename | 8. name/location of hardlinks |
| 2. name of containing directory | 9. access rights |
| 3. file size | 10. timestamps (last access/modify) |
| 4. file type | 11. file contents |
| 5. number of symlinks to file | 12. ordered list of blocks occupied by file |
| 6. name/location of symlinks | |
| 7. number of hardlinks | |

Which of the following data are typically stored in an inode?

- | | |
|---------------------------------|---|
| 1. filename | 8. name/location of hardlinks |
| 2. name of containing directory | 9. access rights |
| 3. file size | 10. timestamps (last access/modify) |
| 4. file type | 11. file contents |
| 5. number of symlinks to file | 12. ordered list of blocks occupied by file |
| 6. name/location of symlinks | |
| 7. number of hardlinks | |

Virtual File System

What is the purpose of the VFS layer in an OS?

What is the purpose of the VFS layer in an OS?

Abstraction. Provide a high-level API that works no matter what file system you use (XFS, ext4, NTFS, SMB, IliasFUSE, ...).

What is the purpose of the VFS layer in an OS?

Abstraction. Provide a high-level API that works no matter what file system you use (XFS, ext4, NTFS, SMB, IliasFUSE, ...).

What are some drawbacks / problems?

What is the purpose of the VFS layer in an OS?

Abstraction. Provide a high-level API that works no matter what file system you use (XFS, ext4, NTFS, SMB, IliasFUSE, ...).

What are some drawbacks / problems?

- Lowest common denominator: Might hide special FS features

What is the purpose of the VFS layer in an OS?

Abstraction. Provide a high-level API that works no matter what file system you use (XFS, ext4, NTFS, SMB, IliasFUSE, ...).

What are some drawbacks / problems?

- Lowest common denominator: Might hide special FS features
- ⇒ Allow accessing underlying FS (e.g. via `ioctl`)
- ⇒ Lose compatibility with other file systems

What is the purpose of the VFS layer in an OS?

Abstraction. Provide a high-level API that works no matter what file system you use (XFS, ext4, NTFS, SMB, IliadFUSE, ...).

What are some drawbacks / problems?

- Lowest common denominator: Might hide special FS features
- ⇒ Allow accessing underlying FS (e.g. via `ioctl`)
- ⇒ Lose compatibility with other file systems

What happens when you mount a filesystem?

E.g. `mount /dev/sda1 /mnt.`

What is the purpose of the VFS layer in an OS?

Abstraction. Provide a high-level API that works no matter what file system you use (XFS, ext4, NTFS, SMB, IliadFUSE, ...).

What are some drawbacks / problems?

- Lowest common denominator: Might hide special FS features
- ⇒ Allow accessing underlying FS (e.g. via `ioctl`)
- ⇒ Lose compatibility with other file systems

What happens when you mount a filesystem?

E.g. `mount /dev/sda1 /mnt`.

Makes files from that file system accessible. It will make the given path (e.g. `/mnt`) the *root* of the new file system. Any access to `/tmp/*` is stripped of the `/tmp/` prefix and then searched in the mounted file system.

FUSE

File System in **US**erspace.

- Allows users to write their own file system without writing kernel code

FUSE

File System in **US**erspace.

- Allows users to write their own file system without writing kernel code
- Runs as an unprivileged user process

FUSE

File System in **US**erspace.

- Allows users to write their own file system without writing kernel code
- Runs as an unprivileged user process
- Is really awesome! You suddenly can use all your normal tools on whatever the FUSE filesystem exposes!

File System Cache

Is data persisted after a call to `write`? If not, why and when?

Is data persisted after a call to `write`? If not, why and when?

- Data is probably buffered by libc first

Is data persisted after a call to `write`? If not, why and when?

- Data is probably buffered by libc first
- Even after that, it doesn't directly get written to disk. Enter...

Is data persisted after a call to `write`? If not, why and when?

- Data is probably buffered by libc first
- Even after that, it doesn't directly get written to disk. Enter...The *File System Cache*!

Is data persisted after a call to `write`? If not, why and when?

- Data is probably buffered by libc first
- Even after that, it doesn't directly get written to disk. Enter...The *File System Cache*!
- Why? Writing to disk every time is *reeeaaally slow*, so it is probably cached in memory and flushed in chunks

Is data persisted after a call to `write`? If not, why and when?

- Data is probably buffered by libc first
- Even after that, it doesn't directly get written to disk. Enter...The *File System Cache*!
- Why? Writing to disk every time is *reeeaaally slow*, so it is probably cached in memory and flushed in chunks

The standard question: Why is a cache even helpful?

Is data persisted after a call to **write**? If not, why and when?

- Data is probably buffered by libc first
- Even after that, it doesn't directly get written to disk. Enter...The *File System Cache*!
- Why? Writing to disk every time is *reeeaaally slow*, so it is probably cached in memory and flushed in chunks

The standard question: Why is a cache even helpful?

- Temporal and spatial locality

Is data persisted after a call to `write`? If not, why and when?

- Data is probably buffered by libc first
- Even after that, it doesn't directly get written to disk. Enter...The *File System Cache*!
- Why? Writing to disk every time is *reeeaaally slow*, so it is probably cached in memory and flushed in chunks

The standard question: Why is a cache even helpful?

- Temporal and spatial locality
- When exactly is it flushed to disk then? Is that even important to know/control?

Is data persisted after a call to `write`? If not, why and when?

- Data is probably buffered by libc first
- Even after that, it doesn't directly get written to disk. Enter...The *File System Cache*!
- Why? Writing to disk every time is *reeeaaally slow*, so it is probably cached in memory and flushed in chunks

The standard question: Why is a cache even helpful?

- Temporal and spatial locality
- When exactly is it flushed to disk then? Is that even important to know/control?
- Using `flush` (for buffers) and `fsync` (for the page cache) and after some time by a daemon

Let's talk about the page cache

Userspace buffer



Page cache




Physical disk



Let's talk about the page cache

Userspace buffer



Page cache




Physical disk



Let's talk about the page cache

Userspace buffer



write


Page cache




Physical disk




Let's talk about the page cache



Let's talk about the page cache



Let's talk about the page cache



How large would you make it? Do you give it a fixed size or let it grow?

How large would you make it? Do you give it a fixed size or let it grow?

Fix Easy to implement

How large would you make it? Do you give it a fixed size or let it grow?

Fix Easy to implement

Fix Can give hard guarantees (real-time systems anyone?)

Sizing the page cache

How large would you make it? Do you give it a fixed size or let it grow?

Fix Easy to implement

Fix Can give hard guarantees (real-time systems anyone?)

Fix Can't adapt to different workloads

Sizing the page cache

How large would you make it? Do you give it a fixed size or let it grow?

Fix Easy to implement

Fix Can give hard guarantees (real-time systems anyone?)

Fix Can't adapt to different workloads

When do you load data in the page cache?

Sizing the page cache

How large would you make it? Do you give it a fixed size or let it grow?

Fix Easy to implement

Fix Can give hard guarantees (real-time systems anyone?)

Fix Can't adapt to different workloads

When do you load data in the page cache?

Read-ahead: Read more data in the cache than requested

Sizing the page cache

How large would you make it? Do you give it a fixed size or let it grow?

Fix Easy to implement

Fix Can give hard guarantees (real-time systems anyone?)

Fix Can't adapt to different workloads

When do you load data in the page cache?

Read-ahead: Read more data in the cache than requested

- + Good sequential performance

Sizing the page cache

How large would you make it? Do you give it a fixed size or let it grow?

Fix Easy to implement

Fix Can give hard guarantees (real-time systems anyone?)

Fix Can't adapt to different workloads

When do you load data in the page cache?

Read-ahead: Read more data in the cache than requested

- + Good sequential performance
- + Improved throughput if files are sequential

Sizing the page cache

How large would you make it? Do you give it a fixed size or let it grow?

Fix Easy to implement

Fix Can give hard guarantees (real-time systems anyone?)

Fix Can't adapt to different workloads

When do you load data in the page cache?

Read-ahead: Read more data in the cache than requested

- + Good sequential performance
- + Improved throughput if files are sequential
- Wasted time and memory if not needed

How could you implement mmap with the file system cache?

How could you implement `mmap` with the file system cache?

Just map the file system cache in the process's memory (and handle faults!)

⇒ Acts as a shared memory segment

You synchronize accesses yourself :(

How could you implement mmap with the file system cache?

Just map the file system cache in the process's memory (and handle faults!)

⇒ Acts as a shared memory segment

You synchronize accesses yourself :(

read() and write() and the file system cache

How could you implement `mmap` with the file system cache?

Just map the file system cache in the process's memory (and handle faults!)

⇒ Acts as a shared memory segment

You synchronize accesses yourself :(

`read()` and `write()` and the file system cache

- `read()`: Copy data from cache to your application buffers

How could you implement `mmap` with the file system cache?

Just map the file system cache in the process's memory (and handle faults!)

⇒ Acts as a shared memory segment

You synchronize accesses yourself :(

`read()` and `write()` and the file system cache


- `read()`: Copy data from cache to your application buffers
- `write()`: Copy data from your application buffers to the file system cache

File system synchronizes access!

Modern File Systems


Your computer crashes (e.g. your power fails). What horrible death does your file system die?

A fun game for the whole family and your „why would I need backups“ crowd




Your computer crashes (e.g. your power fails). What horrible death does your file system die?

A fun game for the whole family and your „why would I need backups“ crowd




Your computer crashes (e.g. your power fails). What horrible death does your file system die?

A fun game for the whole family and your „why would I need backups“ crowd




Your computer crashes (e.g. your power fails). What horrible death does your file system die?

A fun game for the whole family and your „why would I need backups“ crowd




The inode has a refcount of 1 but is no longer referenced!


Where is Data?



Where is Data?




Where is Data?




Both files might point to this block!


Where is Data?



Where is Data?




Where is Data?



Block still referenced but marked free

Where is Data?



Block still referenced but marked free

Or multiple directory entries with the same name, as another problem

What might be the outcome?

```
1 // create a file
2 echo "Hey" > a.txt
3 // And a second
4 echo "Hello" > b.txt
5 // CRASH
```


What might be the outcome?

```
1 // create a file
2 echo "Hey" > a.txt
3 // And a second
4 echo "Hello" > b.txt
5 // CRASH
```

- Both files exist

What might be the outcome?

```
1 // create a file
2 echo "Hey" > a.txt
3 // And a second
4 echo "Hello" > b.txt
5 // CRASH
```

- Both files exist
- No files exist

What might be the outcome?

```
1 // create a file
2 echo "Hey" > a.txt
3 // And a second
4 echo "Hello" > b.txt
5 // CRASH
```

- Both files exist
- No files exist
- Only a exists

What might be the outcome?

```
1 // create a file
2 echo "Hey" > a.txt
3 // And a second
4 echo "Hello" > b.txt
5 // CRASH
```

- Both files exist
- No files exist
- Only **a** exists
- Only **b** exists !!

Reordering is *allowed* (unless you take precautions)

How could you detect those inconsistencies?

Use the **fsck** (file system consistency check) program. And what could that do?

How could you detect those inconsistencies?

Use the **fsck** (file system consistency check) program. And what could that do?

- Check whether all data blocks are referenced by *exactly one* inode

How could you detect those inconsistencies?

Use the **fsck** (file system consistency check) program. And what could that do?

- Check whether all data blocks are referenced by *exactly one* inode
- Do all directories contain valid entries (and just one with a given name)

How could you detect those inconsistencies?

Use the **fsck** (file system consistency check) program. And what could that do?

- Check whether all data blocks are referenced by *exactly one* inode
- Do all directories contain valid entries (and just one with a given name)
- Is every directory / file connected to the directory tree

How could you detect those inconsistencies?

Use the **fsck** (file system consistency check) program. And what could that do?

- Check whether all data blocks are referenced by *exactly one* inode
- Do all directories contain valid entries (and just one with a given name)
- Is every directory / file connected to the directory tree
- Is the refcount consistent with the number of hardlinks? What do you do if not?

refcount > 0 but no links?

How could you detect those inconsistencies?

Use the **fsck** (file system consistency check) program. And what could that do?

- Check whether all data blocks are referenced by *exactly one* inode
- Do all directories contain valid entries (and just one with a given name)
- Is every directory / file connected to the directory tree
- Is the refcount consistent with the number of hardlinks? What do you do if not?

refcount > 0 but no links? \Rightarrow Attach to **lost+found**

Update the refcount to be consistent

How could you detect those inconsistencies?

Use the **fsck** (file system consistency check) program. And what could that do?

- Check whether all data blocks are referenced by *exactly one* inode
- Do all directories contain valid entries (and just one with a given name)
- Is every directory / file connected to the directory tree
- Is the refcount consistent with the number of hardlinks? What do you do if not?

$refcount > 0$ but no links? \Rightarrow Attach to **lost+found**

Update the refcount to be consistent

Can this fix your application data?

How could you detect those inconsistencies?

Use the **fsck** (file system consistency check) program. And what could that do?

- Check whether all data blocks are referenced by *exactly one* inode
- Do all directories contain valid entries (and just one with a given name)
- Is every directory / file connected to the directory tree
- Is the refcount consistent with the number of hardlinks? What do you do if not?

refcount > 0 but no links? \Rightarrow Attach to **lost+found**

Update the refcount to be consistent


Can this fix your application data?

No! It just tries to keep the filesystem internally consistent, it won't find corrupted data blocks

General principle

Walk over the journal and execute any outstanding entries.

Let's crash



What to do

The happy path, everything's nice

General principle

Walk over the journal and execute any outstanding entries.

Let's crash



What to do

⇒ We didn't write anything!


⇒ Operation **failed**

Journal File System - Recovery

General principle

Walk over the journal and execute any outstanding entries.

Let's crash



What to do


- ⇒ Invalid checksum
- ⇒ Skip entry
- ⇒ Operation **failed**

Journal File System - Recovery

General principle

Walk over the journal and execute any outstanding entries.

Let's crash



What to do


- ⇒ Non-terminated journal entry
- ⇒ Retry and complete operation
- ⇒ Operation **successful**

Journal File System - Recovery

General principle

Walk over the journal and execute any outstanding entries.

Let's crash



What to do

- ⇒ Consistent state
- ⇒ Execute operation again
- ⇒ Operation **successful**

Journal File System - Recovery

General principle

Walk over the journal and execute any outstanding entries.

Let's crash



What to do

- ⇒ Consistent state
- ⇒ Don't do anything
- ⇒ Operation **successful**

Physical vs logical logging

What data could you log in the journal?

Physical vs logical logging

What data could you log in the journal?

- **Logical logging:** Store a high level entry (like: „rename a to b“)

Physical vs logical logging

What data could you log in the journal?

- **Logical logging:** Store a high level entry (like: „rename a to b“)
- **Physical logging:** Store the file system blocks that will be modified

What might be problems with the journal presented before?

What might be problems with the journal presented before?

You write *every block twice*!

What might be problems with the journal presented before?

You write *every block twice*! How could you make that less painful?

What might be problems with the journal presented before?

You write *every block twice*! How could you make that less painful? Only journal *metadata*.

What might be problems with the journal presented before?

You write *every block twice*! How could you make that less painful? Only journal *metadata*.

How could you design a safe system that only needs to write once?

What might be problems with the journal presented before?

You write *every block twice*! How could you make that less painful? Only journal *metadata*.

How could you design a safe system that only needs to write once?

A Log Structured File System (Copy on Write). What is the core idea there?

What might be problems with the journal presented before?

You write *every block twice*! How could you make that less painful? Only journal *metadata*.

How could you design a safe system that only needs to write once?

A Log Structured File System (Copy on Write). What is the core idea there?

Only write to unused blocks.

What might be problems with the journal presented before?

You write *every block twice*! How could you make that less painful? Only journal *metadata*.

How could you design a safe system that only needs to write once?

A Log Structured File System (Copy on Write). What is the core idea there?

Only write to unused blocks.

Then you can *atomically* update indexing datastructures to point to the new blocks. If it crashes, you either have *all* of the new state or *exactly* the old state.

Das letzte Tut ist *nächste Woche*

Das wird ein Wiederholungstut von Sachen, die euch noch irgendwo unklar sind oder über die ihr nochmal reden wollt.

⇒ Daher bitte bis nächste Woche Fragen an mich per Mail
peter.bohner@student.kit.edu